

Loop Engineering: The Operating Model Behind Modern Agentic AI Systems

How Claude Code, ChatGPT, Replit, Gemini, Devin, and Cursor design, run, and scale continuous AI loops.

OBSERVE → ACT → VERIFY → ADJUST → REPEAT



Lamhot Siagian

with AI Engineering Insider

THE AGENTIC AI SYSTEMS SERIES

Loop Engineering

The Operating Model Behind Modern Agentic AI Systems

How Claude Code, ChatGPT, Replit, Gemini, Devin, and Cursor design, run, and scale continuous AI loops.

A practitioner's reference for AI Engineers · ML Engineers · Agentic AI Engineers · Software Engineers · AI Architects

Lamhot Siagian

with AI Engineering Insider

aiengineeringinsider.substack.com

First Edition · 2026

Loop Engineering: The Operating Model Behind Modern Agentic AI Systems

How Claude Code, ChatGPT, Replit, Gemini, Devin, and Cursor design, run, and scale continuous AI loops.

Author: **Lamhot Siagian**, with AI Engineering Insider.

Publisher: AI Engineering Insider — aiengineeringinsider.com

First Edition (Version 1.0), 2026.

Companion code: github.com/lamhotsiagian/loop-engineering

© 2026 Lamhot Siagian and AI Engineering Insider. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means without the prior written permission of the author, except for brief quotations in reviews and for scholarly citation with proper attribution.

Disclaimer. This book is provided for educational purposes. The system-design descriptions of third-party products in the bonus chapter are reconstructions based on public information and observable behaviour, not internal specifications; product names and trademarks belong to their respective owners. The author and publisher make no warranties regarding completeness or accuracy and accept no liability for losses arising from use of this material. Verify vendor-specific details against current official documentation.

How to cite this book

Siagian, L. (2026). *Loop Engineering: The operating model behind modern agentic AI systems* (1st ed.). AI Engineering Insider.

BibTeX:

```
@book{siagian2026loop,  
  author = {Siagian, Lamhot},  
  title = {Loop Engineering: The Operating Model Behind Modern Agentic AI Systems},  
  year = {2026},  
  edition = {1st},  
  publisher = {AI Engineering Insider}  
}
```

Preface

Subscribe → aiengineeringinsider.substack.com/subscribe

Sometime in the last two years, the best engineers stopped typing the next instruction.

They noticed that the valuable signal in real work — the failing test, the type error, the reviewer’s objection, the server that won’t start — almost never arrives with the first response. It arrives *after*: after the model acts, after the code runs, after reality answers back. A single prompt, however well crafted, has already stopped listening by then. So these engineers stopped holding the model turn by turn and started building a small system around it — one that finds the work, hands it to the model, checks the result, records what happened, and decides what to do next. Then they let that system drive.

That system is a loop, and designing it is its own discipline. This book is about that discipline.

The thesis

The loop is the operating model of modern agentic AI. Not the model weights, not the prompt, not the framework — the loop: the cycle of *observe, act, verify, adjust, repeat* that turns a clever text generator into something that can pursue a goal across many steps and correct itself along the way. Every agentic product worth studying, from autonomous software engineers to metadata-governance agents, is at its core the same primitive: a model calling tools in a loop until a verifier is satisfied. What separates a demo from a system you can trust in production is never the cleverness of the model. It is the care taken with everything *around* it.

The four surfaces

If there is one idea to carry through the whole book, it is this. Every loop that survives contact with production gets four surfaces right, and every loop that fails in the field failed on at least one of them:

1. **Grounding** — the loop acts in a real environment, so its observations are true consequences of true actions, not guesses.
2. **Verification** — an external source of truth, not the model’s own confidence, decides whether the work is done.
3. **Budget** — explicit ceilings on iterations, tokens, time, and cost, so a stuck loop stops instead of spending without end.
4. **Human boundary** — a designed point where human judgment re-enters, for the decisions no automated check can make.

These four surfaces are the spine of this book. The chapters build them up one at a time; the bonus part shows ten shipping products that are, almost entirely, different answers to the same four questions: *where does the verifier live, what grounds the loop in reality, what bounds its cost, and how does a human get back in?*

Who this is for, and how to read it

This is a practitioner’s reference for the people who build these systems: AI and ML engineers, agentic-AI and software engineers, and the architects who have to make them reliable at scale. It assumes you can read Python and have used an LLM in anger; it does not assume you have built a loop before.

Read it however suits you. Part I is the foundation and rewards a linear read. After that, each part stands on its own — reach for verification when quality is slipping, for the scaling part when one agent is no longer enough, for the economics chapter when the bill arrives. Every chapter ends with ten exercises and, a page later, worked answers, so you can test your understanding before checking it. A complete, runnable companion project, [loop_engineering_lab](https://github.com/lamhotsiagian/loop-engineering), accompanies the book and is free to download:

github.com/lamhotsiagian/loop-engineering

It runs with no API key — the model is a deterministic mock — so every example in these pages is something you can execute, break, and modify yourself. Code, after all, is the one part of this subject that refuses to be hand-waved. That is exactly why loops were born there first.

Lamhot Siagian

2026

How to Use This Book

This book is organized as a reference, not a narrative you must read start to finish. Part I is the foundation; Parts II through V can be read in order or pulled individually once you have the core loop down. Six recurring box types flag specific kinds of content:

★ Key Insight

A core idea worth remembering. If you forget everything else in the section, remember this.

! Pitfall / Warning

A pitfall that shows up repeatedly in production systems. These are failure modes, not theoretical risks.

✓ Best Practice

A concrete, actionable recommendation you can apply directly.

From the Field

An opinionated, first-person aside: a hard-won lesson, a hot take, or a war story from running these systems in production.

Definition

A precise definition of a term used consistently throughout the book.

Project Code — `loop_engineering_lab`

A pointer to runnable reference code that accompanies a chapter. The complete companion project, `loop_engineering_lab`, is available to download at github.com/lamhotsiagian/loop-engineering.

Contents

Preface	iv
Part I Foundations	1
1 From Prompting to Looping	2
1.1 Loop vs. Chain: The Foundational Distinction	2
1.2 What Is a Loop? Five-Stage Anatomy	3
1.3 Prompt Engineering vs. Loop Engineering	4
1.4 Why Coding Is the Proving Ground	4
1.5 Intellectual Lineage	4
1.6 Reference Implementation: The Five Stages in Code	5
1.7 Exercises	5
1.8 Chapter Summary	7
1.9 Answers	8
2 The Core Agent Loop	11
2.1 Observe, Reason, Act, Observe	11
2.2 The ReAct Pattern	11
2.3 Plan-and-Execute vs. Reactive Loops	12
2.4 The Agent Loop as a State Machine	12
2.5 Two Nested Loops	13
2.6 Synchronous vs. Asynchronous Execution	13
2.7 Plan-and-Execute vs. Reactive: A Decision Table	13
2.8 Reference Implementation: The Loop Driver	14
2.9 Exercises	15
2.10 Chapter Summary	16
2.11 Answers	17
3 Loop Control Structures	20
3.1 Termination Conditions	20
3.2 Stopping Rules as First-Class Design Requirements	20
3.3 Loop Budgets	21
3.4 Early-Exit and Short-Circuit Logic	21
3.5 Recursive vs. Iterative Structures	21
3.6 Nested Loops	22
3.7 Loop Re-Entry and Resumability	22
3.8 Budget Dimensions at a Glance	23
3.9 Reference Implementation: Budgets and Stopping Rules	23
3.10 Exercises	24
3.11 Chapter Summary	25
3.12 Answers	26

Part II	State, Reasoning, and Tools Inside the Loop	29
4	State and Memory Across Iterations	30
4.1	The Context Window as Loop Memory	30
4.2	Scratchpad Patterns	30
4.3	State Compaction	31
4.4	Context Window Overflow Handling	31
4.5	Persistent vs. Ephemeral State	31
4.6	State Serialization for Resumable Loops	31
4.7	Tracking What's Been Tried	31
4.8	External Memory as the Real Persistence Layer	32
4.9	Skills as Externalized Project Intent	32
4.10	Overflow Strategies Compared	33
4.11	Reference Implementation: A Compactable Scratchpad	33
4.12	Exercises	34
4.13	Chapter Summary	35
4.14	Answers	36
5	Reasoning Within the Loop	39
5.1	Intra-Step vs. Loop-Level Reasoning	39
5.2	Reflection and Self-Critique	39
5.3	Plan Revision Mid-Loop	40
5.4	Tree-of-Thought and Graph-of-Thought	40
5.5	Verifier-in-the-Loop	40
5.6	Self-Consistency Sampling	41
5.7	Reasoning Techniques Compared	41
5.8	Reference Implementation: Framed Self-Critique	41
5.9	Exercises	42
5.10	Chapter Summary	43
5.11	Answers	44
6	Tool Use Within the Loop	47
6.1	The Tool-Call Cycle	47
6.2	Parallel vs. Sequential Tool Calls	47
6.3	Tool Error Handling and Retry Logic	48
6.4	Tool Result Processing Before Re-Injection	48
6.5	Tool Selection Logic	48
6.6	Idempotency	49
6.7	The Tool Set as the Ceiling on Loop Quality	49
6.8	Failure-Response Policy at a Glance	50
6.9	Reference Implementation: Structured Tool Feedback	50
6.10	Exercises	51
6.11	Chapter Summary	52
6.12	Answers	53
Part III	Failure, Verification, and Robustness	56
7	Failure Handling and Loop Robustness	57
7.1	Infinite Loop and Repetition Detection	57
7.2	Stuck-State Detection	57
7.3	Retry Strategies	57
7.4	Graceful Degradation	58
7.5	Hallucinated and Malformed Actions	58

7.6	Timeout and Deadlock Prevention	58
7.7	Circuit Breakers	58
7.8	Error Taxonomy: Recoverable vs. Hard Blocker	59
7.9	Named Failure Modes	59
7.10	Failure-Handling Decision Flow	60
7.11	Reference Implementation: Circuit Breaker and Stuck Detection	60
7.12	Exercises	61
7.13	Chapter Summary	62
7.14	Answers	63
8	The Verification Loop	66
8.1	Verification as a Distinct Architectural Layer	66
8.2	Grader Design	66
8.3	Rubric-Based Scoring	67
8.4	Feedback Injection	67
8.5	The Maker/Checker Split	67
8.6	The Cost/Latency Tradeoff	68
8.7	Worked Example	68
8.8	Grader Types Compared	69
8.9	Reference Implementation: A Rubric Grader as Verifier	69
8.10	Exercises	70
8.11	Chapter Summary	71
8.12	Answers	72
9	Loop Pattern Catalog	75
9.1	The Retry Loop	75
9.2	The Plan-Execute-Verify Loop	75
9.3	The Explore-Narrow Loop	76
9.4	The Test-Driven Agent Loop	76
9.5	The Compiler-Driven Loop	76
9.6	The Review-Driven Loop	76
9.7	The Runtime Debugging Loop	77
9.8	The Product Iteration Loop	77
9.9	The Human-in-the-Loop Pattern	77
9.10	Choosing a Pattern	78
9.11	Patterns Compose: A Real Bug-Fix Pipeline	78
9.12	Reference Implementation: Patterns Are Configuration	78
9.13	Exercises	79
9.14	Chapter Summary	80
9.15	Answers	81
Part IV	Scaling Beyond One Agent	84
10	Multi-Agent Loop Architectures	85
10.1	Orchestrator-Worker Loops	85
10.2	Peer-to-Peer Agent Loops	85
10.3	Hierarchical Loop Decomposition	85
10.4	The Planner / Parallel-Executors / Reviewer Pattern	86
10.5	Loop Synchronization	86
10.6	Shared State vs. Isolated State	86
10.7	Worktree-Based Isolation	87
10.8	Coordination Topologies Compared	87

10.9	Reference Implementation: Orchestrator with Isolated Workers	87
10.10	Exercises	88
10.11	Chapter Summary	89
10.12	Answers	90
11	Event-Driven and Autonomous Loops	93
11.1	The Agent as a Persistent Background Component	93
11.2	Trigger Types	93
11.3	The Heartbeat Pattern	94
11.4	Time-Based vs. Condition-Based Looping	94
11.5	Decoupling Invocation from Interaction	94
11.6	Infrastructure Concerns	94
11.7	Trigger Semantics Compared	95
11.8	Reference Implementation: An Idempotent Heartbeat	95
11.9	Exercises	96
11.10	Chapter Summary	97
11.11	Answers	98
12	The Hill-Climbing Loop	101
12.1	Traces as a First-Class Artifact	101
12.2	Trace-Analysis Agents	101
12.3	The Return Arrow Reaches Inside	101
12.4	What Gets Improved	102
12.5	Loop Stacking as a Design Principle	102
12.6	Human Review of Harness Changes	103
12.7	The Loop Stack at a Glance	103
12.8	Reference Implementation: Mining Traces for Harness Signal	103
12.9	Exercises	104
12.10	Chapter Summary	105
12.11	Answers	106
Part V	Observability, Safety, and Economics	109
13	Observability and Evaluation of Loops	110
13.1	Loop Tracing	110
13.2	Step-Level Evaluation	110
13.3	Trajectory Evaluation	111
13.4	Identifying Loop Pathologies	111
13.5	A/B Testing Loop Strategies	111
13.6	Replay and Debugging Tools	111
13.7	Loop Robustness Testing	112
13.8	The Observability Pipeline	112
13.9	Evaluation Granularity Compared	113
13.10	Reference Implementation: Runaway Rate as a First-Class Metric	113
13.11	Exercises	113
13.12	Chapter Summary	115
13.13	Answers	116
14	Safety, Control, and Human Oversight	119
14.1	Human-in-the-Loop by Architectural Layer	119
14.2	Permission Boundaries Per Iteration	120
14.3	Sandboxing Iterative Tool Execution	120
14.4	Guardrails Re-Evaluated at Every Step	120

14.5	Audit Logging Per Iteration	120
14.6	Kill Switches	120
14.7	The Orchestration Tax	121
14.8	Comprehension Debt	121
14.9	Cognitive Surrender	121
14.10	Where Human Oversight Attaches	122
14.11	Kill Switch vs. Circuit Breaker	122
14.12	Reference Implementation: Kill Switch and Per-Step Guard	122
14.13	Exercises	123
14.14	Chapter Summary	124
14.15	Answers	125
15	Cost, Latency, and Efficiency Engineering	128
15.1	Cost-Per-Iteration Accounting	128
15.2	Latency Optimization	128
15.3	Adaptive Loop Depth: Model Routing	129
15.4	Batching and Parallelization	129
15.5	Context Pruning	129
15.6	Tool-Call Budgets as a Failure Signal	129
15.7	Putting It Together: A Cost-Aware Loop	129
15.8	A Cost-Aware Loop	130
15.9	Cost Levers Compared	130
15.10	Reference Implementation: Routing by Difficulty	131
15.11	Exercises	131
15.12	Chapter Summary	132
15.13	Answers	134
16	Loop Engineering as an Organizational Discipline	137
16.1	Repository Instructions	137
16.2	Branch and Pull Request Rules for Agent-Generated Changes	137
16.3	Safe Defaults Around Destructive Commands and Credentials	138
16.4	Standard, Reusable Loops as Organizational Assets	138
16.5	CI Checks Both Agents and Humans Trust	138
16.6	Clear Ownership for Final Decisions	138
16.7	Capturing Reusable Loops	139
16.8	The Loop Engineering Maturity Model	139
16.9	The Future of Loop Engineering	140
16.10	From Personal Script to Team Standard	140
16.11	Individual Skill vs. Repository Instruction	141
16.12	Reference Implementation: Reusable Loops as Org Assets	141
16.13	Exercises	141
16.14	Chapter Summary	143
16.15	Answers	144
Part VI	Loop Engineering in Production	147
17	Bonus: Ten Companies Running Loop Engineering in Production	148
17.1	The Ten at a Glance	149
17.2	Cognition AI — Devin	149
17.3	Anthropic — Claude Code	150
17.4	Anysphere — Cursor Agent Mode	152
17.5	OpenAI — Codex & Code Interpreter Loops	153

17.6 Factory AI — Droids	154
17.7 Sweep AI — Sweep	156
17.8 Replit — Replit Agent	157
17.9 Google — Gemini Code Assist & Developer Agents	158
17.10 Atlan — Atlan AI Agents	159
17.11 OpenHands — OpenHands Framework	160
17.12 Cross-Cutting Lessons	162
The Loop Engineering Pre-Launch Checklist	164
Loop Pattern Quick-Reference	166
17.13 Matching Patterns to Feedback Signals	166
Failure Mode Glossary	168
Framework and Primitive Mapping	170
17.14 Core Concepts to Framework Primitives	170
17.15 Building Without a Framework	171
References	172
About the Author	175
Index	176

PART I

Foundations

Before a system can scale, coordinate, or improve itself, one loop has to work. Part I builds that primitive from the ground up: why single-shot prompting hit a ceiling, what a loop actually is, and how to decide when it continues, stops, or escalates. Everything in the rest of the book — verification, multi-agent coordination, autonomy, self-improvement — is stacked on the three chapters here. Get this foundation wrong and nothing above it will hold.

From Prompting to Looping

Subscribe → aiengineeringinsider.substack.com/subscribe

We stopped writing prompts and started building the systems that write them.

Picture an engineer at the end of a long day, watching a terminal. A test is failing. She types a description of the bug, an agent reads the stack trace, edits a file, and reruns the suite. It fails differently this time. The agent reads the new error, forms a fresh hypothesis, changes one line, and reruns. Green. She did not write the fix. She did not even write the second prompt. She built the thing that wrote the prompts — the small system that found the failure, handed it to the model, checked the result, and decided what to do next. Then she let it run.

That shift, from holding the tool to building the system that holds the tool, is what this book is about.

For roughly two years, the standard way to get useful work out of a coding agent was simple: write a good prompt, supply enough context, read what came back, and type the next instruction. The agent was a tool, and the engineer held it for every turn. It worked, within limits, and those limits are now the whole story. A single well-crafted prompt produces a single response, and a single response cannot run itself, test itself, notice it was wrong, and try again. The most valuable signal in real engineering work — the failing test, the type error, the reviewer's objection — arrives *after* that first response, exactly where a one-shot prompt has already stopped listening.

So the pattern is ending, and a different one is taking its place. The unit of work is no longer the prompt; it is the loop. The skill is no longer phrasing a request well; it is designing a system that pursues a goal across many steps, reads the consequences of its own actions, and corrects course without a human retyping the instruction each time. This chapter draws the line between the two, names the parts every loop is made of, and sets up the rest of the book — which is, start to finish, a practitioner's guide to building that system well.

Loop Engineering

The discipline of designing, operating, and improving the iterative cycles that let an AI system pursue a goal across multiple steps, rather than produce a single response. A loop engineer builds the system that prompts the agent, instead of prompting the agent directly.

The replacement is easy to state and hard to do well: stop prompting the agent turn by turn, and instead design a small system that finds the work, hands it to the agent, checks the result, records what was done, and decides the next action — then let that system drive the agent. The engineer's job moves up one level of abstraction: from authoring individual prompts to authoring the loop that authors prompts.

1.1 Loop vs. Chain: The Foundational Distinction

Before going further, it is worth fixing a distinction that gets blurred constantly in casual usage.

Chain

A fixed, linear sequence of steps: step A leads to step B leads to step C, in a predetermined order, regardless of intermediate results.

Loop

A dynamic, repeating cycle in which the outcome of one iteration determines whether and how the next iteration proceeds. The agent might go from step A to step B, discover that B failed, retry with a modified approach, and only then move forward — or it might stop entirely.

Chains are predictable and easy to trace, which makes them attractive for well-understood, low-uncertainty tasks. Loops are more flexible and far better suited to tasks where the right path is not known in advance: debugging an unfamiliar error, exploring an undocumented API, or making a code change in a codebase whose constraints are only partially visible at the start.




Figure 1.1: Chain versus loop. A chain runs straight through and cannot recover from a wrong early step; a loop feeds each observation back into the next iteration, so a failure becomes the input that corrects the next attempt.

! Pitfall / Warning

A pipeline built entirely from chained steps cannot recover from a wrong assumption made at step one. It will faithfully execute steps two through ten on top of a broken foundation. This is the single most common reason teams reach for loop engineering after a chain-based system quietly produces confidently wrong output.

1.2 What Is a Loop? Five-Stage Anatomy

Every agentic loop, regardless of domain, reduces to five stages:

1. **Intent** — the developer or system defines the target outcome.
2. **Context** — the agent gathers relevant code, documentation, errors, logs, and constraints.
3. **Action** — the agent edits files, runs commands, calls tools, or drafts a plan.
4. **Observation** — the system captures test results, compiler errors, runtime output, diffs, or review comments.
5. **Adjustment** — the agent updates its plan and repeats the loop until the work is accepted or blocked.

The value is never in any single stage. It is in *closing* the loop: a test failure is not merely an error message, it is new context; a type error is not merely a blocker, it is a signal that an assumption was wrong; a code review comment is another observation that drives the next action.

★ Key Insight

A strong first answer is not the same as a correct final change. In production code, the signal that actually matters often appears *after* the first action — a failing type check, a missing import, a broken test, a reviewer flagging an edge case. Loop engineering makes those signals part of the system instead of treating them as after-the-fact cleanup.

1.3 Prompt Engineering vs. Loop Engineering

These two disciplines are complementary, not competing. Clear instructions still improve the first plan and reduce unnecessary exploration — prompt engineering still matters inside a loop. The difference is what each discipline assumes and optimizes for.

Dimension	Prompt Engineering	Loop Engineering
Primary unit	A prompt or instruction	A repeated workflow
Goal	Better first response	Better final outcome
Feedback	Usually human evaluation	Tests, tools, diffs, logs, reviews, and humans
Failure mode	The model gives a weak answer	The system loops poorly, stops early, or ignores evidence
Best for	Isolated answers and generation	Real tasks with uncertainty

Table 1.1: Prompt engineering shapes the input to a model; loop engineering shapes the entire process around it.

1.4 Why Coding Is the Proving Ground

Coding is a naturally iterative domain. Even experienced engineers rarely write correct code on the first attempt — they run it, read the error, fix it, run it again. An AI coding agent that generates code once and stops is fundamentally limited: it cannot catch runtime errors, cannot adapt to environment-specific quirks, and cannot verify that what it produced actually works.

This is precisely why loop engineering matured first in coding agents and is now generalizing to every other agentic domain: research assistants, customer-support automation, data pipelines, and operations tooling all share the same underlying need to act, observe, and revise.

1.5 Intellectual Lineage

Loop engineering did not appear from nothing. It draws on several established lines of work:

- **OODA loop** (Observe–Orient–Decide–Act) — a military decision-cycle framework later adopted broadly in systems and control theory.
- **ReAct** (Reason + Act) — research from Princeton and Google showing that interleaving explicit reasoning steps with tool-use actions produces more reliable agent behavior than either reasoning or acting alone.
- **Plan-act-verify** — a pattern from classical software engineering and robotics, where a plan is generated, executed incrementally, and checked against the real world at each step.

What is genuinely new is not any single one of these ideas, but the fact that large language models can now execute all three roles — planner, actor, and verifier — inside the same system, often within the same conversation.

From the Field

The term “loop engineering” is recent and still settling, but the underlying practice is not new. Anyone who has built a CI/CD pipeline with automatic retries, or a control loop in robotics, has already done a version of this work. What changed is that the actor in the loop can now read a stack trace, form a hypothesis about the cause, and write the fix itself.

1.6 Reference Implementation: The Five Stages in Code

The five-stage anatomy is not just a mental model — it maps directly onto the data structures of the reference project. `loop_engineering_lab/core/state.py` encodes each stage as an explicit phase of a state machine, and the loop driver in `core/agent_loop.py` walks them in order, recording every transition so a run can be replayed.

loop_engineering_lab/core/state.py — the loop’s serialisable memory

```
class LoopPhase(str, Enum):
    """The states an agent loop can occupy (the five-stage anatomy)."""
    PLANNING = "planning" # Intent: define the target outcome
    GATHER_CONTEXT = "gather_context" # Context: pull code, errors, logs
    REASONING = "reasoning" # decide the next move
    ACTING = "acting" # Action: a tool call is in flight
    OBSERVING = "observing" # Observation: process a tool result
    VERIFYING = "verifying" # Adjustment: was the action correct?
    SUCCEEDED = "succeeded" # terminal -- few and explicit
    BLOCKED = "blocked"
    FAILED = "failed"

    @property
    def is_terminal(self) -> bool:
        return self in {LoopPhase.SUCCEEDED, LoopPhase.BLOCKED, LoopPhase.FAILED}
```

The crucial design choice is that *adjustment* is a first-class phase, not an afterthought. A chain has no OBSERVING or VERIFYING state; a loop is defined precisely by the edges that lead out of those two phases back into reasoning. In production, the difference between “the model wrote plausible code” and “the change actually works” is exactly the edge from OBSERVING back to REASONING when a test fails.

★ Key Insight

If you can draw your system as a directed graph and find no cycle, you have built a chain, no matter how many LLM calls it contains. Loop engineering begins the moment an observation can send control *backward*.

1.7 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 1.1

What is loop engineering, and how is it different from prompt engineering?

Hint. Think about what each optimises — a single input versus the whole process around the model.

Exercise 1.2

Explain the difference between a chain and a loop with a concrete production example.

Hint. Ask whether an early error can be detected and reversed; use the OCR-then-summarise example.

Exercise 1.3

Walk through the five-stage anatomy of a loop. Which stage carries the value?

Hint. The value is in closing the loop, not in any one stage.

Exercise 1.4

Why did loop engineering mature first in coding agents rather than, say, customer support?

Hint. Consider which domain has a cheap, automatic, trustworthy verifier.

Exercise 1.5

A teammate says “we already have a multi-step LangChain pipeline, so we do loop engineering.” How do you respond?

Hint. Look for a feedback edge — does any output send control backward?

Exercise 1.6

What intellectual lineage does loop engineering draw on, and what is genuinely new?

Hint. OODA, ReAct, plan-act-verify; what is new is one model playing all three roles.

Exercise 1.7

Debugging case: a loop produces confidently wrong output and never errors. Where do you look first?

Hint. A missing or ignored verifier lets confident output through.

Exercise 1.8

When is a chain the right choice and a loop the wrong one?

Hint. Pick the cheapest structure that meets the reliability bar; low uncertainty favours chains.

Exercise 1.9

How would you measure whether moving from a chain to a loop actually improved the system?

Hint. Compare cost per successful task on an offline eval set, not by feel.

Exercise 1.10

Where does prompt engineering still matter once you have adopted loop engineering?

Hint. It still shapes the first plan and the quality of tool-call arguments inside the loop.

1.8 Chapter Summary

- Loop engineering is the discipline of designing the system that prompts the agent, not the act of prompting it directly.
- Loops differ from chains in one essential way: the outcome of an iteration changes what happens next.
- Every loop reduces to five stages: Intent, Context, Action, Observation, Adjustment.
- Prompt engineering and loop engineering are complementary; the former optimizes a single input, the latter optimizes the system around repeated inputs and outputs.
- Coding is the domain where loop engineering matured first, because code is naturally self-verifying through tests, compilers, and runtime behavior.

The next chapter takes the five-stage anatomy introduced here and builds it out into a complete, concrete agent loop — the smallest unit every more advanced pattern in this book is built on top of.

1.9 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 1.1

What is loop engineering, and how is it different from prompt engineering?

Loop engineering is the discipline of designing, operating, and improving the iterative system that drives an agent toward a goal across many steps, rather than authoring the individual prompts the agent receives. Prompt engineering optimises a single input–output exchange: a better instruction yields a better first response. Loop engineering optimises the *process* wrapped around the model — how work is found, how context is assembled, how results are verified, how failures feed back, and when the system stops. The two are complementary: clearer prompts still reduce wasted exploration inside a loop. But they fail differently. A prompt fails when the model gives a weak answer; a loop fails when the system stops too early, ignores evidence from a failed verification, or burns its budget without converging. In practice the tell is whether an engineer reaches for “write a better prompt” or “add a verification step and feed the failure back” when the first attempt is wrong.

Answer 1.2

Explain the difference between a chain and a loop with a concrete production example.

A chain is a fixed linear sequence: extract → transform → load, executed in order regardless of intermediate results. A loop is a cycle in which the outcome of one iteration determines the next. Concretely: a document-processing pipeline that OCRs a PDF, summarises it, and emails the summary is a chain — if the OCR step silently produces garbage, every downstream step faithfully processes garbage. Convert it to a loop by adding an observation and an edge: after OCR, a verifier checks character-confidence and language-model perplexity; if the text looks like noise, control returns to the OCR step with a different engine or a deskew preprocessing pass. The loop can recover from a bad assumption at step one; the chain cannot. The single most common reason teams adopt loop engineering is discovering that a chain confidently produced wrong output because nothing in it could detect and reverse an early error.

Answer 1.3

Walk through the five-stage anatomy of a loop. Which stage carries the value?

The five stages are Intent (define the target outcome), Context (gather relevant code, docs, errors, logs, constraints), Action (edit files, run commands, call tools), Observation (capture test results, compiler errors, runtime output, diffs), and Adjustment (update the plan and repeat until accepted or blocked). The value is never in any single stage — it is in *closing* the loop. A test failure is not merely an error; it is new context the model did not have before. A type error is not a blocker; it is a signal that an assumption was wrong. Loop engineering’s job is to make those signals part of the system so the next action is conditioned on them, rather than discarding them and treating cleanup as a separate human task afterward.

Answer 1.4

Why did loop engineering mature first in coding agents rather than, say, customer support?

Code is uniquely self-verifying. A code change can be compiled, type-checked, linted, and run against a test suite, all of which produce unambiguous machine-readable observations with no human in the loop. That gives a coding agent a dense, cheap, trustworthy reward signal at every iteration — exactly what a loop needs to decide whether to continue, retry, or stop. Customer support, by contrast, has a sparse and subjective signal (was the customer satisfied?) that is expensive to measure and often only available long after the action. The lesson generalises: the feasibility of a reliable loop in any

domain is governed by the quality and latency of its verification signal. Before building a loop, the first design question is “what is my cheap, automatic verifier?” If the honest answer is “a human reads it,” the loop will be slow and expensive.

Answer 1.5

A teammate says “we already have a multi-step LangChain pipeline, so we do loop engineering.” How do you respond?

I would ask whether any step’s output can send control *backward* to an earlier step. Most LangChain pipelines are directed acyclic graphs — sophisticated chains. They have many LLM calls but no cycle, so a wrong result at an early node propagates unchecked. Having multiple steps is necessary but not sufficient for a loop; the defining property is a feedback edge gated by an observation. The practical test I use on a whiteboard: draw the system as a graph and look for a cycle whose back-edge is conditioned on a verification result. If there is none, it is a chain, and it will exhibit the classic failure mode of executing late steps on top of a broken early assumption.

Answer 1.6

What intellectual lineage does loop engineering draw on, and what is genuinely new?

It draws on the OODA loop (Observe–Orient–Decide–Act) from military decision theory and control systems; ReAct (Reason + Act), which showed that interleaving explicit reasoning with tool use beats either alone; and plan–act–verify from classical software engineering and robotics, where a plan is executed incrementally and checked against the real world. None of these ideas is new in isolation. What is new is that a single large language model can now occupy all three roles — planner, actor, and verifier — inside one system, often within one conversation, and can read a stack trace, hypothesise a cause, and write the fix itself. That collapses what used to require three separate subsystems (and often three teams) into one orchestrated loop, which is why the discipline crystallised only recently.

Answer 1.7

Debugging case: a loop produces confidently wrong output and never errors. Where do you look first?

Confident wrong output with no error is the signature of a missing or weak verifier — the loop has no observation that can contradict the model, so it treats its own plausible text as ground truth. I look first at the Observation and Adjustment stages: is there a verification step at all? Is its signal actually being fed back into the next prompt, or logged and ignored? A common bug is a verifier that runs but whose result is discarded because the loop’s termination condition keys off “model produced a final answer” rather than “verifier passed.” The fix is to make the success condition depend on the verifier, not on the model declaring victory. In the reference project this is why `AgentLoop.run` calls `_verify` before transitioning to `SUCCEEDED`, and feeds the failure reason back as a new user message on failure.

Answer 1.8

When is a chain the right choice and a loop the wrong one?

When the task is well-understood, low-uncertainty, and every step’s success is effectively guaranteed given valid input — a chain is simpler, cheaper, more predictable, and easier to trace. Formatting a known JSON payload into an email template, or running a fixed nightly report over clean data, gains nothing from a loop and pays for the extra model calls, latency, and non-determinism a loop introduces. Loops earn their cost only when the right path is not known in advance or when early steps can fail in ways that must be detected and reversed. A mature engineer reaches for the cheapest structure that meets the reliability bar: chain first, add loops exactly where uncertainty and recoverable failure actually live.

Answer 1.9

How would you measure whether moving from a chain to a loop actually improved the system?

I would hold the task fixed and compare on an offline eval set with three families of metric: *quality* (task success rate / pass@1 against a ground-truth checker), *cost* (tokens, dollars, and wall-clock per completed task), and *reliability* (variance of outcomes, rate of confidently-wrong outputs, and escalation rate). A loop almost always wins on quality and loses on raw cost-per-attempt, so the honest comparison is cost *per successful task*, which folds in the chain's silent failures and their downstream remediation cost. I would also track the distribution of iteration counts: if most tasks converge in one iteration, the loop is overkill; if many hit the budget ceiling, the verifier or the action space needs work. Without an eval set, "it feels better" is not a measurement.

Answer 1.10

Where does prompt engineering still matter once you have adopted loop engineering?

Everywhere inside the loop. A clearer system prompt and task framing reduce unnecessary exploration, lower the average iteration count, and make the model's tool-call arguments better-formed — all of which directly cut cost and latency. Prompt engineering also shapes the *quality of the model's self-reports*, which matters because those reports become observations the loop reasons over. The shift is not that prompts stop mattering; it is that they stop being the top-level control surface. The engineer now tunes the architecture (tool selection, context assembly, verification, termination) first, because that is where the leverage is, and tunes the prompt second, as one component inside that architecture rather than the whole system.

The Core Agent Loop

Subscribe → aiengineeringinsider.substack.com/subscribe

Everything is built on a loop that fits on an index card. Get it wrong and nothing above it holds.

Strip away the frameworks, the orchestration diagrams, and the vendor demos, and every agentic system reduces to the same dozen lines: a model decides, a tool runs, a result comes back, repeat. It fits on an index card. That is exactly why it is dangerous — a primitive this small invites you to bolt features on top before the primitive itself is sound, and a wobble at the center becomes a collapse at the edges.

At its core, an agent is a model calling tools in a loop until a task is complete. Everything else in this book — verification layers, multi-agent coordination, event-driven triggers, self-improving harnesses — is built on top of this one primitive. Get this loop wrong and nothing built above it will be reliable.

2.1 Observe, Reason, Act, Observe

The core cycle is a closed feedback loop: the agent observes the current state of its environment, reasons about what to do next, takes an action, and observes the new state that results. This repeats until a termination condition fires.




Figure 2.1: The core agent loop as a closed cycle. The agent reasons, acts, and observes the result, then loops — exiting only when the goal is met and verified. The feedback edge from observation back into reasoning is what makes it a loop rather than a chain.

2.2 The ReAct Pattern

Most modern agent loops trace back to the ReAct pattern (Reason + Act): interleave explicit reasoning steps with action steps, rather than asking the model to act in one shot. In code-generation contexts, this looks like:

1. Understand the goal.
2. Write some code.

3. Run the code and observe the output, or the error.
4. Reason about what went wrong.
5. Revise and re-run.
6. Repeat until the tests pass or the task is complete.

★ Key Insight

A test failure is not just an error message — it is new context the model did not have before. The loop's entire value comes from feeding that new context back into the next reasoning step, rather than discarding it.

2.3 Plan-and-Execute vs. Reactive Loops

There are two broad strategies for sequencing actions inside the loop.

Plan-and-execute

The agent first produces a multi-step plan, then executes it step by step, optionally re-verifying after each step. This works well when the order of operations matters and early mistakes compound — refactoring a module, setting up a new service, or building a feature with several interdependent components.

Reactive (improvisational)

The agent decides only its next single action based on the current state, with no explicit upfront plan. This works well for open-ended exploration — debugging an unfamiliar error, or operating in an environment too uncertain to plan against in advance.

! Pitfall / Warning

Plan-and-execute loops can over-commit to a bad plan. If step two reveals that the original plan was wrong, the agent needs to revise the plan, not push mechanically through the remaining steps. A plan-and-execute loop without a replanning mechanism is just a chain wearing a loop's clothing.

2.4 The Agent Loop as a State Machine

It is useful to think of the agent loop formally as a state machine: a set of states, a set of transitions between them, and one or more terminal states.

- **States** typically include: planning, gathering context, executing a tool call, awaiting tool result, reasoning over a result, awaiting human input, and terminated (success or failure).
- **Transitions** are triggered by events: a tool call returning, a verifier passing or failing, a budget being exhausted, a human approving or rejecting an action.
- **Terminal states** should be explicit and few: success, blocked (needs human input), and failed (exhausted retries or budget).

Treating the loop as a formal state machine, even informally on a whiteboard, surfaces gaps early: states with no exit transition, transitions with no clear trigger, and terminal states that are never actually reachable from certain branches.

2.5 Two Nested Loops

It is easy to conflate two genuinely different loops that operate at different timescales inside the same system:

1. **The token-generation loop** — the model’s own autoregressive decoding process, generating one token at a time. This loop is internal to the model and is not something the loop engineer designs.
2. **The step-orchestration loop** — the outer loop that decides which tool to call, what context to assemble, and when to stop. This is the loop this book is about.

From the Field

New practitioners sometimes try to “optimize the loop” by changing decoding parameters like temperature or top-p. Those knobs affect the inner token-generation loop. The outer orchestration loop — tool selection, context assembly, termination logic — is almost always where the actual leverage is. Tune the architecture before tuning the sampler.

2.6 Synchronous vs. Asynchronous Execution

A loop can run synchronously, where the engineer or calling system blocks and waits for each iteration to complete before the next begins, or asynchronously, where the loop runs in the background and the calling system is notified, polled, or interrupted when it completes or needs input.

Synchronous loops are simpler to reason about and debug. Asynchronous loops are necessary for any loop that runs longer than an interactive session tolerates, or that needs to run on a schedule independent of a human being present — the subject of Chapter 11.

✓ Best Practice

Start every new agentic system with a synchronous loop, even if the long-term plan is asynchronous, event-driven execution. Debugging synchronous failures is dramatically easier than debugging a loop that failed three hours ago in the background. Only move to asynchronous execution once the synchronous version is reliable.

Project Code — `loop_engineering_lab`

The reference implementation accompanying this chapter, [loop_engineering_lab/core/agent_loop.py](#), implements the five-stage anatomy from Chapter 1 as an explicit state machine with logged transitions, intended as a minimal, readable starting point rather than a production framework.

2.7 Plan-and-Execute vs. Reactive: A Decision Table

The choice between the two sequencing strategies is one of the most common system-design decisions in agent design. The table below frames it as an engineering trade-off rather than a preference.

Dimension	Plan-and-Execute	Reactive (Improvisational)
Decision granularity	Full multi-step plan up front	One next action at a time
Best for	Ordered, interdependent tasks (migrations, multi-file features)	Open-ended exploration (debugging, unfamiliar APIs)
Failure mode	Over-commits to a wrong plan	Wanders without global direction
Recovery mechanism	Requires explicit re-planning trigger	Naturally re-decides each step
Token cost	Higher up front, lower per step	Lower up front, can drift higher
Observability	Plan is an auditable artifact	Intent is implicit in the trace

Table 2.1: Neither strategy dominates; mature loops often start reactive and promote to plan-and-execute once the task structure is understood.

2.8 Reference Implementation: The Loop Driver

The entire core loop is about 120 lines. What matters is not its length but that every stop is a *named decision* and every transition is observable. The driver below (from [loop_engineering_lab/core/agent_loop.py](#)) shows the observe–reason–act cycle with an explicit control gate at the top of each iteration and verification before any success.

loop_engineering_lab/core/agent_loop.py — the run driver

```
def run(self, goal: str) -> LoopResult:
    state = LoopState(goal=goal)
    state.add_message("system", self.system_prompt)
    state.add_message("user", goal)

    while not state.phase.is_terminal:
        # 1) Control gate -- may we take another iteration at all? (Ch. 3)
        terminal = self.policy.check_terminal(state)
        if terminal is not None:
            phase, reason = terminal
            state.transition_to(phase, reason)
            break

        state.iteration += 1

        # 2) Reason -- ask the model for the next move.
        state.transition_to(LoopPhase.REASONING, "request next action")
        response = self.llm.complete(state.messages, tools=self.tools.schemas())
        state.tokens_used += response.total_tokens

        # 3a) Final-answer path: verify BEFORE declaring success.
        if not response.wants_tool:
            state.add_message("assistant", response.text)
            if self._verify(state):
                state.transition_to(LoopPhase.SUCCEEDED, "verified final answer")
                return LoopResult(LoopPhase.SUCCEEDED, response.text, state,
                                   ↪ self.tracer)
            self._record_no_progress(state, made_progress=False)
```

```

        continue # feed the failure back and loop again

# 3b) Action path: a stopping rule may veto the tool call (Ch. 3).
for call in response.tool_calls:
    veto = self.policy.veto_action(state, call.name, call.arguments)
    if veto is not None:
        state.transition_to(LoopPhase.BLOCKED, veto)
        return LoopResult(LoopPhase.BLOCKED, veto, state, self.tracer)
    state.transition_to(LoopPhase.ACTING, f"invoke {call.name}")
    result = self.tools.invoke(call.name, call.arguments)
    state.tool_calls_made += 1
    state.transition_to(LoopPhase.OBSERVING, f"observe {call.name}")
    state.add_message("tool", result.render(), tool_call_id=call.id)

```

✓ Best Practice

Notice the control gate is the *first* thing inside the loop body, before the model is even consulted. Checking budgets and termination before spending a model call — rather than after — is what prevents a loop from making one extra expensive call past its ceiling on every run.

2.9 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 2.1

Describe the core agent loop in one sentence, then explain what each part is doing.

Hint. A model calling tools in a loop until done — name reason, act, observe.

Exercise 2.2

Why model the loop as an explicit state machine rather than a while-loop with if-statements?

Hint. It makes hangs and unreachable terminal states structurally visible.

Exercise 2.3

Distinguish the token-generation loop from the step-orchestration loop. Why does it matter?

Hint. Inner decoding versus outer orchestration; leverage lives in the outer loop.

Exercise 2.4

When would you choose plan-and-execute over a reactive loop, and how do you prevent over-commitment?

Hint. Ordered, compounding tasks need a plan plus a replanning trigger.

Exercise 2.5

Why verify before transitioning to a success state? Show where this lives in code.

Hint. The model declaring “done” is a claim; verify against ground truth first.

Exercise 2.6

What are the terminal states of a well-designed loop, and why keep them few?

Hint. Few and explicit: succeeded, blocked, failed.

Exercise 2.7

Production scenario: an agent loop hangs and never returns. Walk through your diagnosis.

Hint. Check the transition log and look for a tool call with no timeout.

Exercise 2.8

What is the ReAct pattern and why does interleaving reasoning with action beat one-shot generation?

Hint. Interleaving conditions each step on a real observation.

Exercise 2.9

Synchronous vs. asynchronous loops: how do you decide, and what changes?

Hint. Build synchronously first; async adds durable state and idempotency.

Exercise 2.10

How would you make the core loop testable without spending money on model calls?

Hint. Depend on an `LLMClient` abstraction and inject a deterministic mock.

2.10 Chapter Summary

- The core agent loop is observe, reason, act, observe again — repeated until termination.
- The ReAct pattern interleaves explicit reasoning with action, rather than asking for one-shot output.
- Plan-and-execute loops suit ordered, interdependent tasks; reactive loops suit open-ended exploration. Both need a way to recover from a wrong initial assumption.
- Modeling the loop as a formal state machine surfaces missing transitions and unreachable terminal states before they cause production incidents.
- The token-generation loop (inside the model) and the step-orchestration loop (around the model) are different loops at different timescales; only the latter is what loop engineering designs.
- Build synchronously first; move to asynchronous execution only once the loop is reliable.

The next chapter addresses the question this one deliberately deferred: how does the loop know when to stop?

2.11 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 2.1

Describe the core agent loop in one sentence, then explain what each part is doing.

An agent is a model calling tools in a loop until a task is complete. Unpacked: the model *reasons* over the current message history to choose a next action; if that action is a tool call, the loop *acts* by dispatching it and *observes* the result by appending it to the history; the cycle repeats until a termination condition fires or the model emits a final answer that passes verification. Everything else in an agentic system — multi-agent coordination, event triggers, self-improvement — is built on this primitive. Getting it wrong (no verification, no budget, observations discarded) makes everything above it unreliable, which is why it is the first thing to get right.

Answer 2.2

Why model the loop as an explicit state machine rather than a while-loop with if-statements?

A state machine makes two whole classes of bug structurally visible: a phase with no exit transition (a hang), and a terminal phase that is unreachable from some branch (a loop that can never legitimately stop on that path). It also gives you a clean place to log every transition, which is the foundation of replay and debugging — you can reconstruct exactly why a run ended where it did. With ad-hoc if-statements, the “state” is implicit in local variables and control flow, so the same questions (can this ever terminate? what states are reachable?) require reading the whole function. In the reference project, `LoopState.transition_to` records every edge, so a finished run carries its own audit trail.

Answer 2.3

Distinguish the token-generation loop from the step-orchestration loop. Why does it matter?

The token-generation loop is the model’s own autoregressive decoding — generating one token at a time — and is internal to the model; the engineer does not design it. The step-orchestration loop is the outer loop that decides which tool to call, what context to assemble, and when to stop. They operate at completely different timescales and the engineer controls only the outer one. It matters because new practitioners try to “fix the loop” by tuning temperature or top-p, which only touches the inner decoding loop. The actual leverage — reliability, cost, correctness — lives almost entirely in the outer orchestration loop. The rule of thumb: tune the architecture before tuning the sampler.

Answer 2.4

When would you choose plan-and-execute over a reactive loop, and how do you prevent over-commitment?

Plan-and-execute suits tasks where the order of operations matters and early mistakes compound: a multi-file feature, a service migration, a refactor with interdependent steps. The danger is over-committing to a plan that step two reveals to be wrong. You prevent that with an explicit replanning trigger: after each step, the loop checks whether the observed reality still matches the plan’s assumptions, and if not, it regenerates the remaining plan rather than mechanically executing it. A plan-and-execute loop without a replanning mechanism is just a chain wearing a loop’s clothing — it will faithfully execute steps three through ten on top of a falsified assumption from step two.

Answer 2.5

Why verify before transitioning to a success state? Show where this lives in code.

Because the model declaring “done” is a claim, not a fact. If success keys off the model emitting a final answer, the loop will happily terminate on confident-but-wrong output — the most expensive failure mode because nothing downstream catches it. Verifying first makes success depend on an external source of truth (tests, a checker, a schema) and, critically, turns a verification *failure* into new context fed back into the next iteration rather than a silent rejection. In `AgentLoop.run`, the final-answer branch calls `self.verify(state)` and only transitions to `SUCCEEDED` if it returns true; on false it records no-progress and continues the loop with the failure reason appended as a message.

Answer 2.6

What are the terminal states of a well-designed loop, and why keep them few?

Three: `SUCCEEDED` (verified completion), `BLOCKED` (needs human input — missing credentials, a product judgment call, a vetoed destructive action), and `FAILED` (exhausted retries or budget). Keeping terminal states few and explicit forces every exit path to declare which of these it is, which makes downstream handling tractable: a caller knows exactly whether to retry, escalate to a human, or alert. Proliferating terminal states (“partially done,” “maybe done,” “done but unsure”) pushes the ambiguity onto every consumer of the loop’s result and is a frequent source of integration bugs in multi-agent systems where one loop’s output is another’s input.

Answer 2.7

Production scenario: an agent loop hangs and never returns. Walk through your diagnosis.

A hang means the loop never reached a terminal phase and never hit a budget ceiling — so either the control gate is not being consulted, or a budget is effectively infinite, or a phase has no exit edge. I would first check the transition log: where did the last transition land, and what phase is it stuck in? If it is oscillating between `REASONING` and `ACTING` without progress, the no-progress detector or tool-call budget is misconfigured. If it is sitting in `ACTING`, a tool call is blocking with no timeout — the most common real cause, a synchronous tool (HTTP, subprocess) without a deadline. The structural fixes: put the control gate first in the loop body (so budgets are checked before every model call), give every tool a hard timeout, and ensure the no-progress counter can drive the loop to `BLOCKED`. This is also the case for building synchronously first — a synchronous hang is trivial to reproduce; the same hang three hours into a background run is not.

Answer 2.8

What is the ReAct pattern and why does interleaving reasoning with action beat one-shot generation?

ReAct (Reason + Act) interleaves explicit reasoning steps with action steps instead of asking the model to produce a final answer in one shot. In a coding context: understand the goal, write some code, run it, observe the output or error, reason about what went wrong, revise, re-run, repeat until the tests pass. Interleaving wins because each action produces an observation the model did not have before, and conditioning the next reasoning step on real feedback is far more reliable than reasoning in a vacuum. One-shot generation has to get everything right with no chance to see the consequences of its choices; ReAct turns each consequence into input. The research finding that motivated the pattern is that reasoning-only and acting-only agents both underperform the interleaved version on the same tasks.

Answer 2.9

Synchronous vs. asynchronous loops: how do you decide, and what changes?

A synchronous loop blocks the caller until each iteration completes; an asynchronous loop runs in the background and notifies, polls, or interrupts the caller when it finishes or needs input. Synchronous loops are dramatically easier to reason about and debug, so the rule is to build every new system synchronously first and only move to asynchronous execution once the synchronous version is reliable.

Async becomes necessary when a loop runs longer than an interactive session tolerates, or must run on a schedule with no human present (Chapter 11). Moving to async forces new requirements the sync version did not need: durable, resumable state (so a crash does not restart from zero), idempotent actions (so a retry does not double a side effect), and out-of-band notification. Those are the real cost of going async — not the concurrency primitives, but the state and idempotency discipline.

Answer 2.10

How would you make the core loop testable without spending money on model calls?

Depend on an abstraction, not a provider. The loop should accept an `LLMClient` protocol with a single `complete` method; production injects a real adapter, tests inject a deterministic mock that returns scripted responses or runs a small policy over the message history. That seam lets you write fast, hermetic tests for the orchestration logic — “does the budget stop a runaway loop at three tool calls?”, “does a failing verifier keep the loop running until it passes?” — with zero network, zero cost, and full determinism. The reference project does exactly this: `MockLLM` accepts either a queue of responses or a `policy(messages)` function, so the entire test suite runs offline. The architectural principle is dependency inversion: the expensive, non-deterministic dependency sits behind an interface the test can replace.

Loop Control Structures

Subscribe → aiengineeringinsider.substack.com/subscribe

A loop that cannot stop is not autonomous. It is expensive.

The most expensive agent run is not the one that fails. It is the one that never stops — the loop that, at three in the morning, is on its four-hundredth attempt to fix a test that cannot pass, burning a dollar a minute against an API while no human is watching. It did not crash. It did everything you told it to. You simply never told it when to quit.

A loop without disciplined control logic is not an agent — it is a resource sink. This chapter covers the mechanisms that decide whether the next iteration happens, and what happens when it does not.

3.1 Termination Conditions

Every loop needs a concrete, evaluable definition of “done.” Three categories matter:

- **Success conditions** — tests pass, output matches a specification, a human approves.
- **Failure conditions** — maximum iterations reached, repeated errors with no progress, a required tool call fails irrecoverably.
- **Escalation paths** — hand off to a human or a different, more capable agent when the loop is stuck.

! Pitfall / Warning

Vague goals produce infinite loops or meaningless output. “Improve the dashboard” gives the loop nothing observable to optimize against. “Reduce initial dashboard load time by deferring non-critical charts while keeping existing filters working” contains scope, behavior, and a constraint — it produces a tight, checkable feedback loop.

3.2 Stopping Rules as First-Class Design Requirements

Stopping rules are not an afterthought bolted onto a working loop; they need to be designed before the loop’s action logic is written.

✓ Best Practice

Write your termination and failure conditions *before* writing any loop logic. “All tests pass and lint is clean” is a termination condition you can check programmatically. “The code looks good” is not a termination condition at all — it is a wish.

Beyond simple success/failure, mature loops include several specific categories of stopping rule:

- **Destructive-action guards** — stop before any irreversible command (database drops, force-pushes, credential rotation) unless a human has explicitly approved that specific action.
- **Scope-creep guards** — stop when the next step would touch files or systems outside the originally defined task boundary, even if the agent believes the change is beneficial.
- **Blocker escalation** — stop and hand off to a human when continuing requires missing credentials, missing data, or a product judgment call the agent has no basis to make.
- **Unrelated-change guards** — stop when the next step would affect changes a human or another agent made concurrently, rather than overwriting them silently.

3.3 Loop Budgets

A budget is a hard resource ceiling placed on the loop, independent of whether the loop believes it is making progress. Common budget dimensions:

- **Token budget** — maximum tokens consumed across all iterations.
- **Time budget** — maximum wall-clock duration.
- **Tool-call budget** — maximum number of tool invocations.
- **Cost budget** — maximum dollar spend, particularly relevant when iterations route across models of different price points.

★ Key Insight

An exhausted budget without progress is itself a signal, not just a limit. If a loop burns through its tool-call budget without converging, that is evidence the current strategy is wrong — treat it as a trigger to escalate or switch strategy, not merely as a stop.

3.4 Early-Exit and Short-Circuit Logic

Not every loop needs to run to its budget ceiling. A loop should be able to exit early when:

- the success condition is met well before the budget is exhausted;
- an early observation makes the remaining plan provably unnecessary (for example, discovering the bug does not exist in the version being tested);
- a hard blocker is detected immediately, with no value in continuing to attempt workarounds.

3.5 Recursive vs. Iterative Structures

A loop can repeat by looping in place (iteration) or by spawning a new sub-agent to handle a sub-problem and waiting for its result (recursion). Recursive structures are useful when a sub-task is genuinely independent and benefits from its own fresh context window, its own tool budget, and its own termination logic — rather than accumulating inside the parent's growing context.

3.6 Nested Loops

Most production agentic systems are not a single loop but a loop containing loops: an outer planning loop that decides *what* to do next, containing an inner tool-execution loop that decides *how* to do it.




Figure 3.1: Loop nesting: an outer planning loop decides *what* to do next; an inner tool-execution loop decides *how*. This stacking recurs at every scale in the book.

This structure recurs throughout the rest of this book: Chapter 8’s verification loop wraps the core agent loop; Chapter 12’s hill-climbing loop wraps an entire population of agent runs. Nesting is the load-bearing idea behind “loop stacking” as a design principle, introduced fully in Chapter 12.

3.7 Loop Re-Entry and Resumability

Long-running loops — particularly those that run asynchronously or on a schedule (Chapter 11) — must be able to pause and later resume from exactly where they left off, rather than restarting from scratch. This requires:

- serializing the loop’s current state (what has been tried, what succeeded, what is pending) to durable storage;
- a clear resumption protocol that reconstructs the loop’s working memory from that serialized state;
- idempotent actions, so that resuming after an interruption does not repeat a side effect that already completed (covered further in Chapter 6).

From the Field

The single most common production failure I have seen in long-running loops is not a bad model decision — it is a loop that crashes mid-iteration with no durable record of what it had already done, and on restart, redoes a half-finished action a second time. Resumability is not a nice-to-have for loops that run longer than a few minutes; it is a prerequisite.

3.8 Budget Dimensions at a Glance

Budget	Protects against	Typical signal on exhaustion
Iteration	Pathological retry loops	Strategy is wrong; escalate
Token	Context bloat, runaway cost	Compaction or summarisation needed (Ch. 4)
Tool-call	Thrashing on a stuck sub-task	Action space or verifier is inadequate
Time (wall-clock)	Blocking tool hangs, SLAs	A dependency is slow or down
Cost (USD)	Cross-model routing surprises	Route to a cheaper model or stop

Table 3.1: Each budget is a hard ceiling *and* a diagnostic: which budget you hit tells you which subsystem to fix.

3.9 Reference Implementation: Budgets and Stopping Rules

Control logic should be a separate, testable object — not `if`-statements buried in the loop body. The reference project models the budget, the success predicate, and the stopping rules as one `TerminationPolicy` the loop consults; you can unit-test “does this budget stop the loop at five tool calls?” without ever invoking a model.

`loop_engineering_lab/core/control.py` — policy as a first-class object

```
@dataclass
class Budget:
    """A hard resource ceiling, independent of perceived progress."""
    max_iterations: int = 12
    max_tokens: int = 100_000
    max_tool_calls: int = 40
    max_seconds: float = 300.0

    def exceeded(self, state: LoopState) -> str | None:
        if state.iteration >= self.max_iterations:
            return f"iteration budget exhausted ({self.max_iterations})"
        if state.tokens_used >= self.max_tokens:
            return f"token budget exhausted ({self.max_tokens})"
        if state.tool_calls_made >= self.max_tool_calls:
            return f"tool-call budget exhausted ({self.max_tool_calls})"
        if state.elapsed_seconds >= self.max_seconds:
            return f"time budget exhausted ({self.max_seconds}s)"
        return None

    def destructive_action_guard(dangerous=None):
        """Block irreversible tools unless a human approved them in scratchpad."""
        dangerous = dangerous or {"drop_table", "force_push", "rotate_credentials"}
        def rule(state, tool_name, args):
            approved = state.scratchpad.get("human_approved", set())
            if tool_name in dangerous and tool_name not in approved:
                return f"destructive action '{tool_name}' requires human approval"
            return None
```

```
return rule
```

The success predicate is injected, not hard-coded, so the same loop machinery serves a coding task (“tests pass”), a research task (“three independent sources agree”), and a data task (“schema validates”). Decoupling the *policy* from the *mechanism* is what lets one 120-line loop be reused across domains.

! Pitfall / Warning

A stopping rule that only *logs* a violation instead of vetoing the action is theatre. In the reference loop, `policy.veto_action` returns a reason *before* the tool is dispatched, and a non-null reason transitions the loop straight to BLOCKED. A guard that fires after the irreversible command has already run is not a guard.

3.10 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 3.1

Why must termination conditions be designed before the loop’s action logic?

Hint. “Done” must be observable before you write any action logic.

Exercise 3.2

Name the categories of stopping rule a mature loop should have beyond simple success/failure.

Hint. Destructive-action, scope-creep, blocker-escalation, and unrelated-change guards.

Exercise 3.3

A loop exhausts its tool-call budget without converging. Is that just a stop?

Hint. It is a signal the strategy is wrong — escalate, do not just stop.

Exercise 3.4

How do you set budget values? Walk through your reasoning for a coding agent.

Hint. Anchor on the eval distribution, above the 95th percentile of successes.

Exercise 3.5

Explain the no-progress detector. Why is it different from a budget?

Hint. It caps wasted (repeated) effort, not absolute consumption.

Exercise 3.6

When should a loop recurse (spawn a sub-agent) versus iterate in place?

Hint. Recurse for genuinely independent sub-tasks that need fresh context.

Exercise 3.7

Describe the nested-loop structure of a typical production agent and why it recurs.

Hint. An outer planning loop wrapping an inner tool-execution loop at different timescales.

Exercise 3.8

Production case: a long-running loop crashes mid-iteration and on restart repeats a side effect. What went wrong and how do you fix it?

Hint. Missing durable state plus non-idempotent actions — serialize and key by id.

Exercise 3.9

How do you implement a destructive-action guard so it cannot be bypassed?

Hint. Fail-closed, per-action, out-of-band human approval, checked before dispatch.

Exercise 3.10

What does early-exit / short-circuit logic add, and when is it wrong to use?

Hint. Exit on verified signals only, never on the model's self-assessment.

3.11 Chapter Summary

- Every loop needs explicit success conditions, failure conditions, and escalation paths — defined before the loop's action logic is written.
- Specific stopping-rule categories worth designing deliberately: destructive-action guards, scope-creep guards, blocker escalation, and unrelated-change guards.
- Budgets (token, time, tool-call, cost) are hard ceilings; exhausting a budget without progress is a signal to change strategy, not just a stop.
- Loops nest: an outer planning loop typically contains an inner tool-execution loop, a structure that recurs at every scale in this book.
- Long-running loops need durable, resumable state so an interruption does not force a restart from zero or a repeated side effect.

With control structures established, Part II turns to what happens *inside* each iteration: how state is carried forward, how the agent reasons, and how it uses tools.

3.12 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 3.1

Why must termination conditions be designed before the loop's action logic?

Because a loop without a concrete, evaluable definition of “done” is a resource sink — it either runs forever or stops on a meaningless heuristic. Writing the termination conditions first forces you to answer “what observable signal means success?” before you write a single line of action logic, and that question often reshapes the whole design. “All tests pass and lint is clean” is a programmatically checkable termination condition; “the code looks good” is a wish, not a condition. Designing stopping rules last, after the loop already runs, almost always produces a system that can act but cannot reliably decide to stop — which in production manifests as runaway cost and confidently-wrong terminal output.

Answer 3.2

Name the categories of stopping rule a mature loop should have beyond simple success/failure.

Four matter in production. *Destructive-action guards* stop before any irreversible command (database drops, force-pushes, credential rotation) unless a human has explicitly approved that specific action. *Scope-creep guards* stop when the next step would touch files or systems outside the originally defined task boundary, even if the agent believes the change is beneficial. *Blocker escalation* stops and hands off when continuing requires missing credentials, missing data, or a product judgment call the agent has no basis to make. *Unrelated-change guards* stop when the next step would overwrite changes a human or another agent made concurrently. Each is a specific, codifiable veto on the next action, and each maps to a real incident class — they are not hypothetical.

Answer 3.3

A loop exhausts its tool-call budget without converging. Is that just a stop?

No — it is a signal, and treating it as merely a limit wastes the most valuable diagnostic the loop produces. If a loop burns its entire tool-call budget without reaching success, that is strong evidence the current strategy is wrong: the action space is inadequate, the verifier is giving an unusable signal, or the task is mis-scoped. The right response is to escalate or switch strategy, not to silently raise the ceiling and try again. Which budget was exhausted is itself informative: hitting the token budget points to context bloat; hitting the time budget points to a slow or hung dependency; hitting the tool-call budget points to thrashing on a stuck sub-task. The budget is a hard ceiling and a diagnostic at the same time.

Answer 3.4

How do you set budget values? Walk through your reasoning for a coding agent.

I anchor on the offline eval distribution, not a guess. I run the loop over a representative task set and look at the iteration-count and token histograms for *successful* runs: the budget should sit above the 95th percentile of successful completions so it does not truncate legitimately hard tasks, but below the point where a non-converging loop wastes serious money. If most successes finish in two iterations but the ceiling is twenty, the ceiling is doing nothing useful and a stuck loop will burn eighteen wasted iterations before stopping — so I would add a no-progress detector that fires far earlier. Budgets are also defence-in-depth: the token and cost budgets protect against pathologies the iteration budget misses, like a single iteration that balloons the context. I tune them per task class, not globally.

Answer 3.5

Explain the no-progress detector. Why is it different from a budget?

A budget caps absolute resource consumption; a no-progress detector caps *wasted* consumption by noticing the loop is repeating itself. The classic implementation counts consecutive iterations that produced no new information — the same error, the same failing test, a near-identical action — and forces the loop to **BLOCKED** once that count crosses a threshold. It fires earlier than a budget on the specific pathology that matters most: a loop confidently retrying a strategy that cannot work. In the reference project, **TerminationPolicy** tracks **consecutive_no_progress** in scratchpad and drives the loop to **BLOCKED** at **no_progress_limit**. The subtlety is defining “progress” — a good signal is a change in the verifier’s output or the set of failing tests, not merely the model saying something different.

Answer 3.6

When should a loop recurse (spawn a sub-agent) versus iterate in place?

Recurse when a sub-task is genuinely independent and benefits from its own fresh context window, its own tool budget, and its own termination logic — rather than accumulating inside the parent’s growing context. A good example: the parent is implementing a feature and needs a focused investigation of why one library function behaves unexpectedly; spawning a sub-agent with just that question keeps the parent’s context clean and gives the sub-task a bounded budget. Iterate in place when the steps are tightly coupled and share state that would be expensive to serialise and pass down. The cost of recursion is coordination overhead and context-passing; the benefit is isolation and independent budgets. The anti-pattern is recursing for trivially coupled steps, which adds latency and failure surface for no isolation benefit.

Answer 3.7

Describe the nested-loop structure of a typical production agent and why it recurs.

Most production systems are a loop containing loops: an outer planning loop that decides *what* to do next, wrapping an inner tool-execution loop that decides *how* to do it. The verification loop (Chapter 8) wraps the core agent loop; the hill-climbing loop (Chapter 12) wraps an entire population of agent runs. Nesting recurs because each layer has a different timescale and a different termination signal: the inner loop terminates on “this tool call succeeded,” the outer on “the overall goal is verified.” Conflating them — using one budget and one stopping rule for both — is a common design error that makes the system either too eager to stop (the outer goal abandoned because one inner step failed) or unable to stop (inner thrashing never bubbling up to the outer budget). Loop stacking is the load-bearing idea behind composing complex behaviour from simple loops.

Answer 3.8

Production case: a long-running loop crashes mid-iteration and on restart repeats a side effect. What went wrong and how do you fix it?

Two things were missing: durable state and idempotent actions. The crash lost the loop’s working memory — what had been tried, what succeeded, what was pending — so on restart it redid a half-finished action whose side effect had already committed (a duplicate payment, a doubled file write, a re-sent notification). The fix has three parts. First, serialise the loop’s state to durable storage at each iteration boundary, not just at the end, so a restart can reconstruct exactly where it was. Second, make side-effecting actions idempotent — key them by a deterministic request id so a repeat is a no-op. Third, define a clear resumption protocol that rebuilds working memory from the serialised state rather than starting over. Resumability is not optional for any loop that runs longer than a few minutes; it is a prerequisite, and the most common real production failure I have seen is its absence.

Answer 3.9

How do you implement a destructive-action guard so it cannot be bypassed?

The guard must run *before* the action is dispatched and must veto by default, requiring an explicit, specific human approval recorded out of band — not a flag the model itself can set. In the reference project, `destructive_action_guard` checks whether the tool name is in a dangerous set and whether that specific tool is present in `state.scratchpad["human_approved"]`, which only a human approval step can populate; a non-null return blocks the call and transitions the loop to `BLOCKED`. Three properties make it un-bypassable: it is consulted on the action path for every tool call (not just some), it fails closed (absence of approval means block), and approval is per-action, so approving one `drop_table` does not approve the next. The failure mode to avoid is a guard that logs a warning and proceeds, or one the model can satisfy by emitting “I have approval” in its own text.

Answer 3.10

What does early-exit / short-circuit logic add, and when is it wrong to use?

Early exit lets a loop stop before its budget ceiling when continuing is provably pointless: the success condition is already met, an early observation makes the remaining plan unnecessary (the bug does not exist in the version under test), or a hard blocker is detected immediately with no value in attempting workarounds. It saves cost and latency on the common case. It is wrong — or rather, dangerous — when the “proof” that continuing is pointless is itself the model’s unverified claim. Short-circuiting on “the model thinks it is done” reintroduces the confident-wrong-output failure mode that verification exists to prevent. The discipline is: early-exit on *verified* signals (a checker passed, a deterministic precondition failed), never on the model’s self-assessment alone. Early exit and verification are complementary, not substitutes.

PART II

State, Reasoning, and Tools Inside the Loop

With the outer loop in place, Part II turns inward to the three things every iteration manages: the *state* it carries forward, the *reasoning* it does between actions, and the *tools* it uses to touch the world. These are the levers that decide whether a loop converges quickly or thrashes — and the tool set, in particular, is the ceiling on how good any loop can be.

State and Memory Across Iterations

Subscribe → aiengineeringinsider.substack.com/subscribe

A model has no memory. The loop is where memory has to live.

A model has no memory. This is the fact every loop engineer eventually runs into, usually the hard way: the agent that fixed a bug brilliantly an hour ago has no idea it ever happened, and the agent on iteration twenty is drowning in the transcript of iterations one through nineteen. Between total amnesia and total recall lies the entire craft of this chapter — deciding what a loop carries forward, what it compresses, and what it must write down somewhere the model can never forget it.

Every iteration of a loop generates new context: code written, errors encountered, decisions made. Left unmanaged, this accumulates until it either exceeds the context window or drowns the model in irrelevant history. This chapter covers how to carry forward exactly the state that matters.

4.1 The Context Window as Loop Memory

By default, an agent's only memory of what it has already tried is whatever remains inside its context window. This is a fragile foundation for anything beyond a short loop: it is bounded in size, costs money and latency proportional to its length, and degrades the model's attention as it grows.

Working Memory

The subset of accumulated loop history that is actively useful for deciding the next action — as distinct from the full transcript of everything that has happened, most of which is no longer relevant by iteration ten.

4.2 Scratchpad Patterns

A scratchpad is an explicit, structured piece of intermediate state, kept separate from the raw conversation history, that the agent reads and updates each iteration. Rather than relying on the model to locate relevant facts buried in a long transcript, the scratchpad surfaces exactly what is needed: current plan, what has been tried, what is still open.

4.3 State Compaction

✓ Best Practice

Log everything, summarize often. Keep a running log of every action and its outcome. Before each new iteration, summarize that log into a compact working memory rather than re-feeding the full transcript. A summary like “attempted fix A (failed: `TypeError`), attempted fix B (failed: same error), attempted fix C (partially successful: error resolved but tests still failing on line 47)” is far more useful to the model than the full transcript of three failed attempts.

4.4 Context Window Overflow Handling

Three standard strategies handle the inevitable growth of loop history:

- **Truncation** — drop the oldest content first. Simple, but risks discarding an early decision that turns out to matter later.
- **Rolling summaries** — periodically compress older history into a summary, keeping recent history verbatim. The most common approach in production loops.
- **Sliding windows** — keep only the last N iterations in full, with no summary of what came before. Appropriate for loops where only recent state is ever relevant.

4.5 Persistent vs. Ephemeral State

Ephemeral state

Exists only for the duration of a single loop run — the current plan, the current scratchpad, the current set of tool results. Discarded when the loop terminates.

Persistent state

Survives across separate loop runs — project conventions, prior decisions, accumulated learnings about what tends to fail in this codebase. This is the substance of Chapter 16’s discussion of repository-level instructions.

4.6 State Serialization for Resumable Loops

As established in Chapter 3, any loop expected to pause and resume needs its state serialized to durable storage in a form that can fully reconstruct the loop’s working memory on resumption — not just the final answer, but the scratchpad, the attempt history, and the current plan.

4.7 Tracking What’s Been Tried

! Pitfall / Warning

A loop that retries the exact same action after the same error is not learning — it is spinning. The single most effective defense against this failure mode is an explicit, structured record of prior attempts and their outcomes, checked before each new action, rather than relying on the model to remember it spotted this error three iterations ago.

4.8 External Memory as the Real Persistence Layer

The most reliable form of loop memory is not a context-engineering trick at all — it is a plain external artifact that lives outside the conversation entirely: a markdown file, a ticket-tracker board, a structured log committed to the repository.

★ Key Insight

The model forgets everything between separate runs; the memory has to be on disk, not in the context. This is the same mechanism every long-running, multi-session agentic system depends on. A markdown progress file or a ticket board sounds almost too simple to matter, but it is the single most load-bearing piece of infrastructure in any loop that needs to survive longer than one sitting.

4.9 Skills as Externalized Project Intent

Intent Debt

The recurring cost paid every time an agent starts a session with no memory of project-specific context, and fills that gap with a confident but possibly wrong guess — conventions, build steps, the reason a particular pattern was avoided after a past incident.

A skill — a small, well-scoped file containing project-specific instructions, conventions, and worked examples — is intent written down once, on the outside, so the agent reads it fresh every run instead of re-deriving the project from zero.

From the Field

Without externalized skills, a loop re-derives your entire project's conventions from scratch every cycle, and that cost compounds with every run. With skills in place, the opposite happens: each run starts from accumulated project knowledge instead of a blank slate, and the loop's output quality compounds upward instead of staying flat.

Project Code — `loop_engineering_lab`

`loop_engineering_lab/state/scratchpad.py` implements a minimal scratchpad with attempt-tracking and rolling-summary compaction, used by the reference agent loop from Chapter 2.

4.10 Overflow Strategies Compared

Strategy	Keeps / drops	Best when / risk
Truncation	Drops oldest verbatim	Cheapest; risks losing an early decision that matters later
Rolling summary	Compresses old, keeps recent verbatim	Default for production loops; risk is a lossy or wrong summary
Sliding window	Keeps last N only, no summary	Only recent state matters; risk is amnesia about earlier constraints
External artifact	Moves memory out of context entirely	Multi-session work; risk is staleness between context and disk

Table 4.1: Memory management is a cost/fidelity trade-off; most real loops combine a rolling summary with an external artifact of record.

4.11 Reference Implementation: A Compactable Scratchpad

The scratchpad is the structured working memory that sits beside the raw transcript. The implementation below (`loop_engineering_lab/state/scratchpad.py`) records every attempt, detects when an action has already been tried, and folds older attempts into a rolling summary so the context stays small without losing the substance of what failed.

loop_engineering_lab/state/scratchpad.py — attempt tracking + compaction

```
def has_tried(self, action: str, detail: str = "") -> bool:
    """True if this exact action+detail was already attempted -- the single
    most effective guard against a loop spinning on a repeated failure."""
    sig = Attempt(action, "", detail).signature()
    return any(a.signature() == sig for a in self.attempts)

def compact(self) -> None:
    """Fold all but the most recent `keep_verbatim` attempts into the rolling
    summary. 'fix A (failed: TypeError); fix B (failed: same error)' is far
    more useful to the model than the full transcript of three failures."""
    if len(self.attempts) <= self.keep_verbatim:
        return
    old, self.attempts = self.attempts[:-self.keep_verbatim],
    ↪ self.attempts[-self.keep_verbatim:]
    folded = "; ".join(
        f"{a.action} ({a.outcome}): ' + a.detail if a.detail else ''}" for a in
    old
    )
    self.summary = (self.summary + " " + folded).strip() if self.summary else folded
```

The `has_tried` check is the load-bearing line: before the loop issues an action, it asks whether this exact action with this exact failure signature has already been attempted, and if so refuses to spin on it. That single guard eliminates the most common “stuck loop” pathology more reliably than any prompt instruction to “remember what you tried.”

4.12 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 4.1

Why is the context window a fragile foundation for loop memory?

Hint. Bounded, costly, and the model's attention degrades as it grows.

Exercise 4.2

Distinguish working memory from the full transcript. Why separate them?

Hint. Distil the actively-useful subset out of the noisy transcript.

Exercise 4.3

Compare truncation, rolling summaries, and sliding windows. When does each fail?

Hint. Each trades cost against fidelity; the rolling summary is the default.

Exercise 4.4

A loop keeps retrying the same failing fix. Diagnose and fix it.

Hint. This is spinning — add a checked attempt log consulted before acting.

Exercise 4.5

What state must be serialized for a loop to be resumable, and in what form?

Hint. The full working memory (plan, attempts, scratchpad), not the final answer.

Exercise 4.6

Persistent vs. ephemeral state: give examples and explain the boundary.

Hint. Facts about this attempt versus facts about the world.

Exercise 4.7

Why is an external artifact (a file, a board) often the most reliable memory layer?

Hint. The model forgets between runs, so memory must live on disk.

Exercise 4.8

What is “intent debt” and how do skills pay it down?

Hint. Re-deriving project context each run; skills write that intent down once.

Exercise 4.9

How would you decide *when* to trigger compaction in a long loop?

Hint. Trigger on a measurable threshold while preserving failure signatures.

Exercise 4.10

Production case: a multi-day migration loop resumes after a crash and corrupts a file it had already migrated. What memory failures are at play?

Hint. A missing completion ledger plus a non-idempotent migration step.

4.13 Chapter Summary

- Relying on the raw context window as the only memory is fragile; it is bounded, costly, and degrades model attention as it grows.
- Scratchpads separate actively useful working memory from the full, noisy transcript.
- Compact loop history through rolling summaries that preserve the substance of failed attempts without the full transcript.
- Persistent state (project conventions, accumulated learnings) should outlive any single loop run; ephemeral state should not.
- The most reliable memory layer is an external artifact — a file, a board, a log — that exists outside the model’s context entirely.
- Skills externalize project-specific intent so the agent does not pay “intent debt” by re-deriving context from zero every run.

4.14 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 4.1

Why is the context window a fragile foundation for loop memory?

Because it is bounded, costly, and degrades the model's attention as it grows. Every token of history costs money and latency on every subsequent call, so a loop that simply appends raw history pays linearly more per iteration while getting *worse* at attending to the relevant parts — the “lost in the middle” effect, where information buried in a long context is attended to less reliably than information at the ends. Past a short loop, relying on the window as the only memory means the loop either hits a hard limit or quietly starts ignoring early-but-important context. The fix is to treat the window as a cache of recent, relevant state and push durable memory to an external store.

Answer 4.2

Distinguish working memory from the full transcript. Why separate them?

The full transcript is everything that has happened; working memory is the subset actively useful for deciding the *next* action. By iteration ten, most of the transcript is noise — superseded plans, resolved errors, verbose tool dumps — and forcing the model to re-locate the relevant facts inside it wastes tokens and degrades decisions. A scratchpad holds the distilled working memory explicitly: current plan, what has been tried and with what outcome, what is still open. Separating them lets you feed the model a small, high-signal state every turn while keeping the full transcript on disk for audit and replay. The transcript is for humans and post-hoc analysis; the scratchpad is for the next decision.

Answer 4.3

Compare truncation, rolling summaries, and sliding windows. When does each fail?

Truncation drops the oldest content first — cheapest, but it can discard an early decision or constraint that turns out to matter, causing the loop to violate something it “knew” fifteen iterations ago. A sliding window keeps only the last N iterations verbatim with no summary of before; fine when only recent state is ever relevant, but it produces amnesia about earlier constraints. Rolling summaries compress older history into a running summary while keeping recent history verbatim — the production default — but its failure mode is a lossy or subtly wrong summary that confidently misstates what happened, which is worse than truncation because the model trusts it. The practical answer is usually a rolling summary plus an external artifact of record so the summary is recoverable if it drifts.

Answer 4.4

A loop keeps retrying the same failing fix. Diagnose and fix it.

This is “spinning”: the loop has no durable, checked record that it already tried this action and got this error, so each iteration re-derives the same plausible-looking fix. The diagnosis is a missing attempt log consulted *before* acting — not a model-capability problem. The fix is an explicit, structured record of prior attempts keyed by a stable signature (the action plus an error fingerprint), checked before each new action; if the signature matches a past failure, the loop must change strategy or escalate rather than repeat. In the reference project, `Scratchpad.has_tried` provides exactly this guard, and `structured_feedback` flags a repeated error fingerprint as “REPEAT — change approach” in the very feedback the model reads. Relying on the model to “remember” it saw this error three turns ago is unreliable; making repetition structurally visible is not.

Answer 4.5

What state must be serialized for a loop to be resumable, and in what form?

Not the final answer — the full working memory needed to reconstruct the loop mid-flight: the current plan, the attempt history with outcomes, the open items, the scratchpad summary, and the loop’s phase and budget counters. It must be serialized to durable storage at iteration boundaries (not only at the end) in a form that a resumption protocol can fully rehydrate, so a restart continues from where it stopped rather than from zero. The reference project’s `LoopState` is deliberately a plain dataclass for this reason: phase, iteration, messages, transitions, token and tool counters, and scratchpad are all serialisable. The test of “did I serialize enough?” is whether a fresh process, given only the serialized blob, can make the same next decision the original would have.

Answer 4.6

Persistent vs. ephemeral state: give examples and explain the boundary.

Ephemeral state exists only for one run — the current plan, the current scratchpad, this run’s tool results — and is discarded at termination. Persistent state survives across runs — project conventions, prior architectural decisions, accumulated knowledge of what tends to fail in this codebase. The boundary is whether the information is a fact about *this attempt* or a fact about *the world the loop operates in*. Getting the boundary wrong is costly in both directions: persisting ephemeral state pollutes future runs with stale plans, while treating persistent state as ephemeral makes the agent re-derive the project from scratch every run (“intent debt”). The discipline is to write durable, world-level facts to an external store and let run-level facts die with the run.

Answer 4.7

Why is an external artifact (a file, a board) often the most reliable memory layer?

Because the model forgets everything between separate runs, so the memory has to live on disk, not in the context. An external markdown progress file, a ticket board, or a structured log committed to the repository persists across sessions, is human-inspectable, survives crashes, and is the same artifact a human teammate would use, which makes the hand-off between agent and human a non-event. It sounds almost too simple to be “architecture,” but it is the single most load-bearing piece of infrastructure in any loop that must survive longer than one sitting. The context window is a cache; the external artifact is the system of record. The only real risk is drift between the two, managed by treating the artifact as authoritative and rehydrating context from it.

Answer 4.8

What is “intent debt” and how do skills pay it down?

Intent debt is the recurring cost paid every time an agent starts a session with no memory of project-specific context and fills the gap with a confident but possibly wrong guess — the build command, a naming convention, the reason a pattern was banned after a past incident. Each run re-derives the project from zero and sometimes derives it wrong. A skill — a small, well-scoped file of project-specific instructions, conventions, and worked examples — is that intent written down once, on the outside, so the agent reads it fresh every run instead of guessing. The effect compounds: with skills, each run starts from accumulated knowledge and output quality trends upward; without them, quality stays flat and every run re-pays the same debt. Skills are externalised, persistent, world-level memory in the sense of the previous question.

Answer 4.9

How would you decide when to trigger compaction in a long loop?

I trigger on a measurable threshold, not a vibe: when the assembled context crosses a fraction of the model’s window (say 60–70%), or when the attempt log exceeds a fixed count, compact

the oldest entries into the rolling summary while keeping the most recent N verbatim. Threshold-based compaction keeps per-iteration cost roughly flat instead of growing linearly, which matters because uncompact growth raises both price and latency on *every* remaining call. I also compact opportunistically at natural boundaries — after a sub-task completes, its detailed history can collapse to a one-line outcome. The thing to avoid is compacting so aggressively that the summary loses the failure signatures the loop needs to avoid repeating itself; compaction must preserve “what was tried and why it failed,” which is exactly the substance the model needs.

Answer 4.10

Production case: a multi-day migration loop resumes after a crash and corrupts a file it had already migrated. What memory failures are at play?

Two, compounding. First, a serialization failure: the loop did not durably record, at an iteration boundary, that this specific file was already migrated, so on resume it re-selected it as pending. Second, an idempotency failure (Chapter 6): the migration action was not safe to apply twice, so re-running it on already-migrated content corrupted it. The memory fix is to persist a per-unit completion ledger (file X: migrated, verified, committed) as part of resumable state, and to consult it before acting — the same “what have I already done” guard as the spinning case, but at the persistence layer. The action fix is to make the migration idempotent or to gate it on the ledger. Together they turn “resume” from a hazard into a no-op for completed work. The root lesson: resumable loops need a durable record of *completed side effects*, not just intended plans.

Reasoning Within the Loop

Subscribe → aiengineeringinsider.substack.com/subscribe

Same model, same error, opposite outcomes — the difference is the thinking in between.

Two agents face the identical failing test. The first reads the error, concludes “probably an import problem,” and starts changing imports. The second reads the same error, asks what would have to be true for it to appear, forms two competing hypotheses, and runs the cheapest experiment that tells them apart. Same model, same tools, same observation — wildly different outcomes. The difference is entirely in the thinking that happens in the half-second between seeing the result and choosing the next move.

Acting and observing are necessary but not sufficient. What the agent does *between* an observation and its next action — how it reasons about what just happened — determines whether the loop converges efficiently or wanders.

5.1 Intra-Step vs. Loop-Level Reasoning

It is useful to separate two different scopes of reasoning that are often conflated:

- **Intra-step reasoning** — chain-of-thought *within* a single model call, reasoning toward one action or one response.
- **Loop-level reasoning** — reasoning that spans multiple iterations, incorporating the outcomes of prior actions to decide whether the overall strategy is still sound.

A model can reason beautifully within a single step and still pursue a bad overall strategy across the loop, because intra-step reasoning has no visibility into whether earlier steps actually worked. Loop-level reasoning is what catches that.

5.2 Reflection and Self-Critique

A reflection loop adds an explicit generate-then-critique cycle: the agent produces an output, then is prompted to critique that output against stated criteria before either accepting it or revising it.

★ Key Insight

Reflection is most effective when the critique step uses a different prompt, and ideally different context framing, than the generation step. Asking the same call that just produced an answer “are you sure?” tends to produce confirmation rather than genuine critique. Framing the critique as “find the flaw” rather than “check your work” produces measurably more useful self-correction.

5.3 Plan Revision Mid-Loop

A plan generated at the start of a loop is a hypothesis, not a commitment. The loop needs an explicit mechanism for detecting when an observation invalidates the current plan, and for replanning rather than mechanically continuing.

Signals that typically indicate a plan needs revision:

- a step fails in a way the plan did not anticipate;
- an observation reveals a constraint the plan was not built around;
- the cost of completing the remaining plan now clearly exceeds its expected value.

! Pitfall / Warning

A plan-and-execute loop with no replanning trigger will mechanically execute steps three through eight on top of a step two that already failed silently, producing output that looks complete but is built on a broken foundation.

5.4 Tree-of-Thought and Graph-of-Thought

Rather than committing to a single reasoning path, tree-of-thought and graph-of-thought approaches let the loop explore multiple candidate reasoning branches, evaluate them, and prune the weaker ones — effectively turning a single linear loop into a small search over possible loops.




Figure 5.1: Tree-of-thought: the loop branches into candidate reasoning paths, evaluates each, and prunes the weaker — a small search over possible loops.

This is expensive — each branch consumes its own tokens and tool calls — so it is best reserved for points in the loop where the right approach is genuinely unclear, not applied uniformly across every decision.

5.5 Verifier-in-the-Loop

A verifier is a separate reasoning step — sometimes a separate model call entirely — whose only job is to check a candidate output or action against explicit criteria before the loop is allowed to proceed. This is distinct from, and a precursor to, the dedicated verification *loop* architecture covered fully in Chapter 8; here it refers to verification as a single gating step inside an otherwise ordinary agent loop.

5.6 Self-Consistency Sampling

Rather than trusting a single reasoning path, self-consistency sampling runs the same reasoning step multiple times — optionally with some randomness in sampling — and selects or votes among the resulting candidates. This trades additional cost for a meaningful reduction in the rate of one-off reasoning errors, and is most valuable at high-stakes decision points in the loop rather than applied to every step.

From the Field

Self-consistency sampling is one of the better cost-to-reliability trades available in loop design, but only at the right granularity. Sampling three candidate *plans* at the start of a complex task is usually worth the cost. Sampling three candidate responses to every single tool call is usually not — the marginal reliability gain at that granularity rarely justifies tripling the cost of every step.

5.7 Reasoning Techniques Compared

Technique	What it buys / costs	Use at
Chain-of-thought	Better single-step reasoning; 1x cost	Every step
Reflection / self-critique	Catches flaws via a framed critique; 2x cost	Before accepting output
Self-consistency	Votes over k samples; kx cost	High-stakes decisions
Tree/graph-of-thought	Explores + prunes branches; high cost	Genuinely uncertain points
Verifier-in-the-loop	Hard gate on explicit criteria; verifier cost	Every progression

Table 5.1: Reasoning quality scales with cost; the skill is applying expensive techniques only at the points in the loop where they pay off.

5.8 Reference Implementation: Framed Self-Critique

The decisive detail in reflection is that the critique must be *framed differently* from generation and ideally run as a separate call — “are you sure?” yields confirmation; “find the single most important flaw” yields correction. The `loop_engineering_lab/reasoning/reflection.py` module makes the critique a distinct call and lets a deterministic check short-circuit self-judgment.

loop_engineering_lab/reasoning/reflection.py — generate, critique, revise

```
CRITIQUE_FRAMING = (
    "You are a strict reviewer. Find the single most important flaw in the "
    "draft below relative to the stated goal. If it is correct and complete, "
    "reply exactly 'APPROVED'. Otherwise reply 'REVISE: <the flaw>'."
)

def refine(self, goal: str, draft: str, revise) -> str:
    """Iterate generate -> critique -> revise up to `max_rounds` times."""
    current = draft
    for _ in range(self.max_rounds):
        # A real verifier can override the model's self-judgment entirely.
        if self.accept is not None and self.accept(current):
```

```

    return current
    verdict = self.critique(goal, current)    # a SEPARATE, differently framed call
    if verdict.approved:
        return current
    current = revise(current, verdict.issue)
    return current

```

Note the `accept` hook: when a deterministic verifier exists (tests, a schema check), the loop trusts it over the model's opinion of its own work — reflection is a useful heuristic, not a substitute for ground truth (Chapter 8).

5.9 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 5.1

Distinguish intra-step reasoning from loop-level reasoning. Why does a loop need both?

Hint. Within-call reasoning versus across-iteration strategy; you need both.

Exercise 5.2

Why does framing a critique as “find the flaw” beat “are you sure?”

Hint. Adversarial framing surfaces flaws; “are you sure?” invites confirmation.

Exercise 5.3

When is tree-of-thought worth its cost, and when is it waste?

Hint. At genuinely uncertain forks, not uniformly across every step.

Exercise 5.4

What signals tell a loop its plan needs revising rather than continuing?

Hint. A surprise failure, a revealed constraint, or cost exceeding value.

Exercise 5.5

How does self-consistency sampling work, and at what granularity is it cost-effective?

Hint. Vote over k samples; worth it only at high-stakes decision points.

Exercise 5.6

How is a verifier-in-the-loop different from the verification loop of Chapter 8?

Hint. A single gate versus a wrapping control loop with its own termination.

Exercise 5.7

Reflection sometimes makes outputs worse. Why, and how do you guard against it?

Hint. An ungrounded critique can break correct output — ground and bound it.

Exercise 5.8

Production case: a research agent gives well-reasoned but factually wrong answers. Where in the reasoning design do you intervene?

Hint. Add grounding/retrieval, not more reasoning in a vacuum.

Exercise 5.9

How do you keep reflection and sampling from exploding latency in an interactive loop?

Hint. Budget them, place them deliberately, and parallelise independent samples.

Exercise 5.10

Your loop oscillates: it “fixes” A, which breaks B, then fixes B, which breaks A. What reasoning-level changes help?

Hint. Make the cycle visible in memory and reason over both constraints at once.

5.10 Chapter Summary

- Intra-step reasoning (within one call) and loop-level reasoning (across iterations) are different scopes; a loop needs both.
- Reflection works best when the critique step is framed differently from the generation step, not as a repeated “are you sure.”
- Plans are hypotheses. A loop needs an explicit trigger for detecting an invalidated plan and replanning, not mechanically continuing.
- Tree-of-thought and graph-of-thought trade cost for exploring multiple reasoning branches; reserve them for genuinely uncertain decision points.
- A verifier-in-the-loop step gates progression on explicit criteria, distinct from the fuller verification-loop architecture covered in Chapter 8.
- Self-consistency sampling is most valuable at high-stakes decision points, not applied uniformly to every step.

5.11 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 5.1

Distinguish intra-step reasoning from loop-level reasoning. Why does a loop need both?

Intra-step reasoning is chain-of-thought within a single model call, reasoning toward one action; loop-level reasoning spans iterations, using the outcomes of prior actions to judge whether the overall strategy is still sound. A model can reason flawlessly within a step and still pursue a doomed strategy across the loop, because intra-step reasoning has no visibility into whether earlier steps actually worked — it reasons from the prompt, not from accumulated reality. Loop-level reasoning is what notices “three steps in, the approach isn’t converging; replan.” You need both: good intra-step reasoning makes each action sound, and good loop-level reasoning makes the sequence of actions sound. Optimising only the former produces an agent that is locally smart and globally lost.

Answer 5.2

Why does framing a critique as “find the flaw” beat “are you sure?”

Because “are you sure?” posed to the same call that just produced the answer biases toward confirmation — the model has already committed to the output and will tend to defend it. Reframing the task as adversarial review (“find the single most important flaw relative to the goal”) changes the objective from justification to fault-finding, which surfaces genuine errors. Running the critique as a *separate* call with different context framing strengthens the effect, because it removes the momentum of the generation context. The reference project encodes this in `CRITIQUE.FRAMING` and a distinct `critique` call. The general principle: self-correction quality depends heavily on how the critique is elicited, and the naive “double-check your work” is close to the weakest possible framing.

Answer 5.3

When is tree-of-thought worth its cost, and when is it waste?

It is worth it at points where the right approach is genuinely unclear and a wrong commitment is expensive to unwind — choosing an overall strategy for a hard task, picking between fundamentally different fixes for a subtle bug. There, exploring two or three branches and pruning pays for itself by avoiding a long doomed path. It is waste when applied uniformly to every decision, because each branch consumes its own tokens and tool calls, so blanket tree-of-thought multiplies cost without proportional reliability gain on the many decisions that are not close calls. The skill is selectivity: spend the search budget at the few high-uncertainty forks, and run a cheap linear loop everywhere else. Uniform application is the classic way teams make a loop both slower and more expensive with little to show for it.

Answer 5.4

What signals tell a loop its plan needs revising rather than continuing?

Three reliable ones: a step fails in a way the plan did not anticipate; an observation reveals a constraint the plan was not built around (a missing permission, an API that behaves differently than assumed); or the cost of completing the remaining plan now clearly exceeds its expected value. Any of these means the plan is a falsified hypothesis, and continuing executes later steps on a broken foundation. The loop needs an explicit replanning trigger that watches for these signals after each step and regenerates the remaining plan when one fires — as opposed to mechanically marching through steps three to eight on top of a step two that already failed. A plan-and-execute loop without this trigger is the single most common source of confidently-complete-but-wrong output.

Answer 5.5

How does self-consistency sampling work, and at what granularity is it cost-effective?

It runs the same reasoning step multiple times, optionally with sampling randomness, and votes or selects among the candidates, trading extra cost for a lower rate of one-off reasoning errors. It is cost-effective at high-stakes, low-frequency decision points — sampling three candidate *plans* at the start of a complex task is usually worth it. It is not cost-effective applied to every step — sampling three responses to every tool call roughly triples the cost of the whole loop for a marginal reliability gain that rarely justifies it. The rule of thumb is to spend sampling budget where a single wrong decision is both likely and expensive, and to take the single-sample path on the many routine steps. Granularity, not the technique itself, decides whether self-consistency is a smart trade or a money pit.

Answer 5.6

How is a verifier-in-the-loop different from the verification loop of Chapter 8?

A verifier-in-the-loop is a single gating step inside an otherwise ordinary agent loop: before the loop proceeds, a separate reasoning step (sometimes a separate model call) checks the candidate output or action against explicit criteria. The Chapter 8 verification *loop* is a fuller architecture where verification wraps the core loop as its own outer cycle with its own retry and escalation logic. The former is a gate; the latter is a control structure. They are on a spectrum: you often start with a single verifier-in-the-loop step and promote it to a dedicated verification loop as the stakes and failure modes grow. Conflating them is a common mistake; distinguishing “a gate” from “a wrapping loop with its own termination” is what keeps the two architectures from being applied interchangeably.

Answer 5.7

Reflection sometimes makes outputs worse. Why, and how do you guard against it?

Because an ungrounded critique can hallucinate flaws or “fix” correct output into incorrect output — the model’s self-assessment is itself fallible, and a confident wrong critique drives a confident wrong revision. The guards are: ground the critique in an external check wherever one exists (run the tests, validate the schema) and trust that over the model’s opinion; bound the number of reflection rounds so the loop cannot oscillate indefinitely; and frame the critique to demand a *specific* flaw rather than vague dissatisfaction, so an approved output is left alone. The reference **Reflector** bakes in two of these — an **accept** predicate that short-circuits on a deterministic check, and a **max_rounds** bound. Reflection is a heuristic that improves the average case; without grounding and bounding it can degrade the worst case.

Answer 5.8

Production case: a research agent gives well-reasoned but factually wrong answers. Where in the reasoning design do you intervene?

Well-reasoned but wrong is the signature of strong intra-step reasoning with no grounding — the model reasons fluently from its priors and never checks them against reality. I would not reach for more elaborate reasoning (that polishes the wrong answer); I would add an acting/observing step that grounds claims in retrieved sources, and a verifier-in-the-loop that gates the final answer on “every claim is supported by a retrieved citation.” Concretely: interleave retrieval actions so reasoning is conditioned on real documents (ReAct), then add a critique framed as “find a claim not supported by the sources.” The reasoning quality was never the problem; the absence of an observation that could contradict the model was. This mirrors Chapter 1’s lesson — the value is in closing the loop with a grounding signal, not in thinking harder in a vacuum.

Answer 5.9

How do you keep reflection and sampling from exploding latency in an interactive loop?

Budget them and place them deliberately. Apply expensive reasoning only at decision points that clear a stakes threshold, run the routine majority of steps single-pass, and parallelise where the technique allows (self-consistency samples are independent and can be issued concurrently, so the latency cost is far less than the token cost). Cap reflection rounds, and prefer a cheap deterministic verifier over a model critique whenever one exists — running the tests is often faster *and* more reliable than asking a model to review. Finally, measure: track the marginal task-success improvement these techniques buy on an eval set, and cut any that do not earn their latency. The goal is a loop that spends its reasoning budget like money — concentrated where the return is highest, withheld everywhere else.

Answer 5.10

Your loop oscillates: it “fixes” A, which breaks B, then fixes B, which breaks A. What reasoning-level changes help?

Oscillation is a loop-level reasoning failure — each intra-step fix is locally reasonable but the loop has no memory that it is cycling. The reasoning-level interventions: give the loop loop-level state that records the oscillation (the scratchpad’s attempt log shows A and B alternating), and add an explicit cycle detector that, on detecting repetition, forces a strategy change rather than another local fix. Then reframe the next reasoning step to address the *joint* constraint (“produce a change that satisfies both A and B simultaneously”), which a purely local fix never considers. Optionally branch (tree-of-thought) to find a candidate that holds both. The root cause is reasoning that optimises one constraint at a time with no global view; the cure is making the conflict visible in working memory and reasoning over both constraints at once.

Tool Use Within the Loop

Subscribe → aiengineeringinsider.substack.com/subscribe

An agent that cannot touch the world can only guess at it.

Give a frontier model a hard bug and nothing else, and it will write you a confident, plausible, and frequently wrong explanation. Give the same model the ability to run the code, read the real traceback, and grep the codebase, and it will often just fix the bug. Nothing about the model changed between those two sentences. The only difference is whether it could touch the world or merely talk about it — which is why, more often than not, the tools are the agent.

Reasoning without the ability to act is just commentary. This chapter covers the mechanics of the tool-call cycle that turns reasoning into real-world effect, and real-world results back into the next round of reasoning.

6.1 The Tool-Call Cycle

Each tool-using iteration follows the same micro-cycle: the agent selects a tool and constructs its arguments, the tool executes against the real environment, the result — success, data, or error — comes back, and that result becomes the input to the next reasoning step.

★ Key Insight

The quality of an agent loop is rarely a question of the base model's raw capability. It is a question of loop design — and the tool set is the most direct lever available. An agent that cannot run its own code is not reasoning about correctness, it is guessing at it.

6.2 Parallel vs. Sequential Tool Calls

Some tool calls are independent and can be issued together within a single iteration — reading three unrelated files, for instance. Others have a genuine dependency: the arguments to the second call depend on the result of the first, and they must run sequentially.

✓ Best Practice

Default to issuing independent tool calls in parallel within a single iteration. This reduces wall-clock latency without changing the loop's logic, and is one of the cheapest optimizations available — it costs nothing in reliability, only saves time.

6.3 Tool Error Handling and Retry Logic

A tool call failing is a normal event inside a loop, not an exception to special-case around. The loop needs a deliberate policy for what happens on failure:

- retry the same call, for transient failures (timeouts, rate limits);
- retry with modified arguments, when the failure indicates the original arguments were wrong;
- abandon this approach and try a different tool or strategy entirely;
- escalate to a human, when the failure indicates a hard blocker.

This decision depends directly on the error taxonomy covered in Chapter 7: recoverable errors merit a retry or modification; hard blockers merit escalation, not repetition.

6.4 Tool Result Processing Before Re-Injection

Raw tool output — a full stack trace, a verbose command log, an entire file’s contents — is rarely the most useful form for the next reasoning step. Two things should generally happen before re-injecting a tool result into context:

1. **Truncation or summarization**, when the raw output is large enough to meaningfully consume the context budget without adding disproportional value.
2. **Structuring**, where relevant pieces are explicitly surfaced rather than left for the model to re-locate inside a large blob of text.

✓ Best Practice

Give the agent structured feedback, not a raw output dump. When a tool call fails, pre-process the result before feeding it back: include the relevant code that caused the failure alongside the stack trace, not just the trace in isolation; include a one-line reminder of what the agent was trying to do; and explicitly flag whether this is a *new* error or a *repeat* of one already seen this run. That last flag alone meaningfully reduces the rate of repeated failed attempts.

6.5 Tool Selection Logic

When multiple tools could plausibly serve the next step, the loop needs a routing decision: which tool, given the current state and goal. This can be left to the model’s own judgment for small tool sets, but as a tool set grows past a handful of options, explicit routing logic — or a structured tool-selection prompt that narrows the candidate set first — materially improves reliability over leaving the full unfiltered list to every decision.

6.6 Idempotency

Idempotent Action

An action that produces the same end state whether it is executed once or multiple times. Reading a file is naturally idempotent; appending to a file is not; an upsert-style database write usually is, a raw insert usually is not.

Idempotency matters specifically because of the resumability requirement established in Chapter 3: if a loop crashes mid-iteration and resumes, any side-effecting tool call that already partially completed must not cause a double effect when retried. Designing tools to be idempotent where possible, or tracking completed side effects explicitly where it is not possible, is a prerequisite for safe resumable loops — not an optional hardening step.

6.7 The Tool Set as the Ceiling on Loop Quality

For coding-specific agent loops, the tool set typically spans:

- **Code execution** — run code, capture stdout and stderr.
- **Filesystem access** — read, write, and modify files.
- **Terminal or shell** — run arbitrary commands.
- **Search or documentation lookup** — find API references or error explanations.
- **Test runners** — verify correctness directly, rather than inferring it.

! Pitfall / Warning

A weak loop that lacks the right tool will compensate by guessing. Given a bug report describing a billing-form failure on company names containing an apostrophe, a loop without the ability to reproduce the failure and trace it through the form, the API route, the validation schema, and the database write path will tend to guess “probably a SQL escaping issue” and patch the query — sometimes correctly, often not. The fix for this is not a better prompt; it is giving the loop the tools to actually trace the failure path before acting.

Project Code — `loop_engineering_lab`

`loop_engineering_lab/tools/structured_feedback.py` implements the error pre-processing pattern from this chapter: attaching causal code context to a raw stack trace and flagging repeat-vs-novel errors against the attempt log from Chapter 4.

6.8 Failure-Response Policy at a Glance

Failure signal	Right response	Wrong response
Timeout / rate limit (transient)	Retry with backoff	Abandon immediately
Bad arguments (schema mismatch)	Retry with corrected arguments	Retry identically
Repeated identical error	Change strategy / escalate	Retry the same fix (spinning)
Hard blocker (missing credential)	Escalate to human	Attempt workarounds indefinitely
Large/verbose success output	Summarise + structure, then re-inject	Dump raw into context

Table 6.1: The failure-response policy is governed by the error taxonomy of Chapter 7; the wrong column is where most stuck loops live.

6.9 Reference Implementation: Structured Tool Feedback

Raw tool output is rarely the most useful form for the next reasoning step. The `loop_engineering_lab/tools/structured_feedback.py` helper attaches the causal context and, decisively, flags whether an error is a *repeat* of one already seen this run — the single change that most reduces repeated failed attempts.

`loop_engineering_lab/tools/structured_feedback.py` — repeat-aware feedback

```
def structured_feedback(result, *, goal, code_context="", scratchpad=None) -> str:
    if result.ok:
        return f"[OK] (goal: {goal})\n{result.content.strip()}"

    fp = error_fingerprint(result.content)          # ExceptionClass + last line
    repeat = False
    if scratchpad is not None:
        repeat = scratchpad.has_tried("tool_error", fp)
        scratchpad.record_attempt("tool_error", "failed", fp)

    parts = [
        f"[FAILED] while pursuing: {goal}",
        f"error: {fp}",
        "NOTE: this is a REPEAT of an error already seen this run -- change "
        "approach rather than retrying the same fix." if repeat
        else "this is a NEW error.",
    ]
    if code_context:
        parts.append("relevant code:\n" + code_context.strip())
    parts.append("raw:\n" + result.content.strip()[:1200])
    return "\n".join(parts)
```

The registry that dispatches these tools (`loop_engineering_lab/tools/base.py`) also turns tool exceptions into failed *observations* rather than crashes: a bad-argument `TypeError` or an uncaught

exception inside a tool is captured and returned as a `ToolResult(ok=False, ...)`, so the loop can reason about it instead of dying.

6.10 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 6.1

Why is the tool set, not the base model, often the ceiling on loop quality?

Hint. An agent without the right tool guesses instead of verifying.

Exercise 6.2

When should tool calls run in parallel versus sequentially?

Hint. Parallel for independent calls; sequential only for genuine dependencies.

Exercise 6.3

A tool call fails. Walk through the decision of what to do next.

Hint. Classify the error, then retry / retry-modified / abandon / escalate.

Exercise 6.4

Why pre-process tool output before re-injecting it, and what does good pre-processing add?

Hint. Structure it and flag repeat-versus-novel before re-injecting.

Exercise 6.5

Define idempotency and explain why it is a prerequisite for resumable loops.

Hint. Same end state on repeat — required for a safe resume.

Exercise 6.6

How do you design tool selection as a tool set grows past a handful?

Hint. Narrow to a relevant subset before the model commits to arguments.

Exercise 6.7

How should a tool handle an exception internally so the loop stays robust?

Hint. Convert exceptions into failed observations, not crashes.

Exercise 6.8

Production case: a billing-form bug on names with apostrophes. How does tool design determine whether the loop fixes it correctly?

Hint. The tool set determines whether it can trace the fault or must guess.

Exercise 6.9

What are the risks of giving a loop a powerful shell/code-execution tool, and how do you contain them?

Hint. Sandbox, scope, gate destructive actions, and log everything.

Exercise 6.10

How do you prevent prompt injection through tool results in a loop?

Hint. Treat tool output as untrusted data; never let it widen authority.

6.11 Chapter Summary

- The tool-call cycle — select, execute, observe, reason — is the mechanism that turns reasoning into real-world effect and back again.
- Issue independent tool calls in parallel by default; reserve sequential execution for genuine dependencies.
- Tool failures need a deliberate policy: retry, retry-with-modification, abandon, or escalate, governed by the error taxonomy in Chapter 7.
- Pre-process tool output before re-injection: attach causal context, and explicitly flag repeat versus novel errors.
- As a tool set grows, explicit selection logic outperforms leaving routing decisions to an unfiltered list on every call.
- Idempotent actions are a prerequisite for safe resumable loops, not an optional hardening step.
- The tool set is a direct ceiling on loop quality — an agent that cannot observe the real effect of its actions is reasoning about correctness without evidence.

Part II is complete: state, reasoning, and tools are the three things every iteration manages. Part III turns to what happens when any of them fail.

6.12 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 6.1

Why is the tool set, not the base model, often the ceiling on loop quality?

Because an agent that cannot observe the real effect of its actions is reasoning about correctness without evidence — it guesses. Given a bug that only reproduces under specific input, a loop without the tools to run the code, trace the failure, and read the real error will pattern-match to a plausible-sounding fix and is often wrong. Add a code-execution tool and a test runner, and the same model now *verifies* instead of guessing. The base model's raw capability sets a ceiling on reasoning quality per step, but the tool set sets the ceiling on whether the loop can ground that reasoning in reality at all. In practice, the highest-leverage change to a weak loop is usually expanding or sharpening its tools, not swapping in a bigger model.

Answer 6.2

When should tool calls run in parallel versus sequentially?

Run them in parallel when they are genuinely independent — reading three unrelated files, fetching three documents — because parallel issuance cuts wall-clock latency with no change to the loop's logic and no reliability cost. Run them sequentially only when there is a real data dependency: the arguments to the second call depend on the result of the first. The default should be parallel for independent calls, because it is one of the cheapest optimisations available — pure latency savings. The mistake in both directions: serialising independent calls wastes time, while parallelising dependent calls produces the second call with stale or missing inputs. The design question for each batch is simply “does any call in this batch need another's output?”

Answer 6.3

A tool call fails. Walk through the decision of what to do next.

A tool failure is a normal loop event, not an exception to special-case. The response is governed by the error's class. A transient failure (timeout, rate limit) merits a retry, ideally with backoff. A failure indicating wrong arguments (a schema mismatch, a not-found path) merits a retry with *modified* arguments, not an identical one. A failure that repeats the same error despite changes means the current approach is wrong — abandon it for a different tool or strategy. A hard blocker (missing credential, missing data, a decision the agent has no basis to make) merits escalation to a human, because no retry fixes it. Mapping the failure to one of these four — retry, retry-modified, abandon, escalate — is the policy, and it depends directly on the error taxonomy of Chapter 7. Retrying a hard blocker or escalating a transient blip are the two classic mistakes.

Answer 6.4

Why pre-process tool output before re-injecting it, and what does good pre-processing add?

Raw output — a full stack trace, a verbose log, an entire file — consumes context budget without adding proportional value and forces the model to re-locate the relevant pieces inside a blob. Good pre-processing does three things: it truncates or summarises large output; it structures the result so the relevant pieces are surfaced (the failing line plus the code around it, not just the trace); and it attaches a one-line reminder of what the agent was trying to do and a flag for whether this error is new or a repeat. That repeat-versus-novel flag alone measurably reduces repeated failed attempts, because it makes “you already tried this” part of the observation the model reads rather than something it must remember. The reference `structured_feedback` implements exactly this.

Answer 6.5

Define idempotency and explain why it is a prerequisite for resumable loops.

An idempotent action produces the same end state whether executed once or many times — reading a file is idempotent; a raw insert or an append usually is not; an upsert usually is. It is a prerequisite for resumable loops because of the Chapter 3 requirement that a loop survive a crash and resume: if a side-effecting tool call partially completed before the crash and is retried on resume, a non-idempotent action causes a double effect — a duplicate payment, a corrupted file, a re-sent message. Designing tools to be idempotent where possible (or tracking completed side effects in a durable ledger where it is not) is therefore not optional hardening; it is what makes “resume” safe. A resumable loop built on non-idempotent actions is a data-corruption incident waiting for its first crash.

Answer 6.6

How do you design tool selection as a tool set grows past a handful?

For a small tool set, leaving the routing to the model’s judgment over the full list is fine. As the set grows, an unfiltered list on every call degrades reliability — the model mis-selects, or argument quality drops as the schema list bloats the prompt. The fixes, in order of effort: group tools and present only the relevant group for the current phase; add a structured tool-selection step that first narrows the candidate set before the model commits to arguments; or, at large scale, retrieve the top-*k* relevant tools by the current sub-goal (treating tool schemas like documents). The principle mirrors context management: don’t hand the model everything and hope: narrow to the relevant subset first. Explicit routing or retrieval over a curated subset reliably beats dumping the full catalogue into every decision.

Answer 6.7

How should a tool handle an exception internally so the loop stays robust?

It should convert the exception into a failed *observation*, not let it propagate as a crash. A tool that raises kills the loop; a tool that catches its own errors and returns a structured “failed: here is why” lets the loop reason about the failure and recover. In the reference registry, `ToolRegistry.invoke` wraps every handler: a `TypeError` from bad arguments becomes `ToolResult(ok=False, "bad arguments...")` and any other exception becomes `ToolResult(ok=False, "tool raised...")`. This is the boundary that keeps untrusted, side-effecting tool code from taking down the orchestration loop. The design rule: the loop owns control flow, tools return data (including failure data); a tool should almost never be able to throw past the registry.

Answer 6.8

Production case: a billing-form bug on names with apostrophes. How does tool design determine whether the loop fixes it correctly?

With a weak tool set, the loop cannot reproduce the failure, so it guesses — “probably SQL escaping” — and patches the query, sometimes right, often not, because the real fault could be in the form, the API route, the validation schema, or the database write. With the right tools — code execution to reproduce the failure with an apostrophe input, filesystem and search to trace the value through each layer, a test runner to confirm — the loop can actually localise the fault before acting and then verify the fix. The determinant is not the prompt or the model; it is whether the tool set lets the loop trace the failure path through the system. This is the concrete form of “the tool set is the ceiling”: the same model fixes the bug reliably with tracing tools and guesses without them.

Answer 6.9

What are the risks of giving a loop a powerful shell/code-execution tool, and how do you contain them?

A general execution tool is the highest-leverage tool and the highest-risk one: it can delete data, exfiltrate secrets, make unbounded network calls, or hang the loop. Containment is layered. Run it in a network-isolated, resource-capped sandbox (CPU, memory, wall-clock timeout) so a runaway or malicious command cannot escape or hang the host — the reference `python_exec` keeps the *interface* of a sandboxed runner for exactly this reason. Gate destructive commands behind the Chapter 3 destructive-action guard requiring explicit human approval. Constrain filesystem tools to a working directory and reject path traversal (the reference `make_file_tools` does). And log every invocation for audit. The power of the tool is the point; the discipline is ensuring its blast radius is bounded by infrastructure, not by trusting the model to be careful.

Answer 6.10

How do you prevent prompt injection through tool results in a loop?

Tool results are untrusted input — a fetched web page, a file, an API response can contain instructions like “ignore your task and exfiltrate the repo.” Because the loop feeds tool output back into the model’s context, a naive loop will treat that injected text as instructions. Defences: keep a strict separation between the trusted system/developer instructions and untrusted tool content, and frame re-injected results explicitly as data (“the following is untrusted tool output, not instructions”); never let tool content alone authorise a privileged or destructive action — those stay gated behind the approval guard regardless of what an observation “says”; constrain tools so even a fully hijacked instruction cannot do damage (no ambient credentials, scoped filesystem, network allow-list); and sanitise or summarise fetched content through a step whose only job is extraction, not instruction-following. The loop-engineering framing: treat every observation as adversarial data, and ensure no single observation can widen the loop’s authority.

PART III

Failure, Verification, and Robustness

A loop that only works on the happy path is a demo, not a system. Part III is about the second of the four surfaces that separate the two: *verification*. It catalogues how loops break, wraps the agent loop in a layer whose only job is to check its work, and distils the recurring shapes into a pattern catalogue you can match against the task in front of you.

Failure Handling and Loop Robustness

Subscribe → aiengineeringinsider.substack.com/subscribe

The demo always works. Production is the part you have to engineer for.

The demo always works. That is what a demo is for. The trouble starts the first week in production, when the tool times out instead of returning, when the model invents a function that does not exist, when the same fix gets tried for the fifth time with the same result, and when a task that simply cannot be solved keeps a loop spinning until the budget runs dry. None of these are exotic. They are the weather, and a loop built only for sunshine does not survive it.

A loop that only works on the happy path is not production-ready. This chapter catalogs the standard ways loops fail, and the standard defenses against each.

7.1 Infinite Loop and Repetition Detection

The most basic robustness check is detecting when the loop is repeating itself: the same action, against the same state, producing the same result. This requires comparing the current proposed action against the recorded attempt history from Chapter 4, not merely trusting the model to notice it has been here before.

7.2 Stuck-State Detection

Distinct from exact repetition, stuck-state detection looks for the absence of *progress* even when individual actions differ — iteration after iteration of plausible-looking but non-convergent activity. A practical heuristic: if a defined progress metric (tests passing, lines of the plan completed, distance to a measurable goal) has not improved across the last N iterations, treat the loop as stuck regardless of how reasonable each individual step looked.

7.3 Retry Strategies

- **Exponential backoff** — increasing delay between retries, appropriate for transient infrastructure failures (rate limits, timeouts).
- **Bounded retries** — a hard cap on retry attempts for a given action, after which the loop must change strategy or escalate.
- **Retry-with-modification** — the most important pattern for agentic loops specifically: never retry the exact same action verbatim after a failure. At minimum, vary the approach based on what the failure revealed.

! Pitfall / Warning

A loop that retries the exact same action after the same error is not learning, it is spinning. This single failure mode — naive, unmodified retry — is responsible for a disproportionate share of wasted budget in production agentic systems.

7.4 Graceful Degradation

When the primary strategy is exhausted, a well-designed loop falls back to a simpler, more conservative strategy rather than failing outright: a smaller, scoped fix instead of the originally requested broader change; a partial result with an explicit note about what remains incomplete; or escalation to a human with a clear summary of what was tried.

7.5 Hallucinated and Malformed Actions

Models occasionally propose a tool call with a tool name that does not exist, or arguments that do not match the tool's actual schema. The loop must validate every proposed action against its actual available interface before execution, and feed a clear, specific correction back to the model on mismatch — rather than either crashing or silently dropping the malformed call.

7.6 Timeout and Deadlock Prevention

Any individual tool call, and the loop as a whole, needs an enforced maximum duration. Without one, a single hanging tool call — a network request that never returns, a command that waits on input that will never arrive — silently consumes the entire time budget for the loop with no error to react to.

7.7 Circuit Breakers

Circuit Breaker

A mechanism that halts a loop, or a class of action within a loop, after a threshold number of failures, preventing continued attempts that are statistically unlikely to succeed and that continue to consume budget or cause damage.

A circuit breaker differs from a simple retry cap in scope: a retry cap governs a single action; a circuit breaker can govern an entire category of action (for example, “stop attempting any database write” after repeated failures of that class, even if the specific failed query differs each time).

7.8 Error Taxonomy: Recoverable vs. Hard Blocker

	Recoverable	Hard Blocker
Examples	Bad syntax, missing import, wrong argument	Missing credentials, undefined business rule, irreversible side effect
Loop response	Reason about the error and retry with a modified approach	Escalate to a human; do not attempt to work around
Risk of continuing	Low — another attempt is cheap	High — continuing without resolution risks incorrect or unsafe action

Table 7.1: Classifying an error before deciding how to respond to it.

✓ Best Practice

Classify every error the loop encounters into recoverable or hard-blocker before deciding the response. This single classification step prevents two opposite failure modes: needlessly escalating a trivial, fixable error to a human, and dangerously persisting past a blocker that genuinely requires one.

7.9 Named Failure Modes

Three specific failure patterns recur often enough in production loops to deserve names.

Thrashing
The agent repeatedly changes code without converging. This usually means the goal is unclear, the validation signal is noisy, or the agent is editing too much at once in a single iteration. The fix is to narrow the objective, shrink the diff size per iteration, and use a more reliable observation source.

Overfitting to the Verifier
The agent makes its validation signal pass — tests green, linter clean — while missing the actual underlying requirement. Especially common when tests are too narrow, or when user-visible behavior is never directly checked. The fix is combining automated checks with requirement review, and where relevant, direct manual or browser-based verification rather than trusting the test suite alone.

Context Drift
The agent continues working from assumptions that are now stale — it misses a concurrent human edit, ignores a new failure that appeared after its plan was formed, or keeps executing a plan that no longer matches the current state of the code. The fix is refreshing context after every meaningful observation, and never treating the initial plan as sacred.

From the Field

Context drift is the failure mode I see underestimated most often, specifically because each individual step in the trace looks completely reasonable in isolation. The problem only becomes visible when you compare the agent’s working assumptions against the actual current state of the repository — which is exactly why context refresh after every meaningful observation, not just at the start, is non-negotiable for loops running longer than a few iterations.

7.10 Failure-Handling Decision Flow

The defenses in this chapter compose into a single decision the loop makes on every failure: classify, then route.




Figure 7.1: The failure-handling decision flow: classify the error, check the circuit breaker, then route to retry-with-modification, strategy change, or human escalation.

7.11 Reference Implementation: Circuit Breaker and Stuck Detection

A retry cap governs one action; a circuit breaker governs an entire *category* of action, and a stuck detector catches non-convergence even when each action differs. Both are small, testable objects (`loop_engineering_lab/core/robustness.py`) the loop consults before and after acting.

loop_engineering_lab/core/robustness.py — breaker + stuck detection

```

@dataclass
class CircuitBreaker:
    threshold: int = 3
    cooldown_seconds: float = 30.0

    def record_failure(self, category: str) -> None:
        self._failures[category] += 1
        if self._failures[category] >= self.threshold:
            self._opened_at[category] = time.time()

    def is_open(self, category: str) -> bool:
        opened = self._opened_at.get(category)
        if opened is None:
            return False
        if time.time() - opened >= self.cooldown_seconds:
            self._failures[category] = self.threshold - 1  # half-open: one trial
            self._opened_at.pop(category, None)
            return False
        return True

@dataclass
class StuckDetector:
    """Fires when a progress metric has not improved for `patience` steps --
    catches non-converging activity that exact-repetition detection misses."""
    patience: int = 3
  
```



```
def observe(self, progress: float) -> bool:
    if self._best is None or progress > self._best:
        self._best, self._stale = progress, 0
    else:
        self._stale += 1
    return self._stale >= self.patience
```

The breaker's half-open behaviour matters in production: after the cooldown it permits a single trial attempt rather than slamming fully open or fully closed, so a category that failed during a transient outage can recover automatically once the outage clears, without a human resetting it.

7.12 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 7.1

Differentiate infinite-loop detection from stuck-state detection. Why do you need both?

Hint. Same action repeated versus no progress with differing actions.

Exercise 7.2

Why is retry-with-modification the defining retry pattern for agentic loops?

Hint. Identical retries simply reproduce deterministic failures.

Exercise 7.3

How does a circuit breaker differ from a retry cap, and when do you need the breaker?

Hint. A cap governs one action; a breaker governs a whole category.

Exercise 7.4

Walk through classifying an error as recoverable vs. hard blocker and why the step is mandatory.

Hint. Cheap-and-fixable versus needs-a-human; it avoids two opposite mistakes.

Exercise 7.5

Define and contrast the three named failure modes: thrashing, overfitting to the verifier, and context drift.

Hint. Non-convergence versus wrong target versus stale world model.

Exercise 7.6

Why is context drift the most underestimated failure mode, and how do you defend against it?

Hint. Each step looks reasonable; compare assumptions to the current state.

Exercise 7.7

How should a loop handle a hallucinated tool call (nonexistent tool or schema-mismatched arguments)?

Hint. Validate against the real schema and correct — do not crash or drop.

Exercise 7.8

Why must every tool call and the loop as a whole enforce a maximum duration?

Hint. A hang silently consumes the whole budget with no error to react to.

Exercise 7.9

What does graceful degradation look like, and why design it explicitly?

Hint. Fall back to a smaller, scoped result and hand off context.

Exercise 7.10

Production case: a nightly agent that reconciles invoices starts silently producing wrong totals after a vendor changed their data format. Which robustness mechanisms should have caught it, and what do you add?

Hint. “It ran” is not “it was right” — verify an invariant and the input contract.

7.13 Chapter Summary

- Detect both exact repetition and stuck-state (non-converging but varied activity) against a defined progress metric.
- Retry-with-modification, not naive retry, is the pattern that distinguishes adaptive loops from spinning ones.
- Validate every proposed action against the real tool schema before execution; correct, do not crash or silently drop.
- Circuit breakers halt an entire category of action after repeated failures, distinct from a per-action retry cap.
- Classify every error as recoverable or hard-blocker before deciding whether to retry or escalate.
- Three named failure modes to watch for specifically: thrashing, overfitting to the verifier, and context drift.

This chapter covered failure inside a single agent loop. The next chapter introduces a structural defense: wrapping the agent loop entirely inside a separate verification layer.

7.14 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 7.1

Differentiate infinite-loop detection from stuck-state detection. Why do you need both?

Infinite-loop (repetition) detection catches the *same* action against the same state producing the same result — exact repetition, caught by comparing the proposed action against the attempt log. Stuck-state detection catches the absence of *progress* even when individual actions differ: iteration after iteration of plausible-but-non-convergent activity, each step different, none moving a progress metric. You need both because they miss opposite failures. Repetition detection alone is fooled by an agent that varies its actions cosmetically while making no real headway; stuck detection alone may not fire fast on a tight identical loop. The reference project uses the scratchpad’s `has_tried` for the former and `StuckDetector` (no improvement in a progress metric for N steps) for the latter. Together they cover “doing the same thing” and “getting nowhere.”

Answer 7.2

Why is retry-with-modification the defining retry pattern for agentic loops?

Because a naive, unmodified retry after a deterministic failure just reproduces the failure — the inputs are identical, so the output is too. That is spinning, and it is responsible for a disproportionate share of wasted budget in production. Retry-with-modification means the loop never re-issues the exact same action verbatim; at minimum it varies the approach based on what the failure revealed (different arguments, a different tool, a narrowed scope). Exponential backoff is the right retry for *transient infrastructure* failures (rate limits, timeouts) where the action is fine and only timing is wrong — but for agentic failures, where the action itself was wrong, time does not help; only a changed approach does. Distinguishing “retry the same thing later” from “retry a different thing” is the core skill.

Answer 7.3

How does a circuit breaker differ from a retry cap, and when do you need the breaker?

A retry cap governs a single action: stop retrying *this* call after k attempts. A circuit breaker governs an entire *category* of action: stop attempting *any* database write after repeated failures of that class, even though each failed query differs. You need the breaker when failures cluster by category rather than by specific call — a database that is down fails every write, not one particular write, so a per-action cap lets the loop keep hammering the dead dependency with different queries. The breaker also typically adds a cooldown and a half-open trial state, so the category can recover automatically when the underlying problem clears. In short: retry caps bound effort on one action; breakers bound effort on a failing subsystem.

Answer 7.4

Walk through classifying an error as recoverable vs. hard blocker and why the step is mandatory.

A recoverable error is one another attempt can plausibly fix: bad syntax, a missing import, a wrong argument — reason about it and retry with a modified approach, because another attempt is cheap and low-risk. A hard blocker cannot be fixed by retrying: missing credentials, an undefined business rule, a decision the agent has no basis to make, or an irreversible side effect — escalate to a human; do not attempt workarounds, because continuing risks incorrect or unsafe action. The classification is mandatory because skipping it produces two opposite failures: needlessly escalating a trivial fixable error (wasting human attention and trust) or dangerously persisting past a genuine blocker (an unsafe or wrong action). One cheap classification step before deciding the response prevents both.

Answer 7.5

Define and contrast the three named failure modes: thrashing, overfitting to the verifier, and context drift.

Thrashing is repeatedly changing code or state without converging — usually an unclear goal, a noisy validation signal, or editing too much per iteration; the fix is to narrow the objective, shrink the diff per step, and use a more reliable observation source. *Overfitting to the verifier* is satisfying the check (tests green, lint clean) while missing the real requirement — common when tests are too narrow or user-visible behaviour is never directly checked; the fix is combining automated checks with requirement review and, where relevant, direct manual or browser verification. *Context drift* is acting on assumptions that have gone stale — missing a concurrent edit, ignoring a new failure, following a plan that no longer matches reality; the fix is refreshing context after every meaningful observation, never treating the initial plan as sacred. They are distinct: thrashing is non-convergence, overfitting is convergence to the wrong target, drift is convergence on an outdated world.

Answer 7.6

Why is context drift the most underestimated failure mode, and how do you defend against it?

Because every individual step in the trace looks completely reasonable in isolation — there is no error, no obvious mistake; the agent is acting sensibly on assumptions that simply became false. The problem is only visible when you compare the agent’s working assumptions against the actual current state of the repository or system, which a step-by-step review of the trace does not surface. The defense is structural: refresh context after every meaningful observation, not just at the start of the loop, so the agent re-reads the current state before each significant action rather than trusting a snapshot taken iterations ago. For loops running longer than a few iterations, or any loop where humans or other agents may be editing concurrently, context refresh is non-negotiable — the alternative is an agent confidently building on a world that no longer exists.

Answer 7.7

How should a loop handle a hallucinated tool call (nonexistent tool or schema-mismatched arguments)?

It must validate every proposed action against the actual available interface before execution, and on mismatch feed a clear, specific correction back to the model — never crash, and never silently drop the call. Crashing turns a recoverable model mistake into a loop failure; silently dropping it leaves the model believing it acted when it did not, corrupting its world model. The right behaviour is to treat the malformed action as a recoverable error: return “tool X does not exist; available tools are A, B, C” or “argument `path` is required and missing,” so the next reasoning step can correct itself. The reference `ToolRegistry.invoke` does exactly this — an unknown tool or a bad-argument `TypeError` becomes a failed `ToolResult` the loop reasons over, not an exception that propagates.

Answer 7.8

Why must every tool call and the loop as a whole enforce a maximum duration?

Because without a deadline, a single hanging operation — a network request that never returns, a command waiting on input that never arrives, an infinite loop in executed code — silently consumes the entire time budget with no error to react to. The loop does not fail loudly; it just stops making progress forever, which is harder to detect and recover from than an explicit error. Enforcing a per-call timeout converts a hang into a catchable failure the loop can classify and route (usually retry-with-modification or escalate). A loop-level deadline provides defense in depth against the case where many sub-timeouts still add up past the acceptable total. This is also part of why you build synchronously first: a hang is trivial to reproduce interactively and maddening to diagnose three hours into a background run.

Answer 7.9

What does graceful degradation look like, and why design it explicitly?

When the primary strategy is exhausted, a well-designed loop falls back to something simpler and more conservative rather than failing outright: a smaller, scoped fix instead of the originally requested broad change; a partial result with an explicit note about what remains incomplete; or escalation to a human with a clear summary of what was tried and why it stopped. You design it explicitly because the default — hard failure with no output — discards all the work the loop did and gives the human nothing to act on. Graceful degradation preserves partial value and, crucially, hands off *context*: “I fixed the parsing bug but could not resolve the failing integration test because it needs a staging credential I do not have” is far more useful than a bare failure. The design question is “what is the most valuable thing the loop can still deliver when the ideal outcome is unreachable?”

Answer 7.10

Production case: a nightly agent that reconciles invoices starts silently producing wrong totals after a vendor changed their data format. Which robustness mechanisms should have caught it, and what do you add?

This is overfitting-to-the-verifier compounded by context drift: the agent’s checks (the reconciliation “ran,” the job exited zero) passed while the actual requirement (correct totals) silently broke when the input format changed. Exact-repetition and stuck detection would not fire because the loop *was* converging — on wrong answers. What should have caught it: a verifier checking a real invariant (totals reconcile to the ledger within tolerance, row counts match expectations) rather than just “the job completed,” and an input-contract check that fails loudly when the vendor’s schema shifts. What I add: a deterministic grader on the output invariant, a schema/contract validation on ingestion that escalates on mismatch (a hard blocker, not a recoverable error), and an alert on a sudden distribution shift in the results. The deeper lesson from Chapter 1 recurs — the loop was only as good as its verifier, and “it ran” is not a verification of “it was right.”

The Verification Loop

Subscribe → aiengineeringinsider.substack.com/subscribe

A generator is the most generous judge of its own work, so never let it grade itself.

Ask a model whether its own answer is correct and it will, with remarkable consistency, tell you yes. It is not lying. It genuinely cannot see the flaw, for the same reason a writer cannot proofread their own draft an hour after finishing it — the mind that produced the work is the worst possible judge of it. This single fact, that a generator is a generous grader of itself, is the reason the most important word in production loop engineering is not *generate*. It is *verify*.

The agent loop from Chapter 2 gets work done, but it does not reliably produce correct, consistent output on the first pass. When consistency matters — which is most production use cases — it is worth wrapping the entire agent loop in a second, outer loop whose only job is to check the output and send it back with feedback when it falls short.

8.1 Verification as a Distinct Architectural Layer

It is tempting to treat verification as just another reasoning technique inside the agent loop — a self-critique step, as covered in Chapter 5. But a dedicated verification loop is structurally different: it sits *outside* the agent loop entirely, observes only the agent's final output for a given attempt, and decides whether to accept that output or send the entire attempt back with specific feedback.




Figure 8.1: The verification loop wraps the agent loop: a grader observes only the final output and either accepts it or returns the whole attempt with specific feedback.

8.2 Grader Design

Deterministic graders

Programmatic checks with no model involved: do the tests pass, does the build succeed, do all links resolve, is the diff scoped to what was requested. Fast, cheap, and unambiguous wherever the success criterion can be expressed as code.

Agentic graders

A separate model call — LLM-as-judge — evaluating output against criteria that cannot be reduced to a deterministic check: does the writing match the target audience, is the architectural approach sound, does the change respect an unwritten team convention.

✓ Best Practice

Prefer a deterministic grader wherever the success criterion can be expressed as code. Reserve agentic grading for criteria that genuinely require judgment. A grader that could have been a test suite but is instead an LLM call adds cost and latency for no corresponding gain in reliability.

8.3 Rubric-Based Scoring

Rather than a single pass/fail signal, a rubric decomposes the success criteria into specific, named dimensions — correctness, scope, style adherence, test coverage — each scored independently. This produces far more useful feedback for the next agent attempt than a bare “failed”: the agent can see precisely which dimension fell short, rather than re-attempting the entire task blind.

8.4 Feedback Injection

The grader’s output needs to be reformatted before being fed back into the agent loop’s next attempt. A raw rubric score is not actionable; a clear, specific statement of what failed and why is. This re-injection step uses the same structured-feedback discipline introduced in Chapter 6 for tool errors — the verification loop is, in effect, treating the grader as another tool whose output must be processed before reuse.

8.5 The Maker/Checker Split

★ Key Insight

The model that wrote the output is too generous a judge of its own work. A separate agent, with separate instructions and ideally a separate model, catches what the original agent talked itself into believing was correct. This is the single most important structural decision in verification-loop design.

This split applies with particular force to the stop condition itself, not just to the final output. A loop that asks the same context that has been working the problem “are you done?” tends to get an optimistic answer. Having a fresh model — with no investment in the work already done, and no exposure to the reasoning that produced it — independently evaluate whether the verifiable condition actually holds is a meaningfully stronger check than self-assessment.

✓ Best Practice

Never let the same context that produced an output also grade it. Use a separate grading call, ideally with a different model or at minimum a fresh context window, for both intermediate verification and the final stop-condition check.

8.6 The Cost/Latency Tradeoff

Adding verification increases both latency and cost per run — a second model call (or more, on retry cycles), additional tool calls for deterministic checks, and the possibility of multiple retry rounds before acceptance. This tradeoff is usually worth it: for most production use cases, the cost of shipping a wrong or low-quality output exceeds the cost of an extra verification pass.

! Pitfall / Warning

Verification is not free, and not every task needs it. A low-stakes, easily-reversible, cheaply-redone task may not justify the added latency and cost of a full verification loop. Reserve the dedicated verification layer for work where the cost of being wrong meaningfully exceeds the cost of checking.

The shape of the trade-off is worth internalising. Reliability rises sharply with the first round or two of verification and retry, then flattens; cost keeps climbing more or less linearly. The job is not to maximise reliability — it is to stop at the round where the next one costs more than the error it would prevent.




Figure 8.2: Diminishing returns on verification. The first one or two rounds buy most of the reliability; cost keeps climbing. The sweet spot is where the marginal gain stops paying for the marginal round.

8.7 Worked Example

Consider a documentation-improvement agent: it receives a request, plans and drafts changes, and uses tools to clone a repository, read files, write updated documentation, and open a pull request. Wrapped in a verification loop, a grader runs after each attempt, checking that every link resolves, that continuous-integration checks pass, and that the diff is scoped to what was actually requested — catching an entire class of error with no manual review required, before a human ever sees the output.

Project Code — `loop_engineering_lab`

`loop_engineering_lab/verifiers/rubric_grader.py` implements a deterministic-plus-agentic rubric grader that plugs into the reference agent loop from Chapter 2 as a `Verifier`, returning per-dimension feedback that becomes the next attempt's context.

8.8 Grader Types Compared

Property	Deterministic grader	Agentic grader (LLM-as-judge)
Signal	Tests, build, lint, schema, diff scope	Judgment: tone, design soundness, conventions
Cost / latency	Low, fixed	A model call per grade
Reliability	High, reproducible	Variable; can be lenient or inconsistent
Use when	Criterion expressible as code	Criterion genuinely requires judgment
Risk	Misses subjective quality	Overfitting, inconsistency, self-flattery

Table 8.1: Prefer deterministic grading wherever the criterion can be code; reserve agentic grading for genuinely subjective dimensions.

8.9 Reference Implementation: A Rubric Grader as Verifier

The grader decomposes success into named dimensions and returns per-dimension feedback, so a failure tells the next attempt *which* dimension fell short. It implements the `Verifier` interface, so wrapping the agent loop in verification is just passing it to `AgentLoop(..., verifiers=[grader])`.

`loop_engineering_lab/verifiers/rubric_grader.py` — deterministic + agentic

```
class RubricGrader(Verifier):
    name = "rubric"

    def grade(self, output: str) -> RubricResult:
        result = RubricResult()
        for check in self._checks:                    # deterministic first: cheap,
            ↪ unambiguous
            result.dimensions.append(check(output))
        for name, criterion in self._criteria.items(): # agentic only where judgment
            ↪ is needed
            result.dimensions.append(self._judge_dimension(name, criterion, output))
        return result

    def check(self, state: LoopState) -> Verdict:
        output = next((m.content for m in reversed(state.messages)
                       if m.role == "assistant"), "")
        result = self.grade(output)
        # On failure, the per-dimension feedback becomes the next attempt's context.
        return Verdict(passed=result.passed, reason=result.feedback())
```

Because `check` returns the rubric’s per-dimension `feedback()` as the verdict reason, the agent loop’s existing failure path (Chapter 2) re-injects “Rubric failures: scope: touched files outside the request” as a new message — closing the verification loop without any special-case wiring.

8.10 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 8.1

Why treat verification as a distinct architectural layer rather than a self-critique step inside the agent loop?

Hint. It sits outside the loop and judges only the output, with no investment in it.

Exercise 8.2

When should a grader be deterministic versus agentic, and what is the cost of getting it wrong?

Hint. Deterministic where the criterion is codeable; agentic only for judgment.

Exercise 8.3

What does rubric-based scoring add over a single pass/fail signal?

Hint. Per-dimension, actionable feedback instead of a bare pass/fail.

Exercise 8.4

Explain the maker/checker split and why it applies to the stop condition, not just the output.

Hint. A fresh checker — applied to the stop condition, not just the output.

Exercise 8.5

How is feedback injection from the grader related to tool-error handling from Chapter 6?

Hint. Both reformat a raw signal into actionable feedback before reuse.

Exercise 8.6

Verification adds cost and latency. How do you decide which tasks get a full verification loop?

Hint. Compare the cost of being wrong to the cost of checking.

Exercise 8.7

An agent's tests are green but the feature is visibly broken for users. Which failure is this and how does verification design address it?

Hint. Overfitting to the verifier — check user-visible behaviour directly.

Exercise 8.8

Walk through the worked documentation-agent example end to end as a verification loop.

Hint. Deterministic checks: links resolve, CI passes, the diff is scoped.

Exercise 8.9

How do you prevent the verification loop itself from running forever?

Hint. It is still a loop — bound the rounds and escalate.

Exercise 8.10

Production case: an LLM-as-judge grader’s pass rate drifts upward over weeks while real quality stays flat. What is happening and what do you do?

Hint. Grader drift / reward hacking — anchor with gold examples and human calibration.

8.11 Chapter Summary

- A verification loop wraps the entire agent loop from the outside, deciding accept-or-retry based on the agent’s final output for an attempt.
- Prefer deterministic graders wherever the criterion can be expressed as code; reserve agentic, LLM-as-judge grading for genuinely subjective criteria.
- Rubric-based scoring gives the next attempt specific, actionable feedback rather than a bare pass/fail.
- The maker/checker split — a separate grader from the agent that produced the output — is the central structural idea, and applies to the stop condition itself, not only the final output.
- Verification adds real cost and latency; reserve it for work where being wrong is more expensive than checking.

Chapter 9 catalogs verification loops, retry loops, and several other named loop architectures as a single reference you can match against the task in front of you.

8.12 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 8.1

Why treat verification as a distinct architectural layer rather than a self-critique step inside the agent loop?

A self-critique step (Chapter 5) runs inside the agent loop, in the same context that produced the output, and is therefore a generous judge of its own work. A verification loop sits *outside* the agent loop, observes only the final output of an attempt, and decides accept-or-return-with-feedback — structurally separate, with its own retry logic and ideally its own model. The separation matters because it removes the conflict of interest: the checker has no investment in the work and no exposure to the reasoning that rationalised it. It also gives you a clean place to put deterministic graders (tests, builds) that have nothing to do with the agent’s reasoning at all. Self-critique improves an attempt; a verification loop decides whether an attempt is acceptable. They operate at different layers and a serious system has both.

Answer 8.2

When should a grader be deterministic versus agentic, and what is the cost of getting it wrong?

Use a deterministic grader wherever the success criterion can be expressed as code — tests pass, build succeeds, links resolve, the diff is scoped to what was requested. They are fast, cheap, reproducible, and unambiguous. Reserve agentic (LLM-as-judge) grading for criteria that genuinely require judgment — does the writing match the audience, is the architecture sound, does the change respect an unwritten convention. The cost of getting it wrong runs both ways: a grader that could have been a test suite but is an LLM call adds cost, latency, and variance for no reliability gain; conversely, forcing a subjective criterion into a brittle deterministic check produces a grader that passes bad work and fails good work. The default is deterministic-first, agentic-only-where-necessary, often combined in one rubric.

Answer 8.3

What does rubric-based scoring add over a single pass/fail signal?

Actionable, targeted feedback. A bare “failed” forces the next attempt to re-approach the entire task blind; a rubric decomposes success into named dimensions — correctness, scope, style, test coverage — each scored independently, so the agent sees precisely which dimension fell short and can fix that without disturbing the rest. “Scope: you modified three files outside the request” is a directive; “failed” is a guess prompt. In the reference grader, `RubricResult.feedback()` emits only the failed dimensions with their reasons, and that string becomes the next attempt’s context. Rubrics also make the grader auditable — you can see the per-dimension history across attempts — and they let you weight or gate dimensions (e.g., correctness is mandatory, style is advisory). The decomposition is what turns verification from a gate into a teacher.

Answer 8.4

Explain the maker/checker split and why it applies to the stop condition, not just the output.

The maker/checker split means the entity that produced the output is never the entity that grades it — a separate grader, ideally a different model or at minimum a fresh context window. The model that wrote the code is too generous a judge of it; a fresh checker, with no investment in the work and no exposure to the reasoning that produced it, catches what the maker talked itself into believing was correct. It applies with particular force to the *stop condition*: asking the same working context “are you done?” reliably yields an optimistic yes. Having a fresh checker independently evaluate

whether the verifiable condition actually holds is a meaningfully stronger termination signal than self-assessment. This is the single most important structural decision in verification design — and it is why the reference loop verifies via a separate **Verifier** object rather than asking the generating call to self-certify.

Answer 8.5

How is feedback injection from the grader related to tool-error handling from Chapter 6?

They use the same discipline: a raw result is not actionable, so it must be processed before reuse. A raw rubric score (“0.6”) tells the agent nothing; a clear statement of which dimension failed and why is what drives a better next attempt — exactly as a raw stack trace must be turned into structured feedback before re-injection. In effect, the verification loop treats the grader as another tool whose output must be reformatted before it re-enters the agent loop. The reference grader makes its **check** return the per-dimension **feedback()** string as the verdict reason precisely so it flows through the same re-injection path the loop already uses for tool errors. The unifying principle: every signal fed back into reasoning — tool error, grader verdict, human comment — should be structured for action, not dumped raw.

Answer 8.6

Verification adds cost and latency. How do you decide which tasks get a full verification loop?

By comparing the cost of being wrong to the cost of checking. For most production use cases the cost of shipping a wrong or low-quality output — a broken deploy, a bad customer-facing change, downstream remediation — exceeds the cost of an extra verification pass, so verification pays for itself. But it is not free, and a low-stakes, easily-reversible, cheaply-redone task may not justify the added latency and retry cycles. The decision is per-task-class: reserve the dedicated verification layer for work where being wrong is meaningfully more expensive than checking, and skip or lighten it for cheap, reversible work. You can also tier it — a cheap deterministic gate on every task, an expensive agentic grade only on high-stakes ones. The anti-pattern is uniform heavyweight verification on everything, which taxes the cheap-and-reversible majority to protect the expensive minority.

Answer 8.7

An agent’s tests are green but the feature is visibly broken for users. Which failure is this and how does verification design address it?

This is overfitting to the verifier (Chapter 7): the agent satisfied the check while missing the actual requirement, typically because the tests are too narrow or user-visible behaviour is never directly checked. Verification design addresses it by broadening what “pass” means: add a rubric dimension that checks user-visible behaviour directly — a browser-based or end-to-end check, a screenshot diff, an acceptance criterion validated against the running app — rather than trusting the unit-test suite alone. It also helps to include a requirement-review dimension (agentic grading against the original request) so “does this actually do what was asked?” is graded explicitly. The deeper point: a verifier is only as good as its coverage of the real requirement, and a loop will optimise to exactly the signal you give it, so an incomplete verifier produces confidently green, behaviourally broken output.

Answer 8.8

Walk through the worked documentation-agent example end to end as a verification loop.

The agent receives a request, plans and drafts documentation changes, and uses tools to clone the repository, read files, write updated docs, and open a pull request — that is the inner agent loop. Wrapping it, the verification loop runs a grader after each attempt with several deterministic dimensions: every link resolves, CI checks pass, and the diff is scoped to what was actually requested. If any dimension fails, the loop returns the attempt to the agent with the specific failure (“link to

/api/v2 is broken”), and the agent revises. Only when all dimensions pass is the output accepted — catching an entire class of error before a human ever reviews it. The design wins here are that the grader is mostly deterministic (links, CI, scope are all checkable as code), so verification is cheap and reliable, and the scope check directly guards against the agent “helpfully” editing unrelated files.

Answer 8.9

How do you prevent the verification loop itself from running forever?

The verification loop is still a loop, so it needs the control structures of Chapter 3: a bounded number of accept-or-retry rounds, and an escalation path when the budget is exhausted without acceptance. If the agent cannot satisfy the grader within k rounds, that is itself a signal — the task may be mis-scoped, the grader may be too strict or contradictory, or the agent lacks a needed tool — and the loop should escalate to a human with the rubric history rather than retrying indefinitely. You also guard against oscillation, where each attempt fixes one dimension and breaks another; detecting that the failing-dimension set is cycling should trigger a strategy change or escalation. The principle: wrapping a loop in another loop does not exempt the outer loop from needing its own budget, no-progress detection, and terminal states.

Answer 8.10

Production case: an LLM-as-judge grader’s pass rate drifts upward over weeks while real quality stays flat. What is happening and what do you do?

This is grader drift / reward hacking combined with the limits of agentic grading: either the agent has learned (through prompt or process changes) to produce outputs that flatter the judge without being better, or the judge itself is inconsistent and trending lenient. The fixes: anchor the agentic grader with a fixed, version-controlled rubric and a held-out set of gold examples scored periodically to detect judge drift; add or strengthen deterministic dimensions so the pass signal is not purely the judge’s opinion; and calibrate the judge against human labels on a sample, treating a divergence between judge pass-rate and human-rated quality as the alarm. Long term, a single LLM judge is a moving target — pair it with deterministic checks and periodic human calibration so “pass rate” is tied to ground truth, not to the judge’s mood. This is the verification analogue of overfitting to the verifier, one level up.

Loop Pattern Catalog

Subscribe → aiengineeringinsider.substack.com/subscribe

Good engineers do not invent loops. They recognise them.

Experienced engineers rarely invent a loop from nothing. They recognise one. The bug report that needs a failing test before a fix, the migration the compiler can drive, the pull request that loops on reviewer comments — each has a shape that has been solved before, and the skill is matching the task in front of you to the pattern that already fits it. This chapter is the catalogue you match against.

The preceding chapters built up the individual mechanisms — control structures, state, reasoning, tools, verification — that any loop is assembled from. This chapter catalogs nine standard, named loop architectures built from those mechanisms, organized by the kind of feedback signal each one is built around. Treat this as a reference to match against the task in front of you, not a sequence to read in order.

9.1 The Retry Loop

The simplest pattern: the agent attempts an action, checks whether it worked, and retries if it did not.

When to Use

Short, atomic tasks with a clear pass/fail criterion — writing a function that must pass one test, generating output that must meet a simple spec, running a query until it returns valid data.

! Pitfall / Warning

Watch for unmodified retries without a strategy change. As established in Chapter 7, if the same approach keeps failing, the loop needs logic to vary its next attempt, not simply repeat the last one.

9.2 The Plan-Execute-Verify Loop

The agent generates a plan, executes it step by step, and verifies each step before proceeding to the next.

When to Use

Multi-step tasks where order matters and early mistakes compound — refactoring a module, standing up a new service, building a feature with multiple interdependent components.

! Pitfall / Warning

Watch for over-commitment to a bad plan. If an early step reveals the plan was wrong, the agent must revise the plan, as covered in Chapter 5, rather than push through the remaining steps regardless.

9.3 The Explore-Narrow Loop

The agent explores multiple solution paths — sequentially or in parallel — then narrows to the most promising one based on intermediate results, closely related to the tree-of-thought approach from Chapter 5.

When to Use

Debugging an unknown error, exploring an unfamiliar API, or optimizing performance — situations where the right approach is genuinely unknown upfront.

! Pitfall / Warning

Watch for context explosion. Running multiple paths in parallel is expensive in both tokens and tool calls; pruning weak branches early and often is essential, not optional.

9.4 The Test-Driven Agent Loop

The agent reproduces a failure or writes a failing test first, then changes implementation code until that test passes.

When to Use

Bug fixes, parser behavior, data transformation logic, and regression prevention — any task where a test can directly encode the expected behavior.

The cycle: reproduce or encode the expected behavior, confirm the test fails for the right reason, implement the smallest fix, rerun the targeted test, then run broader validation if needed.

9.5 The Compiler-Driven Loop

In typed languages, the compiler or type checker itself becomes the feedback source: the agent makes a change, runs the type checker, and uses the resulting errors directly as a precise repair list.

When to Use

Migrations, dependency upgrades, and refactors in typed codebases. The compiler does not prove the product is correct, but it catches structural mistakes quickly, before they consume review time.

9.6 The Review-Driven Loop

Human feedback becomes the observation source directly: a reviewer leaves comments, the agent categorizes them, makes changes where appropriate, responds to non-code feedback, and verifies the result.

When to Use

Any workflow where mechanical follow-through on review comments can be automated while judgment stays with a human.

✓ Best Practice

Treat review comments as requirements to be triaged, not suggestions to be blindly applied. The agent should distinguish bugs, product questions, style preferences, and out-of-scope requests before editing anything.

9.7 The Runtime Debugging Loop

The agent uses logs, stack traces, HTTP responses, browser output, or reproduction steps to narrow the cause of a failure that only appears when the application actually runs.

When to Use

Failures that cannot be reproduced statically — frontend hydration errors and layout bugs, backend request traces, database constraint violations, queue retry storms.

The key discipline: keep every iteration evidence-based. Form a hypothesis, make a targeted change or add a targeted inspection point, observe the actual result, and update the hypothesis — rather than pattern-matching to a plausible-sounding cause.

9.8 The Product Iteration Loop

Not every loop is purely technical. The agent implements a version of copy, a UI state, or an edge case, compares it against design constraints, adjusts, and re-verifies — including semantic HTML, metadata, accessibility, and responsive behavior for frontend work.

When to Use

Iterative product and design work: copy refinement, UI state handling, acceptance-criteria validation against design constraints.

9.9 The Human-in-the-Loop Pattern

Technically a loop variant rather than a separate architecture: the agent runs until it needs clarification or hits genuine ambiguity, pauses for human input, then continues with that input incorporated.

When to Use

Tasks with requirements that cannot be fully specified upfront; production changes where a human should review before execution; any task where the cost of a wrong assumption is high.

! Pitfall / Warning

Watch for interrupting too often. If the agent asks for input on every small decision, it is not actually saving the human any time — it has just moved the work from doing the task to babysitting the loop.

9.10 Choosing a Pattern

Feedback Signal	Pattern
Pass/fail test or check	Retry Loop, Test-Driven Loop
Type checker / compiler output	Compiler-Driven Loop
Human review comments	Review-Driven Loop
Runtime behavior (logs, traces)	Runtime Debugging Loop
Unknown solution space	Explore-Narrow Loop
Multi-step, order-dependent task	Plan-Execute-Verify Loop
Design or product constraints	Product Iteration Loop
Genuine ambiguity requiring judgment	Human-in-the-Loop

Table 9.1: Matching a loop pattern to the dominant feedback signal available for the task.

From the Field

These patterns are rarely used in isolation in a real system. A typical production agent loop for a bug fix is a Runtime Debugging Loop to locate the cause, feeding into a Test-Driven Loop to encode and fix it, wrapped in a Review-Driven Loop for the pull request, with a Human-in-the-Loop pause before anything touches production. Treat the patterns in this chapter as composable building blocks, not mutually exclusive choices.

9.11 Patterns Compose: A Real Bug-Fix Pipeline

The patterns are building blocks, not mutually exclusive choices. A production bug-fix workflow chains several:




Figure 9.1: Patterns compose: a Runtime Debugging loop locates the cause, a Test-Driven loop encodes and fixes it, a Review-Driven loop handles the PR, and a Human-in-the-Loop gate guards production.

9.12 Reference Implementation: Patterns Are Configuration

The named patterns are not nine different engines — they are the same `AgentLoop` wired to a different feedback signal, verifier, and budget. The `loop_engineering_lab/patterns/catalog.py` builders make that explicit.

`loop_engineering_lab/patterns/catalog.py` — one engine, many patterns

```
def retry_loop(llm, tools, *, success_marker, max_iters=5) -> AgentLoop:
    """Attempt, check a pass/fail marker, retry on failure."""
    policy = TerminationPolicy(budget=Budget(max_iterations=max_iters))
    return AgentLoop(llm, tools, policy, verifiers=[KeywordVerifier(success_marker)])

def test_driven_loop(llm, tools, verifier, *, max_iters=8) -> AgentLoop:
    """Failure-encoded-as-test: loop until the supplied verifier passes."""
    policy = TerminationPolicy(budget=Budget(max_iterations=max_iters,
    ↪ max_tool_calls=24))
    return AgentLoop(llm, tools, policy, verifiers=[verifier])

def plan_execute_verify_loop(llm, tools, verifier, *, max_iters=12) -> AgentLoop:
    """Order-dependent work: larger budget, no-progress detection, gated steps."""
    policy = TerminationPolicy(
        budget=Budget(max_iterations=max_iters, max_tool_calls=40),
        no_progress_limit=3)
    return AgentLoop(llm, tools, policy, verifiers=[verifier])
```

The point is architectural economy: once the core loop, control policy, and verifier abstractions exist, a “new pattern” is usually a new *configuration*, not new code. That is why the patterns share the building blocks of Chapters 2–8 rather than each being a bespoke system.

9.13 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 9.1

What organising principle defines these patterns — their internals or something else?

Hint. The dominant feedback signal, not the internal mechanics.

Exercise 9.2

Show how the same core loop expresses both a retry loop and a plan-execute-verify loop.

Hint. Same engine, different budget and verifier.

Exercise 9.3

When is the Test-Driven Agent Loop the right pattern, and what discipline makes it work?

Hint. Encode behaviour as a test, and confirm it fails for the right reason.

Exercise 9.4

Compare the Compiler-Driven Loop and the Runtime Debugging Loop. What signal does each exploit?

Hint. Static structural signal versus dynamic runtime signal.

Exercise 9.5

The Review-Driven Loop turns human comments into observations. What is the critical discipline?

Hint. Triage comments as requirements before editing anything.

Exercise 9.6

What goes wrong with the Human-in-the-Loop pattern, and how do you tune it?

Hint. Interrupting too often — pause only on genuine ambiguity.

Exercise 9.7

How do these patterns compose in a real production workflow? Give a concrete chain.

Hint. Debugging → test-driven fix → review-driven PR → human gate.

Exercise 9.8

The Explore-Narrow Loop is powerful but expensive. How do you control its cost?

Hint. Prune weak branches early and often; bound the branching factor.

Exercise 9.9

How would you pick a starting pattern for a task you have never automated before?

Hint. Start from the cheapest trustworthy signal and expect to compose.

Exercise 9.10

Production case: a team built one big monolithic agent prompt for "fix the bug, write tests, open a PR, and deploy." It is unreliable. How do you re-architect it using these patterns?

Hint. Decompose by feedback signal, each stage with its own verifier.

9.14 Chapter Summary

- Nine named patterns — Retry, Plan-Execute-Verify, Explore-Narrow, Test-Driven, Compiler-Driven, Review-Driven, Runtime Debugging, Product Iteration, and Human-in-the-Loop — cover the great majority of production agentic workflows.
- Each pattern is defined by its dominant feedback signal source, not by its internal mechanics, which mostly reuse the building blocks from Chapters 2 through 8.
- Real systems typically compose several patterns in sequence rather than choosing exactly one.

Part III is complete. Part IV moves beyond a single agent's loop to loops that coordinate across multiple agents, run without a human present, and improve themselves over time.

9.15 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 9.1

What organising principle defines these patterns — their internals or something else?

Their dominant feedback signal, not their internal mechanics. Almost every pattern reuses the same building blocks — a control loop, state, reasoning, tools, verification — and differs only in *what observation it is built around*: a pass/fail test (Retry, Test-Driven), compiler output (Compiler-Driven), human comments (Review-Driven), runtime behaviour (Runtime Debugging), an unknown solution space (Explore-Narrow), order-dependent steps (Plan-Execute-Verify), design constraints (Product Iteration), or genuine ambiguity (Human-in-the-Loop). Recognising that the feedback signal is the organising axis is what lets you pick a pattern by asking “what is the cheapest, most trustworthy signal available for this task?” rather than memorising nine recipes. It also explains why they compose so freely — they are variations on one engine, not incompatible architectures.

Answer 9.2

Show how the same core loop expresses both a retry loop and a plan-execute-verify loop.

Both are an **AgentLoop** with a verifier and a termination policy; they differ in budget and signal, not in engine. A retry loop is a small iteration budget and a simple pass/fail verifier — attempt, check the marker, retry. A plan-execute-verify loop is a larger budget, a higher tool-call ceiling, no-progress detection, and a verifier that gates each step — because the task is multi-step and order-dependent. In the reference project these are literally two builder functions returning the same **AgentLoop** class with different **TerminationPolicy** and verifier arguments. The lesson for system design is economy: invest in clean loop/control/verifier abstractions once, and most “new patterns” become configuration. If adding a pattern requires a new engine, the abstractions are probably wrong.

Answer 9.3

When is the Test-Driven Agent Loop the right pattern, and what discipline makes it work?

It fits any task where expected behaviour can be encoded as a test: bug fixes, parser behaviour, data-transformation logic, regression prevention. The discipline that makes it work is the order: reproduce or encode the expected behaviour as a test *first*, confirm the test fails for the right reason (so you know it actually exercises the bug), implement the smallest fix, rerun the targeted test, then run broader validation. Skipping “confirm it fails for the right reason” is the classic mistake — a test that passes for the wrong reason gives false confidence, and a test that never failed proves nothing. The pattern’s power is that it converts a vague “fix this” into a concrete, machine-checkable target, which is exactly the dense verification signal a loop needs to converge reliably.

Answer 9.4

Compare the Compiler-Driven Loop and the Runtime Debugging Loop. What signal does each exploit?

The Compiler-Driven Loop exploits static structural feedback: in a typed language, the compiler or type checker produces a precise list of errors after each change, and the agent uses that list directly as a repair queue. It is fast and excellent for migrations, dependency upgrades, and refactors, but it only proves structural correctness, not that the product behaves correctly. The Runtime Debugging Loop exploits dynamic behavioural feedback: logs, stack traces, HTTP responses, browser output — evidence that only appears when the application actually runs. It is the right pattern for failures that cannot be reproduced statically (hydration errors, request traces, constraint violations). The key difference is static versus dynamic signal, and the discipline for runtime debugging is to stay evidence-based — hypothesis, targeted change or inspection, observe, update — rather than pattern-matching

to a plausible cause.

Answer 9.5

The Review-Driven Loop turns human comments into observations. What is the critical discipline?

Treat review comments as requirements to be triaged, not suggestions to be blindly applied. The agent should categorise each comment before editing — bug, product question, style preference, out-of-scope request — because they demand different responses: a bug gets fixed, a product question gets answered or escalated rather than guessed, a style preference gets applied if it matches conventions, an out-of-scope request gets deferred. Blindly applying every comment produces scope creep and can introduce errors when a reviewer’s suggestion is itself wrong or based on a misunderstanding. The pattern automates the mechanical follow-through on review while keeping judgment with the human, and the triage step is what keeps it from becoming an obedient but thoughtless change-applier. Done well, it makes the agent behave like a competent contributor responding to review, not a script applying diffs.

Answer 9.6

What goes wrong with the Human-in-the-Loop pattern, and how do you tune it?

The failure mode is interrupting too often. If the agent pauses for input on every small decision, it has not saved the human any work — it has converted “do the task” into “babysit the loop,” which is often worse than doing it manually. The pattern is valuable when requirements cannot be fully specified upfront, when a human must review before a production change, or when the cost of a wrong assumption is high — but the tuning is calibrating *when* to pause. Good practice: let the agent proceed autonomously on reversible, low-stakes decisions and pause only on genuine ambiguity, irreversible actions, or judgment calls it has no basis to make. The destructive-action and blocker-escalation guards from Chapter 3 are the concrete mechanism. The goal is to spend the human’s attention on the few decisions that actually need it, not to seek constant reassurance.

Answer 9.7

How do these patterns compose in a real production workflow? Give a concrete chain.

They are rarely used in isolation. A production bug-fix pipeline typically chains four: a Runtime Debugging Loop to locate the cause from logs and traces; feeding into a Test-Driven Loop to encode the failure as a test and fix it; wrapped in a Review-Driven Loop that opens a pull request and responds to reviewer comments; with a Human-in-the-Loop pause before anything touches production. Each stage’s output is the next stage’s input, and each uses the feedback signal appropriate to its job. Treating the patterns as composable building blocks rather than mutually exclusive choices is the mark of a mature design — the question is never “which one pattern?” but “which sequence of patterns matches the stages of this task, and where are the hand-offs?” The composition diagram in this chapter shows exactly this chain.

Answer 9.8

The Explore-Narrow Loop is powerful but expensive. How do you control its cost?

Explore-Narrow runs multiple solution paths — sequentially or in parallel — then narrows to the most promising based on intermediate results, which is exactly the tree-of-thought trade from Chapter 5 at the architecture level. The cost control is aggressive, early pruning: evaluate branches on cheap intermediate signals and kill weak ones before they consume a full budget, rather than running every branch to completion. Bound the branching factor (two or three candidates, not ten), and reserve the pattern for points where the right approach is genuinely unknown — debugging an unfamiliar error, exploring an undocumented API, optimising performance — not for routine decisions with an obvious path. Running many paths to the end “to be thorough” is the failure mode; the value is

in branching *and pruning*, where the pruning is what keeps the token and tool-call cost bounded. Without disciplined pruning, the pattern is just expensive.

Answer 9.9

How would you pick a starting pattern for a task you have never automated before?

I start from the feedback signal, using the selection table as a map: what is the cheapest, most trustworthy observation available for this task? If a pass/fail test exists or can be written, I start Test-Driven; if the language is typed and the task is structural, Compiler-Driven; if the failure only appears at runtime, Runtime Debugging; if the solution space is unknown, Explore-Narrow; if order matters and steps compound, Plan-Execute-Verify; if judgment is unavoidable, Human-in-the-Loop. Then I expect to *compose* rather than commit to one, and I start with the simplest pattern that has a real verification signal, adding structure only where the task proves it needs it. The mistake is reaching for an elaborate pattern when a Retry Loop with a good verifier would do, or reaching for a Retry Loop on a task whose real signal is human judgment. Signal first, structure second.

Answer 9.10

Production case: a team built one big monolithic agent prompt for "fix the bug, write tests, open a PR, and deploy." It is unreliable. How do you re-architect it using these patterns?

The monolith fails because it collapses four different feedback signals into one undifferentiated loop with no clear verifier at any stage, so errors at one stage silently propagate to the next — the chain-not-a-loop failure from Chapter 1, scaled up. I would decompose it along the feedback signals into composed patterns, each with its own verifier and budget: a Runtime Debugging stage that must produce a confirmed root cause before proceeding; a Test-Driven stage gated on "a test reproduces the bug and now passes"; a Review-Driven stage gated on CI green and scope checks; and a Human-in-the-Loop gate before deploy. Each stage has an explicit terminal condition, so a failure stops and escalates at the right place instead of corrupting downstream work. In the reference vocabulary, that is replacing one over-broad **AgentLoop** with a pipeline of pattern-configured loops, each wired to the signal appropriate to its job. The reliability gain comes from giving each stage a real, checkable definition of done.

PART IV

Scaling Beyond One Agent

One loop has limits. Part IV crosses them three ways: outward to many agents coordinating on one goal, sideways to loops triggered by events rather than by a human pressing go, and upward to the loop whose job is to improve every other loop in the system. This is where loop *stacking* — composing simple loops into layered systems — becomes the organising idea.

Multi-Agent Loop Architectures

Subscribe → aiengineeringinsider.substack.com/subscribe

Ten agents is a coordination problem wearing a reasoning problem's clothes.

The moment a single agent stops being enough, the temptation is to spin up ten. It is a trap, and a seductive one — because the hard part of running ten agents is not getting them to think. It is getting them to not overwrite each other's work, not read each other's half-finished state, and not collectively produce more output than any human can possibly review. Multi-agent systems trade a reasoning problem you understand for a coordination problem you probably do not, and the teams that succeed are the ones that respect that trade before they make it.

Some tasks exceed what a single agent loop can reliably handle within one context window and one line of reasoning. Multi-agent loop architectures coordinate several agents, each running its own loop, to handle far more complexity than any single loop could manage alone.

10.1 Orchestrator-Worker Loops

A controller loop dispatches sub-tasks to one or more worker loops, collects their results, and decides the next round of dispatch. The orchestrator holds the overall plan and progress; workers hold no context beyond their assigned sub-task.

★ Key Insight

The orchestrator-worker split is most valuable specifically because it keeps each worker's context narrow and focused. A worker that only ever sees its one sub-task and the tools it needs for that sub-task reasons more reliably than one large agent juggling the entire task's full context at once.

10.2 Peer-to-Peer Agent Loops

Rather than a strict hierarchy, peer agents interact directly — one agent proposes, another critiques, a third revises — in a debate-style or critique-style exchange with no single controlling orchestrator.

This pattern is closely related to the maker/checker split from Chapter 8, generalized to more than two participants, and is most useful when the value of multiple independent perspectives outweighs the coordination overhead of managing a fixed hierarchy.

10.3 Hierarchical Loop Decomposition

A large goal is broken into sub-goals, each of which gets its own loop — with its own termination condition, its own tool budget, and potentially its own further decomposition into sub-sub-goals.




Figure 10.1: Hierarchical decomposition: a top-level goal splits into sub-goals, each handed to its own loop with its own termination condition and budget.

10.4 The Planner / Parallel-Executors / Reviewer Pattern

A common, specific instance of hierarchical decomposition worth naming directly: one planning agent breaks a large task into independent subtasks, multiple executor agents each handle one subtask in parallel, and a reviewer agent checks each output and routes failures back for correction.

This is multi-agent loop engineering proper: the loops are nested and coordinated, and the system handles far more complexity than a single agent loop — particularly valuable for real-world tasks spanning dozens of files, services, or APIs.

10.5 Loop Synchronization

When multiple agents run concurrently, the loop needs explicit synchronization points: barriers where all parallel workers must complete before the orchestrator proceeds, and handoff protocols defining exactly what state passes from one agent to the next.

! Pitfall / Warning

Implicit synchronization — assuming agents will naturally avoid stepping on each other — does not hold in practice. Without explicit barriers and handoff contracts, concurrent agents will read stale state from one another or overwrite each other's in-progress work.

10.6 Shared State vs. Isolated State

Shared state

All agents read and write to the same state store (the markdown file or ticket board from Chapter 4). Simple, but creates real collision risk under concurrency.

Isolated state

Each agent gets its own scoped state, merged back into shared state only at defined synchronization points. More overhead to design, but eliminates the collision risk entirely between synchronization points.

10.7 Worktree-Based Isolation

Git Worktree

A separate working directory on its own branch, sharing the same underlying repository history. Each agent operating in its own worktree has a checkout that no other agent's edits can touch.

! Pitfall / Warning

Two agents writing to the same file is the exact same problem as two engineers committing to the same lines with no coordination. A git worktree per parallel agent fixes the mechanical collision entirely — but it does not remove the need for a human to review and merge the resulting work. Worktrees solve the file-collision problem; they do not solve the review-bandwidth problem covered in Chapter 14.

From the Field

The moment you run more than one agent against the same codebase, file collision becomes the dominant failure mode, well before any reasoning failure shows up. Solve isolation first. A perfectly-reasoning multi-agent system that lets two workers clobber the same file is still a broken system.

10.8 Coordination Topologies Compared

Topology	Strength	Cost / risk
Single agent	Simplest; one context, one budget	Hits a complexity ceiling
Orchestrator-worker	Narrow worker context, parallelism	Orchestrator is a bottleneck / single point of failure
Peer-to-peer (debate)	Multiple independent perspectives	Coordination overhead; may not converge
Hierarchical	Scales to deep task trees	Deep nesting is hard to trace and budget

Table 10.1: More agents buy capacity and independence at the cost of coordination, isolation, and traceability overhead — add them only when one loop genuinely cannot cope.

10.9 Reference Implementation: Orchestrator with Isolated Workers

The orchestrator's power comes from building each worker *fresh*, so a worker only ever sees its own sub-task — no shared context window, no cross-talk. `loop_engineering_lab/multiagent/orchestrator.py` makes the isolation explicit via a worker factory and an all-succeeded barrier.

loop_engineering_lab/multiagent/orchestrator.py — isolation by construction

```
class Orchestrator:
    def __init__(self, build_worker: Callable[[], AgentLoop]) -> None:
        self._build_worker = build_worker      # a FACTORY, not a shared instance
```

```

def run(self, subtasks: list[str]) -> list[SubtaskResult]:
    results = []
    for subtask in subtasks:
        worker = self._build_worker()      # fresh worker => isolated context
        results.append(SubtaskResult(subtask, worker.run(subtask)))
    return results

    @staticmethod
    def all_succeeded(results) -> bool:
        # A synchronization barrier: proceed only when every worker finished ok.
        return all(r.succeeded for r in results)

```

Passing a *factory* rather than a single shared loop is the whole trick: it guarantees each sub-task gets its own `LoopState`, its own budget, and its own termination logic, which is exactly what keeps one worker's context from polluting another's.

10.10 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 10.1

Why does the orchestrator–worker split improve reliability over one large agent?

Hint. Narrow worker context reasons more reliably than one large agent.

Exercise 10.2

When is peer-to-peer (debate/critique) coordination worth its overhead over a hierarchy?

Hint. When independent perspectives catch errors a single line of reasoning misses.

Exercise 10.3

What does hierarchical decomposition give each sub-goal, and what is the cost of deep nesting?

Hint. Its own budget and termination; deep nesting is hard to trace.

Exercise 10.4

Describe the planner / parallel-executors / reviewer pattern and where it fits.

Hint. Planner, parallel executors, reviewer — for independent sub-tasks.

Exercise 10.5

Why does implicit synchronization fail, and what explicit mechanisms replace it?

Hint. Agents clobber each other — use barriers and hand-off contracts.

Exercise 10.6

Compare shared state and isolated state for multi-agent loops. How do you choose?

Hint. Isolation with controlled merge versus shared mutable state.

Exercise 10.7

What problem do git worktrees solve for parallel agents, and what do they explicitly not solve?

Hint. They fix file collision, not the review-bandwidth bottleneck.

Exercise 10.8

Production case: a 5-agent parallel refactor keeps producing merge conflicts and lost edits. Diagnose and fix.

Hint. Isolation was solved last — give each agent its own worktree.

Exercise 10.9

How do budgets and termination compose across a multi-agent hierarchy?

Hint. Per-agent and global budgets that roll up the hierarchy.

Exercise 10.10

When should you *not* reach for a multi-agent architecture?

Hint. When one loop fits the task — avoid premature multi-agent design.

10.11 Chapter Summary

- Orchestrator-worker loops keep individual worker context narrow, which improves per-worker reliability.
- Peer-to-peer loops generalize the maker/checker split to more than two participants, trading coordination overhead for multiple independent perspectives.
- Hierarchical decomposition gives each sub-goal its own loop, termination condition, and budget.
- The planner / parallel-executors / reviewer pattern is the most common concrete instance of multi-agent loop engineering for real software tasks.
- Concurrent agents need explicit synchronization barriers and handoff protocols; implicit coordination does not hold in practice.
- Worktree-based isolation eliminates file-collision risk between parallel agents, but does not eliminate the need for human review of the merged result.

10.12 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 10.1

Why does the orchestrator-worker split improve reliability over one large agent?

Because it keeps each worker's context narrow and focused. A worker that only ever sees its single sub-task and the tools required for it reasons more reliably than one large agent juggling the entire task's context, plan, and tool surface at once — the model attends better to a small, relevant context than to a sprawling one (the same “lost in the middle” pressure from Chapter 4, at the architecture level). The orchestrator holds the overall plan and progress; workers hold nothing beyond their assignment. The reliability gain is real but it is not free: the orchestrator becomes a coordination bottleneck and a single point of failure, and you pay context-passing overhead at each hand-off. You adopt the split when a single loop's context genuinely cannot hold the whole task reliably, not reflexively.

Answer 10.2

When is peer-to-peer (debate/critique) coordination worth its overhead over a hierarchy?

When the value of multiple independent perspectives outweighs the cost of managing the exchange. Peer agents — one proposes, another critiques, a third revises — generalise the maker/checker split of Chapter 8 to more than two participants, and they shine on tasks where a single line of reasoning is prone to a blind spot a fresh perspective catches: design decisions, ambiguous requirements, adversarial review. The overhead is coordination and convergence: debate can fail to terminate, or amplify a confident wrong consensus if all agents share the same bias. So it is worth it for high-stakes, judgment-heavy decisions where independent critique materially reduces error, and wasteful for mechanical, well-specified work where a hierarchy with a single verifier is cheaper and just as reliable. The deciding question is whether the task's errors are the kind that independent perspectives actually catch.

Answer 10.3

What does hierarchical decomposition give each sub-goal, and what is the cost of deep nesting?

Each sub-goal gets its own loop — its own termination condition, its own tool budget, and potentially its own further decomposition. That isolation lets a large goal be attacked as a tree of independently-bounded, independently-verifiable units, which is how multi-agent systems handle tasks spanning dozens of files or services. The cost is that deep nesting becomes hard to trace and budget: a failure three levels down must surface meaningfully to the top, budgets must compose so a deep subtree cannot silently consume the whole run, and the trace of “what happened” fans out across many loops. The discipline is to keep the tree shallow where possible, give every level explicit budgets and terminal states (Chapter 3), and ensure failures and progress bubble up rather than getting buried. Decomposition is powerful precisely because each node is bounded; lose that and it becomes an untraceable sprawl.

Answer 10.4

Describe the planner / parallel-executors / reviewer pattern and where it fits.

It is the most common concrete instance of multi-agent loop engineering: one planning agent breaks a large task into independent sub-tasks, multiple executor agents each handle one sub-task in parallel, and a reviewer agent checks each output and routes failures back for correction. It fits real software tasks that span many files, services, or APIs and decompose into genuinely independent units — the parallelism is only safe when the sub-tasks do not contend for the same resources. The pattern combines three ideas from earlier chapters: orchestration (planner), isolation (parallel executors with

their own context and, ideally, their own worktree), and the maker/checker split (reviewer). Its failure mode is mis-decomposition — if the planner carves out sub-tasks that are not actually independent, the parallel executors collide, which is why isolation and synchronization are prerequisites, not afterthoughts.

Answer 10.5

Why does implicit synchronization fail, and what explicit mechanisms replace it?

Implicit synchronization — assuming concurrent agents will naturally avoid stepping on each other — fails because they will not: without barriers and contracts, concurrent agents read stale state from one another or overwrite each other's in-progress work, exactly as uncoordinated engineers committing to the same lines would. The explicit mechanisms are synchronization barriers (points where all parallel workers must complete before the orchestrator proceeds) and hand-off protocols (an explicit contract defining precisely what state passes from one agent to the next). In the reference orchestrator, `all_succeeded` is a simple barrier: the orchestrator only proceeds once every worker has finished. At larger scale you add resource locks, isolated state merged only at barriers, and per-agent worktrees. The principle: concurrency correctness must be designed in, never assumed.

Answer 10.6

Compare shared state and isolated state for multi-agent loops. How do you choose?

With shared state, all agents read and write the same store (the markdown file or ticket board of Chapter 4) — simple to set up, but it creates real collision risk under concurrency. With isolated state, each agent gets its own scoped state, merged back into shared state only at defined synchronization points — more design overhead, but it eliminates collisions between barriers entirely. You choose based on concurrency and contention: if agents run sequentially or never touch the same data, shared state is fine and simpler; if they run in parallel against overlapping resources, isolated state with explicit merge points is the safe default. The general rule mirrors database design — prefer isolation with controlled merge over shared mutable state whenever genuine concurrency exists, because debugging a race between two agents after the fact is far more expensive than designing the isolation up front.

Answer 10.7

What problem do git worktrees solve for parallel agents, and what do they explicitly not solve?

A git worktree is a separate working directory on its own branch sharing the same repository history, so each agent has a checkout no other agent's edits can touch. It solves the mechanical file-collision problem completely: two agents writing the same file is the same hazard as two engineers committing to the same lines with no coordination, and a worktree per parallel agent removes that hazard entirely. What it explicitly does not solve is the review-bandwidth problem (Chapter 14): the resulting branches still have to be reviewed and merged by a human, and worktrees do nothing to reduce that load — if anything, more parallel agents produce more branches to review. So worktrees are necessary infrastructure for safe parallelism but not sufficient for a healthy system; you still need a merge and review strategy. Solve isolation first, because file collision becomes the dominant failure mode the moment you run more than one agent against a codebase, well before any reasoning failure appears.

Answer 10.8

Production case: a 5-agent parallel refactor keeps producing merge conflicts and lost edits. Diagnose and fix.

This is the classic “isolation solved last” failure: the agents are sharing a working tree (or shared state) with no isolation, so they overwrite each other's in-progress edits and collide at merge. The

diagnosis is that file collision, not reasoning, is the dominant failure mode here — and it shows up before any model-quality issue. The fix is structural, in order: give each agent its own git worktree on its own branch so no two agents touch the same checkout; have the planner decompose along clean boundaries so sub-tasks do not target the same files (mis-decomposition is the second-order cause); add explicit synchronization barriers so the orchestrator merges only completed, isolated work; and route the merged result through human review. The reasoning quality of the agents is irrelevant until isolation is fixed — a perfectly-reasoning multi-agent system that lets two workers clobber the same file is still broken.

Answer 10.9

How do budgets and termination compose across a multi-agent hierarchy?

Each level needs its own budget and terminal states, and they must compose so the whole system has a bounded total cost. A worker has a per-sub-task budget; the orchestrator has a global budget over the whole goal; and a deep subtree's budget must roll up so that one runaway branch cannot consume the entire run. Termination composes too: a worker reaching **BLOCKED** or **FAILED** must surface that to the orchestrator, which decides whether to retry the sub-task, re-plan, escalate, or fail the whole goal — rather than the failure vanishing at the leaf. The anti-pattern is a single global budget with no per-agent limits (one stuck worker starves everyone) or per-agent budgets with no global cap (the total cost is unbounded). The Chapter 3 nested-budget discipline is exactly this, applied across agents instead of across nested loops in one agent.

Answer 10.10

When should you not reach for a multi-agent architecture?

When a single agent loop can handle the task reliably within one context window — which is more often than multi-agent enthusiasm suggests. Multi-agent systems buy capacity and independent perspective at a real cost: coordination overhead, isolation infrastructure (worktrees, state merging, barriers), harder tracing across many loops, and more failure surface. If the task fits one context, one budget, and one line of reasoning, a single loop is simpler, cheaper, easier to debug, and usually more reliable. The honest test is whether you have actually hit a single-agent ceiling — context overflow, genuinely parallelisable independent sub-tasks, or a need for independent critique — or whether you are adding agents because it sounds sophisticated. Start with one loop, prove it is insufficient, and add exactly the coordination structure the evidence demands. Premature multi-agent architecture is a common and expensive form of over-engineering.

Event-Driven and Autonomous Loops

Subscribe → aiengineeringinsider.substack.com/subscribe

Remove the human who presses go, and you remove the safety net nobody noticed.

Up to now, a human has always pressed go. Someone opened a session, typed a task, and watched. This chapter removes that person — and with them, the safety net nobody realised they were. An interactive loop that fails does so in front of you. A loop triggered by a webhook at 4 a.m. can fail in total silence for days, and the most dangerous word in this chapter is not *autonomous*. It is *unnoticed*.

Every loop covered so far is invoked: a human or a calling system explicitly starts it. This chapter covers loops that instead run as a persistent, continuously-available component of a larger system — triggered by events rather than by a person sitting down to start a session.

11.1 The Agent as a Persistent Background Component

An event-driven loop connects the agent to its surrounding ecosystem. An event fires — a new document lands, a schedule triggers, a webhook arrives — and the loop runs. The agent is no longer something invoked manually; it is a component running continuously inside a larger system.




Figure 11.1: An event-driven loop: an external event triggers the agent loop, whose output lands back in the larger system — no human present at the moment it runs.

11.2 Trigger Types

- **Webhooks** — an external system pushes an event the moment it occurs (a pull request opened, a support ticket filed).
- **Cron / schedule triggers** — the loop runs on a fixed cadence, regardless of whether anything specific has happened.
- **Message-queue or channel triggers** — a message in a specific channel (a Slack channel, a queue topic) starts a run.

11.3 The Heartbeat Pattern

Heartbeat

A periodic wake-up, typically on a cron schedule, that turns an otherwise passive agent into an always-on, proactive system — checking for new work even when no external event has explicitly triggered it.

A heartbeat-driven loop might run every morning, read the prior day’s CI failures, open issues, and recent commits, and decide for itself whether anything is worth acting on — archiving the run silently if nothing is found, and surfacing findings to a human inbox if something is.

11.4 Time-Based vs. Condition-Based Looping

Two genuinely different trigger semantics are worth distinguishing precisely, since they are easy to conflate:

Time-based looping

The loop re-runs on a fixed cadence, regardless of state — every morning, every hour. There is no notion of “done”; it simply runs again next cycle.

Condition-based looping

The loop keeps running, across as many turns as needed, until a specific verifiable condition becomes true — for example, until all tests in a given suite pass and the linter is clean. After every turn, a separate check evaluates whether the condition holds, consistent with the maker/checker split on the stop condition established in Chapter 8.

★ Key Insight

Condition-based looping is structurally a verification loop (Chapter 8) running unattended: the same agent that did the work is not the one that decides it is finished. A fresh, separate check evaluates the stop condition after every turn, which is precisely what makes it safe to let the loop run without a human watching each iteration.

11.5 Decoupling Invocation from Interaction

The defining shift in this chapter is architectural: the loop’s trigger and the loop’s output land in different places, and neither requires a human to be present at the moment the loop runs. Findings land in a triage inbox, a ticket board, or a pull request, to be reviewed on the human’s own schedule — not synchronously, the way an interactive session works.

11.6 Infrastructure Concerns

Three concerns become load-bearing once a loop runs unattended on a trigger, in a way they are not for an interactively-invoked loop:

- **Trigger reliability** — a missed webhook or a failed cron run should not silently disappear; it needs its own monitoring.

- **Idempotent event handling** — the same event firing twice (a common occurrence with webhook delivery) must not cause the loop to take a duplicate action, returning to the idempotency discussion from Chapter 6.
- **Concurrency** — multiple events firing close together may trigger overlapping loop runs against the same resource, requiring the same synchronization discipline as the multi-agent case in Chapter 10.

! Pitfall / Warning

A scheduled loop that silently fails to run is far more dangerous than one that fails loudly. An interactively-invoked loop's failure is immediately visible to the person who invoked it. An unattended, triggered loop's failure can go unnoticed for days unless its own execution is independently monitored.

Project Code — `loop_engineering_lab`

`loop_engineering_lab/automation/heartbeat.py` implements a minimal cron-triggered heartbeat loop, with idempotent event handling and a triage-inbox output pattern, built on the verification loop from Chapter 8.

11.7 Trigger Semantics Compared

Trigger	Fires when	Notion of “done”
Webhook	An external event is pushed (PR opened, ticket filed)	One run per event
Cron / schedule	A fixed cadence elapses	None — runs again next cycle
Message / queue	A message lands on a channel/topic	One run per message
Condition-based	Re-runs until a verifiable condition holds	Stops when the check passes

Table 11.1: Time-based triggers have no notion of done; condition-based looping is a verification loop (Ch. 8) running unattended.

11.8 Reference Implementation: An Idempotent Heartbeat

The two infrastructure concerns that become load-bearing the moment a loop runs unattended — idempotent event handling and a durable processed-ledger — are visible in `loop_engineering_lab/automation/heartbeat.py`.

`loop_engineering_lab/automation/heartbeat.py` — exactly-once + triage inbox

```
class Heartbeat:
    def deliver(self, event: Event) -> Finding | None:
        # Idempotent handling: a redelivered event is a no-op (common with
        # at-least-once webhook delivery).
        if event.id in self._seen:
            self.skipped_duplicates += 1
            return None
```

```
self._seen.add(event.id)
self.processed += 1
finding = self._handler(event)
if finding is not None:
    self.inbox.append(finding)    # surfaced for asynchronous human triage
return finding                  # None -> archived silently
```

The handler returning `None` is the “archive silently” path; a `Finding` is the “surface to the human inbox” path. That asymmetry is what lets an always-on loop wake on every cron tick yet only interrupt a human when it actually found something worth their attention.

11.9 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 11.1

What changes architecturally when a loop becomes a persistent background component rather than an invoked tool?

Hint. Trigger and output decouple from any human’s presence.

Exercise 11.2

Distinguish time-based from condition-based looping. Why is the distinction important?

Hint. Fixed cadence versus run-until-a-condition-holds.

Exercise 11.3

What is the heartbeat pattern and what does it enable?

Hint. A periodic wake-up that makes a passive agent proactive.

Exercise 11.4

Why is idempotent event handling essential for event-driven loops, and how do you implement it?

Hint. At-least-once delivery — dedupe by a durable event-id ledger.

Exercise 11.5

Why is a silently failing scheduled loop more dangerous than a failing interactive one?

Hint. Its failure is invisible until something downstream breaks.

Exercise 11.6

How does condition-based looping relate to the maker/checker split from Chapter 8?

Hint. It is the maker/checker split running unattended.

Exercise 11.7

Two webhook events for the same resource fire milliseconds apart. What can go wrong and how do you handle it?

Hint. Overlapping runs corrupt shared state — lock or serialise per resource.

Exercise 11.8

What does “decoupling invocation from interaction” mean, and why does it matter for UX and safety?

Hint. Findings land in an inbox for asynchronous, reviewable handling.

Exercise 11.9

Production case: a heartbeat agent that triages CI failures starts flooding the team’s inbox with low-value findings. How do you fix it without losing the real signal?

Hint. Raise the surface threshold without losing genuine signal.

Exercise 11.10

What are the minimum infrastructure guarantees you would require before putting an unattended loop into production?

Hint. Trigger reliability, idempotency, concurrency, and absence-monitoring.

11.10 Chapter Summary

- Event-driven loops are triggered by webhooks, schedules, or channel events, and run as a persistent background component rather than a manually-invoked tool.
- A heartbeat is a periodic wake-up that turns an otherwise passive agent into a proactive one.
- Time-based looping (fixed cadence) and condition-based looping (run until a verifiable condition holds) are distinct trigger semantics; the latter is structurally a verification loop running unattended.
- Decoupling invocation from interaction means findings land in an inbox or ticket for asynchronous human review, not a synchronous session.
- Trigger reliability, idempotent event handling, and concurrency control become load-bearing infrastructure concerns the moment a loop runs unattended.

The next chapter introduces the most structurally distinct loop in this book: one whose job is not to complete a task, but to improve every other loop in the system.

11.11 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 11.1

What changes architecturally when a loop becomes a persistent background component rather than an invoked tool?

The loop's trigger and its output decouple from any human's presence: an event fires (webhook, schedule, message), the loop runs, and its findings land somewhere asynchronous — a triage inbox, a ticket board, a pull request — to be reviewed on the human's own schedule. The agent stops being something a person sits down to start and becomes a component running continuously inside a larger system. This shift makes a new set of concerns load-bearing that interactive loops never face: trigger reliability (a missed event must not vanish silently), idempotent event handling (the same event may fire twice), concurrency (overlapping runs against the same resource), and independent monitoring (a silently failed unattended run is far more dangerous than a visible interactive failure). The mechanics of the inner loop are unchanged; what changes is everything around invocation and output.

Answer 11.2

Distinguish time-based from condition-based looping. Why is the distinction important?

Time-based looping re-runs on a fixed cadence regardless of state — every morning, every hour — with no notion of “done”; it simply runs again next cycle. Condition-based looping keeps running, across as many turns as needed, until a specific verifiable condition becomes true (all tests pass, the linter is clean), with a separate check evaluating the stop condition after every turn. The distinction matters because they are different control structures wearing similar clothing: a cron heartbeat is a scheduler, while condition-based looping is structurally a verification loop (Chapter 8) running unattended — the agent that did the work is not the one that decides it is finished. Conflating them leads to bugs like a “scheduled” job that should actually keep retrying until a condition holds, or a condition-based loop that someone caps at a fixed cadence and so abandons before the condition is met. Name which semantics you need before building.

Answer 11.3

What is the heartbeat pattern and what does it enable?

A heartbeat is a periodic wake-up, typically on a cron schedule, that turns an otherwise passive agent into an always-on, proactive system — it checks for new work even when no external event explicitly triggered it. A heartbeat-driven loop might wake every morning, read the prior day's CI failures, newly opened issues, and recent commits, decide for itself whether anything is worth acting on, archive the run silently if nothing is found, and surface findings to a human inbox if something is. It enables proactivity: instead of waiting to be invoked, the system continuously scans its environment and escalates only when warranted. The design subtlety is the silent-archive-versus-surface decision — a heartbeat that interrupts a human on every tick is noise, while one that only surfaces genuine findings is leverage. The reference **Heartbeat** encodes exactly that with a handler returning **None** (archive) or a **Finding** (surface).

Answer 11.4

Why is idempotent event handling essential for event-driven loops, and how do you implement it?

Because webhook and queue delivery is typically at-least-once: the same event will sometimes fire twice, and without idempotency the loop takes a duplicate action — a doubled deploy, a re-sent notification, a duplicate ticket. Idempotent handling makes a redelivered event a no-op. You implement

it with a durable processed-ledger keyed by a stable event id: before acting, check whether the id has been seen; if so, skip; otherwise record it and proceed. The reference `Heartbeat` keeps a `_seen` set and short-circuits redeliveries, counting them as skipped duplicates. In production the ledger must be durable (a database row, not an in-memory set) so it survives restarts, and the check-and-record should be atomic so two concurrent deliveries of the same event do not both slip through. This is the Chapter 6 idempotency requirement applied at the event boundary rather than the tool boundary.

Answer 11.5

Why is a silently failing scheduled loop more dangerous than a failing interactive one?

Because an interactive loop's failure is immediately visible to the person who invoked it — they are sitting there waiting for the result — whereas an unattended, triggered loop's failure can go unnoticed for days unless its own execution is independently monitored. A nightly reconciliation that silently stops running, or a webhook handler that silently stops receiving events, looks exactly like “nothing happened,” which is indistinguishable from “everything is fine” until something downstream breaks. The defense is to monitor the loop's own execution as first-class infrastructure: alert on missed cron runs, track webhook delivery and dead-letter failures, and emit a heartbeat-of-the-heartbeat so the absence of runs is itself an alertable event. The principle is that for unattended systems, you must monitor not just failures but the absence of expected successes — silence is a failure mode.

Answer 11.6

How does condition-based looping relate to the maker/checker split from Chapter 8?

Condition-based looping *is* the maker/checker split running unattended. After every turn, a separate check — not the agent that did the work — evaluates whether the verifiable stop condition holds, and only that independent check can terminate the loop. This is precisely what makes it safe to let the loop run without a human watching each iteration: the agent cannot simply declare itself done and stop, because a fresh checker with no investment in the work decides termination. If you let the working context judge its own completion in an unattended loop, you get the optimistic-self-assessment failure mode (Chapter 8) with no human there to catch it — the loop stops early on confidently-wrong output and nobody notices. So the maker/checker discipline, which is good practice in interactive loops, becomes a safety requirement in unattended ones.

Answer 11.7

Two webhook events for the same resource fire milliseconds apart. What can go wrong and how do you handle it?

Two near-simultaneous events can trigger overlapping loop runs against the same resource, which is the multi-agent concurrency problem (Chapter 10) arriving through the event layer: the two runs read each other's stale state or overwrite each other's work. Handling it requires the same synchronization discipline plus event-layer controls. Options, often combined: a per-resource lock or lease so only one run touches a given resource at a time; debouncing or coalescing rapid events into a single run; serialising runs for the same resource key through an ordered queue; and idempotency so that even if two runs proceed, the net effect is correct. The wrong answer is assuming events arrive neatly spaced and runs never overlap — under real load they do. The design question for every event-driven loop is “what happens when two events for the same target arrive at once?”, answered before launch, not after the first corruption.

Answer 11.8

What does “decoupling invocation from interaction” mean, and why does it matter for UX and safety?

It means the loop's trigger and the loop's output land in different places, and neither requires a human present at the moment the loop runs: an event invokes the loop, and its findings land in

a triage inbox, ticket board, or pull request for review on the human’s own schedule — not in a synchronous session. For UX it matters because it lets the system do work continuously and present results in batch, respecting the human’s attention rather than demanding real-time engagement. For safety it matters because the asynchronous output surface becomes the natural place to insert review gates: findings are proposals a human approves, not actions taken live. The shift is from “human drives each run and sees results immediately” to “system runs autonomously and deposits reviewable output,” which is what makes always-on agents tractable — but it requires designing the output surface (inbox, board, PR) as deliberately as the loop itself.

Answer 11.9

Production case: a heartbeat agent that triages CI failures starts flooding the team’s inbox with low-value findings. How do you fix it without losing the real signal?

The handler’s surface-versus-archive threshold is mis-tuned — it is escalating noise. The fix is to raise the bar for what becomes a **Finding** without losing genuine signal: filter out known-flaky or already-tracked failures (dedupe against open tickets), require a minimum severity or recurrence before surfacing, and cluster related failures into one finding rather than one-per-occurrence. I would also add a feedback signal — track which surfaced findings humans act on versus dismiss — and use that to calibrate the threshold over time (a small hill-climbing loop, Chapter 12). The structural point is that an always-on loop’s value is entirely in the precision of its surface decision: it should wake on every tick but interrupt a human only on genuine, actionable, novel findings. Flooding the inbox is the heartbeat equivalent of the Human-in-the-Loop “interrupting too often” failure from Chapter 9 — it converts leverage back into noise.

Answer 11.10

What are the minimum infrastructure guarantees you would require before putting an unattended loop into production?

Four, each addressing a concern that interactive loops never face. First, trigger reliability with monitoring: missed webhooks or failed cron runs must be detected, retried, and dead-lettered, never silently dropped. Second, idempotent event handling backed by a durable ledger, so at-least-once delivery and restarts do not cause duplicate actions. Third, concurrency control — per-resource locking or serialisation — so overlapping events do not corrupt shared state. Fourth, independent execution monitoring that alerts on the *absence* of expected runs, because silence is a failure mode for unattended systems. I would add resumable, durable state (Chapter 3) so a crash mid-run does not lose or double work, and a reviewable output surface so the loop proposes rather than acts unsupervised on anything irreversible. Without these, an unattended loop is not production-ready no matter how good its reasoning — the reliability lives in the infrastructure around the loop, not the model inside it.

The Hill-Climbing Loop

Subscribe → aiengineeringinsider.substack.com/subscribe

The most valuable loop is the one that improves all the others.

Every loop in this book so far does a job and forgets it. It runs, it finishes, and whatever it learned along the way evaporates. Now imagine a loop whose only job is to read the transcripts of all the others, notice that one tool fails far more often than it should or that a grader has quietly grown too lenient, and rewrite the system so the next thousand runs are better. Most teams build this loop last, if at all. It is the one that compounds.

Every loop covered so far automates work. This chapter covers the loop that automates *improvement* — arguably the most important pattern in this book, and the one most teams build last despite it compounding the value of everything else.

12.1 Traces as a First-Class Artifact

Every agent run produces a trace: a complete record of what the model did, which tools it called, what the verifier returned, and what the final outcome was. Traces are usually treated purely as observability data — something to look at when debugging a specific failure. The hill-climbing loop treats them as something more: training signal for improving the system itself.

★ Key Insight

A trace is not just a debugging artifact. Aggregated across many runs, it is evidence about which parts of the harness — prompts, tools, grader criteria — are systematically underperforming, in a way no single trace can reveal on its own.

12.2 Trace-Analysis Agents

A hill-climbing loop runs a dedicated analysis agent over a batch of production traces, looking for recurring patterns: a tool that fails unusually often, a prompt that produces verifier failures on a specific category of task, a grader rubric that is too lenient or too strict. When the same issue shows up across multiple independent traces, that is a much stronger signal than any single failed run.

12.3 The Return Arrow Reaches Inside

The defining structural feature of this loop, distinct from every other pattern in this book: the feedback does not just loop back to the top of the same run. It reaches inside the system and rewrites the harness itself — the prompts, the tool definitions, or the grader criteria used by every *future* run.




Figure 12.1: The hill-climbing loop: agent runs produce traces, an analysis agent mines them for recurring failures, and the return arrow rewrites the harness used by every future run.

From the Field

Most teams build observability long before they build a hill-climbing loop, which is reasonable — you cannot improve what you cannot see. But observability alone is a dead end if no system ever acts on what it sees. The step that actually compounds value is closing that loop: routing aggregated trace signal into an automated, or at least semi-automated, harness-rewriting step, rather than leaving trace review as a manual, occasional human task.

12.4 What Gets Improved

- **Prompt and tool configuration** — the most straightforward target. A prompt tweak or a tool description rewrite, generated from observed failure patterns and reviewed by a human before deployment.
- **Grader criteria** — if the verification loop from Chapter 8 is systematically too lenient or too strict on a category of task, the hill-climbing loop can surface and propose a rubric correction.
- **The model itself** — for teams running open-weight models, trace and evaluation outcomes can become training signal for reinforcement-learning fine-tuning, closing the loop at the model-weights level rather than only the harness level.
- **Auxiliary context** — memory stores and retrieved skills (Chapter 4) can be improved by the same mechanism: identifying which stored project knowledge is stale, missing, or actively misleading the agent.

★ Key Insight

Prompt and tool configuration are the simplest things to improve through hill-climbing, but they are not the only options. The loop is the pattern; what it optimizes is a design choice, not a limitation of the architecture.

12.5 Loop Stacking as a Design Principle

This chapter completes the picture introduced informally in Chapter 3’s discussion of nested loops. Production agentic systems are best understood not as one loop, but as several loops stacked at increasing scope:

1. The **agent loop** (Chapter 2) automates a single piece of work.
2. The **verification loop** (Chapter 8) wraps the agent loop and ensures quality and correctness.

3. The **event-driven loop** (Chapter 11) wraps both of the above and connects them to the surrounding ecosystem at scale.
4. The **hill-climbing loop** wraps all three and improves them continuously based on aggregated production evidence.

★ Key Insight

Each cycle of the outer loop makes the inner loops more effective. This is the core thesis of loop stacking as a design principle: you do not choose one loop pattern, you compose loops in layers, where each outer layer wraps and improves the one beneath it.

12.6 Human Review of Harness Changes

! Pitfall / Warning

An automated system that rewrites its own prompts, tools, or grading criteria without human review is a meaningfully higher-risk system than one that merely executes a fixed harness. Harness changes proposed by a hill-climbing loop should flow through human review before deployment, exactly as established for the other architectural layers in Chapter 14 — the hill-climbing loop is not an exception to human-in-the-loop oversight, it is the layer where that oversight matters most, because its changes affect every future run, not just one.

Project Code — `loop_engineering_lab`

`loop_engineering_lab/hillclimbing/trace_analyzer.py` implements a batch trace-analysis agent that surfaces recurring failure patterns across a set of logged traces and drafts a proposed harness change for human review.

12.7 The Loop Stack at a Glance

Layer	Job	Timescale
Agent loop (Ch. 2)	Complete one piece of work	Seconds–minutes
Verification loop (Ch. 8)	Ensure quality of each output	Per attempt
Event-driven loop (Ch. 11)	Connect to the ecosystem, run unattended	Continuous
Hill-climbing loop (Ch. 12)	Improve the harness from aggregated traces	Days–weeks

Table 12.1: Loop stacking: each outer layer wraps and improves the one beneath it, at a progressively longer timescale.

12.8 Reference Implementation: Mining Traces for Harness Signal

The hill-climbing loop’s distinctive move is aggregation: a single trace is debugging data, but the same failure recurring across many traces is harness signal. `loop_engineering_lab/hillclimbing/trace_analyzer.py` ranks failure signals across a batch and drafts a change *for human review*.

loop_engineering_lab/hillclimbing/trace_analyzer.py — aggregate, rank, propose

```
def analyze_traces(tracers: list[Tracer]) -> AnalysisReport:
    n = len(tracers)
    signal_counts, traces_with_signal = Counter(), Counter()
    for tracer in tracers:
        seen_here = set()
        for ev in tracer.events:
            signal = None
            if ev.kind == "tool" and ev.detail.get("ok") is False:
                signal = f"tool:{ev.detail.get('name')} failed"
            elif ev.kind == "verify" and ev.detail.get("passed") is False:
                signal = f"verifier:{ev.detail.get('name')} rejected"
            if signal:
                signal_counts[signal] += 1
                seen_here.add(signal)
        for s in seen_here:
            traces_with_signal[s] += 1 # how many DISTINCT runs show this
    findings = [TraceFinding(s, c, traces_with_signal[s] / n)
                for s, c in signal_counts.most_common()]
    return AnalysisReport(n_traces=n, findings=findings)
```

Note the report `draft_harness_change()` never edits anything — it produces a human-reviewable proposal. A system that rewrites its own prompts, tools, or grader is higher-risk than one running a fixed harness, so the human gate is part of the design, not an add-on.

12.9 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 12.1

What makes the hill-climbing loop structurally distinct from every other loop in this book?

Hint. Its return arrow rewrites the harness for every future run.

Exercise 12.2

Why are traces “training signal” rather than just debugging data?

Hint. Aggregated across runs, they reveal systematic harness weaknesses.

Exercise 12.3

What can a hill-climbing loop improve, and which target is the usual starting point?

Hint. Prompts/tools, graders, the model, and memory — prompts first.

Exercise 12.4

Explain loop stacking as a design principle.

Hint. Each outer layer wraps and improves the one beneath it.

Exercise 12.5

Why must harness changes from a hill-climbing loop go through human review?

Hint. Its changes compound across every future run.

Exercise 12.6

How do you measure whether a proposed harness change actually improves the system before rolling it out?

Hint. Offline evaluation on a held-out set before rollout.

Exercise 12.7

What does “hill-climbing” connote here, and what is the risk implied by the name?

Hint. Iterative local improvement — beware local optima.

Exercise 12.8

Most teams build observability before hill-climbing. Why, and where do they get stuck?

Hint. Observability is necessary but not sufficient — close the loop.

Exercise 12.9

Production case: after a hill-climbing-driven prompt change, overall success ticked up but one customer segment regressed badly. What went wrong and how do you prevent it?

Hint. An aggregate win hid a segment regression — stratify the eval set.

Exercise 12.10

For an open-weight model team, how does the hill-climbing loop extend down to the model weights, and what extra caution applies?

Hint. Trace outcomes become RL signal; weight changes are hard to review and revert.

12.10 Chapter Summary

- Traces, aggregated across many runs, are training signal for improving the system — not just debugging data for a single failure.
- A dedicated analysis agent detects recurring patterns across traces that no single trace reveals on its own.
- The hill-climbing loop is structurally distinct: its return arrow rewrites the harness used by every future run, not just the current one.
- Targets for improvement include prompts and tool configuration, grader criteria, the underlying model via RL fine-tuning, and auxiliary context like memory and skills.
- Loop stacking — agent loop, wrapped by verification, wrapped by event-driven execution, wrapped by hill-climbing — is the organizing design principle for production agentic systems.
- Harness changes proposed by a hill-climbing loop require human review before deployment; this is the layer where oversight matters most, since its changes compound across every future run.

Part IV is complete: scaling beyond a single agent, beyond manual invocation, and beyond static improvement. Part V turns to observability, safety, and the economics of running these systems at scale.

12.11 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 12.1

What makes the hill-climbing loop structurally distinct from every other loop in this book?

Its return arrow reaches *inside* the system. Every other loop's feedback loops back to the top of the same run — a failed test drives the next action of the current task. The hill-climbing loop's feedback rewrites the harness itself: the prompts, tool definitions, and grader criteria used by every *future* run, not just the current one. That is a categorical difference: it does not complete a task, it improves the machinery that completes tasks. Aggregated across many runs, traces become evidence about which parts of the harness systematically underperform, and the loop routes that evidence into a harness change. Because its changes compound across all future runs, it is both the highest-leverage loop and the one requiring the most careful human oversight — a bad task output affects one run, a bad harness change affects everything.

Answer 12.2

Why are traces “training signal” rather than just debugging data?

A single trace tells you what happened in one run — useful for debugging that failure. Aggregated across many runs, traces reveal what no single trace can: that a particular tool fails unusually often, that a prompt produces verifier failures on a specific task category, that a grader rubric is systematically too lenient or too strict. That cross-run pattern is evidence about the harness, and evidence about the harness is exactly what you need to improve it. The reference `analyze_traces` makes this concrete by counting both total occurrences and the number of *distinct* runs exhibiting each failure signal — the share across runs is what distinguishes a one-off fluke from a systematic problem worth a harness change. Treating traces only as per-incident debugging data leaves this signal on the floor; treating them as an aggregate dataset is what enables continuous improvement.

Answer 12.3

What can a hill-climbing loop improve, and which target is the usual starting point?

Four targets. Prompt and tool configuration is the usual starting point and the most straightforward — a prompt tweak or a tool-description rewrite generated from observed failure patterns and reviewed before deployment. Grader criteria come next: if the verification loop is systematically too lenient or strict on a category, the loop can propose a rubric correction. The model itself is a target for teams running open-weight models, where trace and evaluation outcomes become training signal for reinforcement-learning fine-tuning, closing the loop at the weights level. And auxiliary context — memory stores and retrieved skills (Chapter 4) — can be improved by identifying which stored knowledge is stale, missing, or actively misleading. The key framing: the loop is the pattern; what it optimises is a design choice, not a limitation. Most teams start with prompts and tools because they are cheap to change and review, and graduate to the others as the system matures.

Answer 12.4

Explain loop stacking as a design principle.

Production agentic systems are best understood not as one loop but as several loops stacked at increasing scope: the agent loop (Chapter 2) automates a single piece of work; the verification loop (Chapter 8) wraps it and ensures quality; the event-driven loop (Chapter 11) wraps both and connects them to the ecosystem at scale, running unattended; and the hill-climbing loop wraps all three and improves them continuously from aggregated production evidence. Each outer layer makes the inner layers more effective, and each operates at a longer timescale — seconds for the agent loop, per-

attempt for verification, continuous for event-driven, days-to-weeks for hill-climbing. The principle is that you do not choose one loop pattern; you compose loops in layers where each outer layer wraps and improves the one beneath. This is the organising architecture for serious systems, and recognising which layer a problem belongs to is half of diagnosing it.

Answer 12.5

Why must harness changes from a hill-climbing loop go through human review?

Because an automated system that rewrites its own prompts, tools, or grading criteria without review is meaningfully higher-risk than one that merely executes a fixed harness — its changes affect every future run, not just one task. A bad task output is a contained, single-run failure; a bad harness change is a systematic regression applied to everything downstream, and it can be subtle (a prompt tweak that improves one category while quietly degrading another). So harness changes should flow through human review before deployment, exactly as the other architectural layers do (Chapter 14) — and this is the layer where that oversight matters most, precisely because the blast radius is every future run. The reference analyzer reflects this: it *drafts* a proposed change and never applies it. Human-in-the-loop here is not a limitation on the loop’s value; it is what makes a self-modifying system safe enough to run at all.

Answer 12.6

How do you measure whether a proposed harness change actually improves the system before rolling it out?

You evaluate it offline on a held-out set before it touches production, the same discipline as any model change. Concretely: maintain a representative eval set of tasks with known-good outcomes; run the current harness and the proposed harness against it; and compare on task-success rate, cost-per-successful-task, and the specific failure category the change targeted — while watching for regressions elsewhere, because a fix for one category often quietly degrades another. Only promote a change that improves the targeted metric without regressing the aggregate. The reference `draft_harness_change` explicitly recommends “adjust, then re-evaluate on a held-out set” for this reason. Rolling a harness change straight to production on the strength of “it should help” is how hill-climbing loops cause systematic regressions; the eval set is the verifier for the meta-loop, and the meta-loop needs a verifier exactly as much as the inner loops do.

Answer 12.7

What does “hill-climbing” connote here, and what is the risk implied by the name?

The name connotes iterative local improvement: each cycle makes a change that improves a metric and keeps it, climbing toward better performance. The risk the name implies is the classic one from optimisation — local optima. A hill-climbing loop that greedily optimises whatever metric it is given can get stuck improving a narrow signal while the broader system plateaus or regresses on dimensions the metric does not capture. Concretely, optimising hard against the grader can produce overfitting to the verifier (Chapter 7) at the harness level: the system climbs the measured hill while real quality stalls. Mitigations are to optimise against a broad, representative eval set rather than a narrow proxy, to keep a human in the loop who can see beyond the metric, and periodically to re-examine whether the metric still reflects real value. The pattern compounds value, but only if the hill it is climbing is the right one.

Answer 12.8

Most teams build observability before hill-climbing. Why, and where do they get stuck?

They build observability first because you cannot improve what you cannot see — traces, metrics, and replay are prerequisites for any analysis. The place teams get stuck is treating observability as

the destination: they accumulate dashboards and trace logs that a human occasionally reviews, but no system ever acts on the aggregate signal, so the value never compounds. Observability alone is a dead end. The step that actually compounds value is closing the loop — routing aggregated trace signal into an automated or semi-automated harness-improvement step, even if the final apply is human-gated. The reference `analyze_traces` plus `draft_harness_change` is the minimal version of that step: it turns a pile of traces into a ranked, actionable proposal. The lesson is that observability is necessary but not sufficient; the return arrow from traces back into the harness is what separates a system that improves from one that is merely well-instrumented.

Answer 12.9

Production case: after a hill-climbing-driven prompt change, overall success ticked up but one customer segment regressed badly. What went wrong and how do you prevent it?

The loop optimised an aggregate metric and found a change that improved the average while degrading a minority segment — a local optimum that the headline number hid. This is the local-optima and overfitting-to-the-metric risk made concrete: the eval set was probably not segment-stratified, so the regression was invisible until production. The immediate fix is to roll back the change (which is why changes are reviewed and reversible, not silently applied). The prevention: stratify the eval set by segment and gate promotion on no-regression *per segment*, not just in aggregate; add the affected segment's cases to the permanent eval set so the same regression cannot recur; and keep the human reviewer empowered to reject an aggregate-positive change that harms a segment. The deeper principle is that the meta-loop needs a verifier broad enough to catch what its optimisation target misses — the same lesson as overfitting to the verifier, now at the level of the system that tunes the verifier.

Answer 12.10

For an open-weight model team, how does the hill-climbing loop extend down to the model weights, and what extra caution applies?

For teams running open-weight models, trace and evaluation outcomes can become training signal for reinforcement-learning fine-tuning, closing the loop at the weights level rather than only the harness level — successful trajectories reinforce desired behaviour, failed ones discourage it. This is the deepest form of the return arrow: instead of rewriting prompts and tools around a fixed model, the loop improves the model itself. The extra caution is substantial. Weight changes are far harder to review, attribute, and roll back than a prompt diff; they can introduce subtle, broad behavioural regressions and even reward-hacking where the model learns to game the trace-derived reward; and the feedback signal must be high-quality, because training on a biased or hackable reward bakes that flaw into the weights. So the same human-review and held-out-eval discipline applies with more force, plus careful reward design, regression testing across many categories, and the ability to revert to a prior checkpoint. The loop pattern is identical; the stakes and irreversibility of the change are much higher.

PART V

Observability, Safety, and Economics

The final two surfaces decide whether a loop is fit to run in production: the *budget* that bounds its cost, and the *human boundary* that keeps judgment where it belongs. Part V covers how to see what a loop did, how to keep a human in the judgment seat as loops scale, how to make loops affordable, and how a team turns one engineer's loop into an organisational discipline.

Observability and Evaluation of Loops

Subscribe → aiengineeringinsider.substack.com/subscribe

A single healthy number is the most reliable place to hide a failure.

Here is a metric that looks healthy: a 92% success rate, flat for a month. Here is what it hides: a tenth of the successful runs took twelve flailing steps to get where three would do, the failing 8% are runaway loops quietly burning the budget, and one customer segment has been broken the whole time. A single number told you everything was fine. It always does. The job of this chapter is to stop trusting the single number and start seeing what the loop actually did.

The hill-climbing loop in Chapter 12 depends entirely on having good traces to analyze. This chapter covers what makes a trace good, and how to evaluate loop behavior at every level from a single step to an entire production fleet.

13.1 Loop Tracing

A complete trace logs every iteration's input, the action taken, the observation received, and the reasoning that connected them — not just the final output. Without step-level logging, debugging a failed run means reconstructing what happened from the final state alone, which is rarely sufficient to identify where the loop actually went wrong.

✓ Best Practice

Log the full decision at every iteration, not just the action taken. Recording *why* the loop chose an action, alongside the action itself, is what makes a trace useful for the hill-climbing analysis in Chapter 12 — a log of actions alone cannot reveal a systematically flawed decision process.

13.2 Step-Level Evaluation

Beyond evaluating whether a run succeeded overall, it is worth evaluating whether each individual *step* made progress. A run can succeed despite several wasted or counterproductive intermediate steps; step-level evaluation surfaces that inefficiency even when the final outcome looks fine.

13.3 Trajectory Evaluation

Trajectory Evaluation

Scoring the entire path a loop took to reach its outcome, not merely the final output. Two runs can reach the identical correct final state by very different paths — one direct, one circuitous and expensive — and trajectory evaluation is what distinguishes them.

This matters because optimizing purely for final-output correctness, without regard to the path taken, can reward loops that happen to stumble into the right answer through an inefficient or fragile process — one that will not generalize reliably to the next, slightly different task.

13.4 Identifying Loop Pathologies

The named failure modes from Chapter 7 — thrashing, overfitting to the verifier, context drift — along with oscillation (alternating between two states without converging to either) and premature termination (stopping before the actual goal is met) are best detected through systematic trace review, not by waiting for a user to report a bad outcome.

! Pitfall / Warning

Runaway loops — those that exhaust their entire budget without ever reaching a terminal state — are often invisible in aggregate success-rate metrics, because a capped, budget-exhausted run can look identical in summary statistics to a deliberately bounded one. Track budget-exhaustion rate as its own explicit metric, separate from success rate.

From the Field

The most expensive bug I have watched a team chase did not throw a single error. Their dashboard showed a healthy, flat success rate, and the cost line crept up so slowly that for weeks everyone blamed traffic growth. The truth was in the trajectory data nobody was looking at: one category of task had started taking nine tool calls to reach an answer that used to take three. Each run still “succeeded,” so the success metric never blinked — it was just quietly paying triple. We only found it when someone finally plotted tool-calls-per-success by category instead of the headline number. The lesson stuck: the metric that looks fine is the one to distrust first.

13.5 A/B Testing Loop Strategies

Different stopping conditions, different planning approaches, and different grader rubrics can all be tested against each other directly, using the same trace-collection infrastructure that feeds the hill-climbing loop. This turns loop design decisions that would otherwise rest on intuition into decisions backed by measured outcomes across a representative sample of real tasks.

13.6 Replay and Debugging Tools

A trace that can be replayed step by step — re-examining the exact context, the exact tool call, and the exact result at any given iteration — is dramatically more useful for debugging than a flat log file. Investing in replay tooling early pays for itself the first time a production failure needs to be root-caused.

13.7 Loop Robustness Testing

★ Key Insight

The hard part of loop engineering is not making it work when everything goes right. It is making it fail gracefully when things go wrong. Before deploying a loop in production, deliberately test it against conditions designed to break it, not just its intended happy path.

This is a distinct discipline from evaluating whether the agent’s task output is correct — it specifically tests whether the *loop’s own control logic* behaves correctly under stress. Useful test conditions include:

- incomplete or genuinely ambiguous task definitions, to confirm the loop escalates rather than guesses confidently;
- tools that return errors or unexpected, malformed output formats, to confirm the error-handling logic from Chapter 7 actually fires;
- tasks that are genuinely unsolvable, specifically to verify the exit condition works and the loop does not run to budget exhaustion silently.

✓ Best Practice

Test your loops on failure cases, not just happy paths. A loop that has never been deliberately pointed at an unsolvable task has an unverified exit condition — you do not actually know it will stop correctly until you have watched it try to stop on a task designed to never succeed.

13.8 The Observability Pipeline




Figure 13.1: Traces flow from runs into per-step, trajectory, run, and fleet evaluation; the aggregate feeds A/B tests and the hill-climbing loop, while replay serves single-run debugging.

13.9 Evaluation Granularity Compared

Level	Question it answers	Catches
Step	Did this iteration make progress?	Wasted / counterproductive steps
Trajectory	Was the whole path efficient?	Right answer via a fragile route
Run	Did the task succeed?	Outright failure
Fleet	What fails systematically?	Recurring harness weaknesses

Table 13.1: Final-output success is the coarsest signal; step, trajectory, and fleet evaluation surface what a pass/fail run metric hides.

13.10 Reference Implementation: Runaway Rate as a First-Class Metric

A runaway run — budget exhausted, no success — looks almost identical to a clean bounded run in aggregate stats, so it must be measured separately. `loop_engineering_lab/observability/metrics.py` computes it explicitly from the trace.

`loop_engineering_lab/observability/metrics.py` — separate the runaway signal

```
def fleet_summary(results: list[LoopResult]) -> dict[str, float]:
    n = len(results) or 1
    rows = [run_metrics(r) for r in results]
    return {
        "n": len(results),
        "success_rate": sum(r.succeeded for r in rows) / n,
        "runaway_rate": sum(r.budget_exhausted for r in rows) / n, # tracked
        ↪ separately!
        "blocked_rate": sum(r.blocked for r in rows) / n,
        "avg_tool_calls": sum(r.tool_calls for r in rows) / n,
        "avg_tokens": sum(r.tokens for r in rows) / n,
    }
```

Because `run_metrics` derives `budget_exhausted` from the transition log’s reason — not merely “not succeeded” — a dashboard built on `fleet_summary` shows the runaway rate as its own line, which is exactly the metric that otherwise hides inside a healthy-looking success rate.

13.11 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 13.1

What distinguishes a good trace from a flat log, and why does it matter downstream?

Hint. It logs the reasoning, not just the action.

Exercise 13.2

Differentiate step-level, trajectory, and run-level evaluation.

Hint. Step progress versus whole-path efficiency versus final success.

Exercise 13.3

Why must budget-exhaustion (runaway) rate be tracked separately from success rate?

Hint. A runaway looks identical to a clean stop in aggregate stats.

Exercise 13.4

What loop pathologies do you look for in systematic trace review, and why not wait for user reports?

Hint. The named failure modes plus oscillation and premature termination.

Exercise 13.5

How does A/B testing of loop strategies work, and what makes it trustworthy?

Hint. Same task sample, vary one thing, compare measured outcomes.

Exercise 13.6

Why invest in replay tooling, and how does it differ from reading logs?

Hint. Re-examine the exact state at any iteration, not a flat log.

Exercise 13.7

What is loop robustness testing, and how is it different from evaluating task-output quality?

Hint. It tests the loop's own control logic, not the task output.

Exercise 13.8

Production case: success rate looks healthy at 92%, but the on-call team is firefighting cost spikes. What do you measure?

Hint. Break the aggregate apart — runaway rate and per-category cost.

Exercise 13.9

How does observability connect the verification loop (Ch. 8) and the hill-climbing loop (Ch. 12)?

Hint. Verify per run, observe across runs, improve the system.

Exercise 13.10

What is the minimum observability you would require before scaling a loop from one task to a fleet?

Hint. Complete traces, separated metrics, per-run cost, and aggregation.

13.12 Chapter Summary

- Complete traces log the reasoning behind each action, not just the action itself, which is what makes them useful for hill-climbing analysis.
- Step-level evaluation surfaces wasted intermediate steps even in runs that ultimately succeed.
- Trajectory evaluation scores the entire path taken, distinguishing an efficient run from one that stumbled into the right answer.
- Track budget-exhaustion rate explicitly; runaway loops can hide inside otherwise normal-looking success-rate metrics.
- A/B testing loop strategies turns design intuitions into measured decisions using the same trace infrastructure as hill-climbing.
- Loop robustness testing — deliberately probing with ambiguous tasks, malformed tool output, and unsolvable problems — verifies the loop’s own control logic, separate from evaluating task output quality.

13.13 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 13.1

What distinguishes a good trace from a flat log, and why does it matter downstream?

A good trace logs every iteration's input, the action taken, the observation received, and crucially the *reasoning* that connected them — not just the final output and not just the actions. A flat log of actions tells you what the loop did; a complete trace tells you why, which is what makes it possible to diagnose a systematically flawed decision process rather than just a single bad action. It matters downstream because the hill-climbing loop (Chapter 12) depends entirely on this: aggregated across runs, the reasoning-level record reveals which prompt or tool is causing a category of failure, something a log of actions alone cannot show. The discipline is to record the full decision at every step, treating “why did the loop choose this?” as first-class data, because you cannot improve a decision process you never recorded.

Answer 13.2

Differentiate step-level, trajectory, and run-level evaluation.

Run-level evaluation asks only whether the task succeeded overall — the coarsest signal. Step-level evaluation asks whether each individual iteration made progress, surfacing wasted or counterproductive steps even in a run that ultimately succeeded. Trajectory evaluation scores the entire path the loop took, distinguishing two runs that reach the identical correct answer by very different routes — one direct, one circuitous and expensive. They matter together because optimising purely for final-output correctness can reward a loop that stumbled into the right answer through a fragile, inefficient process that will not generalise to the next slightly different task. A run that passes but took twelve failing steps is a latent reliability and cost problem that only step and trajectory evaluation expose; the binary success metric calls it a win.

Answer 13.3

Why must budget-exhaustion (runaway) rate be tracked separately from success rate?

Because a runaway run — one that exhausts its entire budget without reaching a terminal success — can look identical in aggregate success-rate statistics to a deliberately bounded run that stopped cleanly. Both are simply “not a success,” so a fleet with a rising runaway problem can show a stable success rate while quietly burning money and latency on loops that never converge. Tracking budget-exhaustion as its own explicit metric makes that failure visible. In the reference `fleet_summary`, `runaway_rate` is a separate line derived from whether the failed terminal's reason mentions budget, so a dashboard distinguishes “failed cleanly” from “ran away.” The general lesson: aggregate metrics hide failure modes that do not change the headline number, so any failure with real cost needs its own dedicated signal rather than being folded into a catch-all.

Answer 13.4

What loop pathologies do you look for in systematic trace review, and why not wait for user reports?

The named failure modes from Chapter 7 — thrashing (changing without converging), overfitting to the verifier (passing the check while missing the requirement), and context drift (acting on stale assumptions) — plus oscillation (alternating between two states) and premature termination (stopping before the goal is met). You look for them in trace review rather than waiting for user reports because by the time a user reports a bad outcome, the loop has already shipped it, and many of these pathologies (a run that succeeded but thrashed, a runaway that was capped) never produce a user-visible complaint at all yet still cost money and erode reliability. Systematic trace review

catches the pattern across runs before it becomes an incident. Waiting for user reports means only detecting the subset of failures that happen to be both user-visible and reported — a biased, lagging, incomplete signal.

Answer 13.5

How does A/B testing of loop strategies work, and what makes it trustworthy?

You run two variants — different stopping conditions, planning approaches, or grader rubrics — against the same representative sample of real tasks, using the same trace-collection infrastructure that feeds the hill-climbing loop, and compare measured outcomes. What makes it trustworthy is the same rigour as any experiment: a representative and sufficiently large task sample (so the result is not noise), holding everything else constant between arms, and comparing on the metrics that matter — task success, cost-per-successful-task, runaway rate, trajectory efficiency — not just the one the change was meant to improve, so you catch regressions. This turns loop-design decisions that would otherwise rest on intuition (“I think replanning more often helps”) into decisions backed by data. The trap is testing on an unrepresentative sample or eyeballing a handful of runs, which produces confident conclusions from noise.

Answer 13.6

Why invest in replay tooling, and how does it differ from reading logs?

Replay lets you re-examine the exact context, the exact tool call, and the exact result at any given iteration of a past run, stepping through it as it happened — which is dramatically more useful for root-causing a failure than reconstructing events from a flat log file. The difference is fidelity and interactivity: a log tells you a sequence of things happened; replay lets you stand at iteration seven and inspect precisely what the model saw and decided, which is usually where the bug actually is. It pays for itself the first time a production failure needs root-causing, because the alternative — inferring what went wrong from final state and scattered log lines — is slow and often inconclusive. Investing early matters because you cannot replay a run you did not trace completely, so the tracing discipline and the replay tooling have to be built before the incident, not after.

Answer 13.7

What is loop robustness testing, and how is it different from evaluating task-output quality?

Loop robustness testing specifically probes whether the loop’s own *control logic* behaves correctly under stress — it tests the loop, not the task output. You deliberately point the loop at conditions designed to break it: incomplete or genuinely ambiguous task definitions (does it escalate rather than confidently guess?), tools that return errors or malformed output (does the Chapter 7 error handling actually fire?), and genuinely unsolvable tasks (does the exit condition work, or does it run to budget exhaustion silently?). This is distinct from evaluating whether the agent’s answer is correct, which assumes the happy path. The key insight is that the hard part of loop engineering is not working when everything goes right; it is failing gracefully when things go wrong — and a loop that has never been pointed at an unsolvable task has an unverified exit condition. You do not actually know it stops correctly until you have watched it try to stop on something that can never succeed.

Answer 13.8

Production case: success rate looks healthy at 92%, but the on-call team is firefighting cost spikes. What do you measure?

A healthy success rate hiding cost pain is the classic runaway-and-trajectory blind spot. I would break the aggregate apart: track runaway rate (budget-exhausted runs) separately, because capped runaway loops can sit inside the 8% non-success without anyone noticing they are expensive; measure cost-per-run and cost-per-successful-run by task category and loop pattern (Chapter 9), to find which

specific pattern is disproportionately expensive; and run trajectory evaluation on the successful 92%, because a run that succeeds in twelve flailing steps costs far more than one that succeeds in three, and the binary success metric treats them identically. The likely culprit is some category thrashing to success or running away to the cap; the fix follows from the data — a stricter no-progress detector, model routing, or a better verifier for that category. The meta-point: when the headline metric looks fine but the bill does not, you are measuring the wrong granularity.

Answer 13.9

How does observability connect the verification loop (Ch. 8) and the hill-climbing loop (Ch. 12)?

Observability is the shared substrate. The verification loop produces a per-attempt accept/reject signal; tracing records that signal, the reasoning around it, and the eventual outcome for every run. The hill-climbing loop then consumes the *aggregate* of those traces to find systematic harness weaknesses — a grader that is too lenient on a category, a prompt that produces verifier failures on a task type. So verification generates the per-run truth, observability records and aggregates it, and hill-climbing acts on the aggregate to improve the harness, after which the cycle repeats with better verification. Without good observability in the middle, the verification signal is consumed once and thrown away, and the hill-climbing loop has nothing to learn from. The three form a pipeline: verify each run, observe across runs, improve the system — which is exactly the loop-stacking picture, with observability as the connective tissue.

Answer 13.10

What is the minimum observability you would require before scaling a loop from one task to a fleet?

Complete per-step traces (input, action, observation, and reasoning) that are durable and replayable; run-level metrics that separate success, blocked, and runaway (budget-exhausted) rather than collapsing them into pass/fail; per-run cost accounting so expensive patterns and categories are identifiable; and a way to aggregate across runs for fleet-level pathology detection and A/B testing. I would also require robustness tests for the loop's control logic — proven exit conditions on unsolvable tasks — before unattended scale. The principle is that scaling multiplies whatever you cannot see: one badly-traced loop is debuggable by a human watching it, but a fleet of them is only as manageable as your ability to reconstruct, measure, and compare runs after the fact. Observability is the prerequisite for every later improvement mechanism in the book, so it has to be in place *before* the fleet, not retrofitted after the first opaque production incident.

Safety, Control, and Human Oversight

Subscribe → aiengineeringinsider.substack.com/subscribe

The deepest risk is not a machine that fails — it is a human who stopped watching because it rarely does.

There is a quiet moment that arrives once a loop becomes reliable. It has worked a hundred times; you stop reading its output and start clicking approve. That moment feels like productivity. It is the most dangerous thing in this book. The deepest risk in mature loop engineering is not a model that fails — it is a human who has stopped paying attention because the machine so rarely does. The whole point of keeping a human in the loop is the judgment they bring, and judgment is exactly the thing that erodes when nothing seems to require it.

Automation does not mean removing humans from the system. As loops scale across every layer covered in this book — single agent, verification, multi-agent, event-driven, hill-climbing — the question is not whether humans stay involved, but precisely where their judgment adds value that no grader or check can substitute for.

14.1 Human-in-the-Loop by Architectural Layer

Every loop layer in this book has a natural point where human oversight belongs:

- **Agent loop** — require explicit human approval before any sensitive action or tool call (a database write, a financial transaction, a credential change).
- **Verification loop** — a human can act as the grader directly for sensitive workflows, rather than an automated rubric.
- **Event-driven loop** — a human approves output before it reaches an end user, even when the loop itself ran fully unattended.
- **Hill-climbing loop** — harness improvements flow through human review before deployment, as established in Chapter 12, because changes here compound across every future run.

★ Key Insight

An automated grader can check whether links resolve. It takes a human to notice the framing is wrong for the audience. That kind of judgment — earned from context, experience, and taste — is exactly where human review earns its place, and it does not disappear as the surrounding system becomes more automated.

14.2 Permission Boundaries Per Iteration

Not every iteration of a loop carries the same risk. Read-only actions (searching, reading files, running a non-destructive test) can generally proceed without per-action approval. Actions with real-world side effects — particularly irreversible ones — need an explicit permission boundary checked at the moment of execution, not assumed from a blanket approval given at the start of the run.

14.3 Sandboxing Iterative Tool Execution

! Pitfall / Warning

A loop that can execute arbitrary code should do so inside an isolated, resource-limited sandbox, not against a production environment or an unrestricted host system. This matters specifically *because* the loop runs iteratively and potentially unattended: a single bad action in an unsandboxed environment is bad enough; a loop that keeps acting against that same unsandboxed environment across many iterations compounds the risk with every cycle.

14.4 Guardrails Re-Evaluated at Every Step

A guardrail checked only once, at the start of a run, is a guardrail that has already gone stale by iteration five. The state of the environment changes with every action the loop takes; permission and safety checks need to be re-evaluated against the *current* state at every step, not assumed to still hold from the initial check.

14.5 Audit Logging Per Iteration

For compliance and for post-hoc review, every iteration's action, the reasoning behind it, and its outcome should be logged in a durable, auditable form — building directly on the loop-tracing discipline from Chapter 13, but framed specifically around the question “can we reconstruct exactly what happened and why, after the fact, for someone who was not watching at the time.”

14.6 Kill Switches

Every loop that can run unattended needs an external mechanism to halt it immediately, independent of the loop's own internal logic. A circuit breaker (Chapter 7) handles automatic halting based on the loop's own observed failure pattern; a kill switch handles the case where a human, observing something the loop's own logic has not detected, needs to stop it regardless.

14.7 The Orchestration Tax

The Orchestration Tax

The principle that a human’s review bandwidth, not the available tooling, is the actual ceiling on how many parallel agent loops can be run responsibly. Worktree-based isolation (Chapter 10) removes the mechanical collision problem between concurrent agents; it does not remove the bottleneck of a human needing to review and trust what each of those agents produced.

From the Field

It is technically straightforward to spin up ten parallel agent loops today. It is not equally straightforward to meaningfully review ten parallel streams of output with the same care you would give to one. The tooling scaled faster than the human review capacity it depends on, and that gap is exactly where quality problems tend to surface in teams that scale loops aggressively without scaling review practice alongside them.

14.8 Comprehension Debt

Comprehension Debt

The growing gap between what exists in a codebase or system and what the responsible engineer actually understands about it, as a direct consequence of loop-driven work shipping faster than a human can absorb and verify it.

A smoother, faster loop does not reduce this risk — it accelerates it, unless the engineer deliberately reads and understands what the loop produced rather than treating “the verifier passed” as equivalent to “I understand this.”

14.9 Cognitive Surrender

! Pitfall / Warning

When a loop runs itself reliably, it becomes tempting to stop forming an independent opinion about its output and simply accept whatever it returns. This temptation, not any single technical failure mode, is the deepest risk in mature loop engineering. Designing and running a loop with genuine judgment is the cure; running it to avoid having to think about the work at all is the same action producing the opposite outcome.

★ Key Insight

Two engineers can build the exact same loop and get opposite results. One uses it to move faster on work they understand deeply. The other uses it to avoid understanding the work at all. The loop cannot tell the difference between these two uses. The engineer has to.

14.10 Where Human Oversight Attaches




Figure 14.1: Each loop layer has a natural human-oversight touchpoint: per-action approval, human grading, output review, and harness-change review — with an external kill switch spanning all of them.

14.11 Kill Switch vs. Circuit Breaker

Property	Circuit breaker (Ch. 7)	Kill switch (Ch. 14)
Trigger	The loop's own observed failure pattern	A human, observing externally
Scope	A category of action	The entire loop
Detects	What the loop's logic can see	What the loop's logic missed
Authority	Automatic, internal	External, overrides the loop

Table 14.1: Both halt a loop, but a circuit breaker is the loop reacting to itself; a kill switch is a human overriding it from outside.

14.12 Reference Implementation: Kill Switch and Per-Step Guard

Two controls unattended loops always need: an external halt independent of the loop's logic, and a guard that re-evaluates permission against the *current* state every step. Both live in `loop_engineering_lab/core/safety.py`.

`loop_engineering_lab/core/safety.py` — external halt + fresh re-check

```

@dataclass
class StepGuard:
    sensitive: set[str]
    approved: set[str]

    def allow(self, action, *, state_is_safe) -> tuple[bool, str]:
        if action not in self.sensitive:
            return True, "non-sensitive"
        if action not in self.approved:
            return False, f"'{action}' requires explicit approval"
        # Re-evaluate against CURRENT state, not the state at approval time.
  
```

```
if not state_is_safe():
    return False, f'"{action}" approved earlier but current state is unsafe'
return True, "approved and state safe"
```

The `state_is_safe` callback is the point: a guardrail checked once at the start of a run has gone stale by iteration five, so the gate takes a *fresh* read of the environment each call. The `KillSwitch` in the same module is checked at the top of every iteration, so a human can stop a loop for a reason its own logic never detected.

14.13 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 14.1

Where does human oversight attach across the loop stack, and what does each touchpoint catch?

Hint. Each layer has a touchpoint requiring judgment no check can replace.

Exercise 14.2

Why must permission boundaries be re-evaluated every iteration rather than approved once at the start?

Hint. State changes each step, so a stale check silently expires.

Exercise 14.3

Distinguish a kill switch from a circuit breaker. Why do you need both?

Hint. Internal automatic reaction versus external human override.

Exercise 14.4

Why sandbox iterative tool execution specifically, beyond ordinary code-execution caution?

Hint. Iteration compounds the risk of a single bad action.

Exercise 14.5

What is the orchestration tax, and what does it imply for scaling parallel agents?

Hint. Review bandwidth, not tooling, is the real ceiling.

Exercise 14.6

Define comprehension debt and explain why a faster loop makes it worse, not better.

Hint. Code outpaces understanding; a faster loop accelerates the gap.

Exercise 14.7

What is cognitive surrender, and why is it called the deepest risk in mature loop engineering?

Hint. Accepting output without forming an opinion — a human-only risk.

Exercise 14.8

Design the human-approval flow for an agent that can run destructive database commands. What are the requirements?

Hint. Fail-closed, per-action, re-evaluated, auditable, no ambient credentials.

Exercise 14.9

Production case: a team scaled to 20 parallel coding agents; velocity rose, then a series of subtle bugs shipped that all passed CI. What happened?

Hint. Orchestration tax plus comprehension debt plus verifier overfitting.

Exercise 14.10

Two engineers build the identical loop and get opposite results. What determines the difference, and what does that imply for how you design oversight?

Hint. The human's use, not the loop, decides — design oversight in by construction.

14.14 Chapter Summary

- Every architectural layer — agent, verification, event-driven, hill-climbing — has a natural human-oversight touchpoint; automation does not remove the need for any of them.
- Permission boundaries and safety guardrails must be checked against current state at every iteration, not assumed to hold from an initial check.
- Sandbox any loop capable of executing arbitrary code, specifically because iterative, unattended execution compounds the risk of a single unsandboxed action.
- A kill switch is an external halt mechanism, independent of the loop's own logic; a circuit breaker is the loop's own automatic response to observed failure.
- The orchestration tax: review bandwidth, not tooling, is the real ceiling on how many parallel loops a team can run responsibly.
- Comprehension debt and cognitive surrender are human-factors risks, not technical failure modes — and they grow specifically as loops get faster and more reliable, not as they get worse.

14.15 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 14.1

Where does human oversight attach across the loop stack, and what does each touchpoint catch?

Each layer has a natural oversight point. In the agent loop, require explicit human approval before any sensitive or irreversible action (a database write, a transfer, a credential change). In the verification loop, a human can serve as the grader directly for sensitive workflows where judgment beats a rubric. In the event-driven loop, a human approves output before it reaches an end user even though the loop ran unattended. In the hill-climbing loop, harness changes flow through human review before deployment, because they compound across every future run. Each touchpoint catches something automation cannot: an automated grader can check that links resolve, but it takes a human to notice the framing is wrong for the audience. Automation does not remove these touchpoints; it makes choosing them deliberately more important, because the question is precisely where human judgment adds value no check can substitute for.

Answer 14.2

Why must permission boundaries be re-evaluated every iteration rather than approved once at the start?

Because the state of the environment changes with every action the loop takes, so a guardrail checked only at the start has already gone stale by iteration five. A blanket approval given at run start cannot account for what the loop did to the world in between — an action that was safe to approve initially may be unsafe against the current state, and a permission assumed to still hold may no longer apply. So permission and safety checks must be evaluated against the *current* state at the moment of execution, every step. The reference `StepGuard` encodes this by taking a fresh `state.is_safe` read on each call rather than trusting the approval timestamp: even an explicitly approved sensitive action is blocked if the current state is unsafe. The principle is that in an iterative system, any check that is not re-evaluated is a check that is silently expiring.

Answer 14.3

Distinguish a kill switch from a circuit breaker. Why do you need both?

A circuit breaker (Chapter 7) is automatic and internal: it halts a category of action based on the loop's *own* observed failure pattern — the loop reacting to what it can see. A kill switch is external: it halts the entire loop because a human, observing something the loop's logic has not detected, needs to stop it regardless. You need both because they cover disjoint cases. The breaker handles failures the loop can recognise (repeated errors of a class); the kill switch handles everything the loop cannot recognise — a subtly wrong-but-passing output, a changed external situation, a judgment call — where a human must override from outside. A system with only a breaker trusts the loop to detect every reason it should stop, which is exactly the assumption that fails. The reference `KillSwitch` is checked at the top of every iteration so the override is always available, independent of the loop's internal logic.

Answer 14.4

Why sandbox iterative tool execution specifically, beyond ordinary code-execution caution?

Because iteration compounds risk. A single bad action in an unsandboxed environment is bad enough, but a loop keeps acting against that same environment across many iterations, potentially unattended — so the risk multiplies with every cycle rather than being a one-shot exposure. A loop that can execute arbitrary code therefore belongs in an isolated, resource-limited sandbox (network-restricted,

CPU/memory/time-capped), not against a production environment or an unrestricted host. The iterative, possibly-unattended nature is the differentiator from one-off code execution: there is no human watching each cycle to catch the third iteration going wrong, and the loop's own momentum can drive it to keep acting on a bad trajectory. Sandboxing bounds the blast radius regardless of how many times the loop acts, which is the only safe assumption for something that runs in a loop.

Answer 14.5

What is the orchestration tax, and what does it imply for scaling parallel agents?

The orchestration tax is the principle that a human's review bandwidth, not the available tooling, is the actual ceiling on how many parallel agent loops can be run responsibly. It is technically trivial to spin up ten parallel loops; it is not trivial to review ten parallel streams of output with the care you would give one. Worktree-based isolation (Chapter 10) removes the mechanical collision problem between concurrent agents, but it does nothing to remove the bottleneck of a human needing to review and trust what each agent produced — if anything more agents mean more to review. The implication is that you scale review practice alongside parallelism, not after: better diffs, scoped changes, reliable automated gates that genuinely reduce what needs human eyes, and a realistic cap on concurrency tied to review capacity. Teams that scale agents aggressively without scaling review are exactly where quality problems surface, because the tooling outran the human capacity it depends on.

Answer 14.6

Define comprehension debt and explain why a faster loop makes it worse, not better.

Comprehension debt is the growing gap between what exists in a codebase or system and what the responsible engineer actually understands about it, as a direct consequence of loop-driven work shipping faster than a human can absorb and verify it. A smoother, faster loop accelerates this rather than relieving it: the faster the loop produces verified-passing changes, the faster code the engineer has not actually read accumulates, and “the verifier passed” quietly substitutes for “I understand this.” The danger is that the debt is invisible until it is called — the system works until a change requires understanding nobody has. The mitigation is deliberate: the engineer reads and understands what the loop produced rather than treating a green check as comprehension, and the team treats review-for-understanding as non-negotiable, not as overhead to be optimised away. Velocity without comprehension is borrowing against future debuggability.

Answer 14.7

What is cognitive surrender, and why is it called the deepest risk in mature loop engineering?

Cognitive surrender is the temptation, once a loop runs itself reliably, to stop forming an independent opinion about its output and simply accept whatever it returns. It is the deepest risk not because it is a technical failure but because it is a human one that no technical control catches: the loop works, the checks pass, and the engineer gradually stops thinking about the work. The same loop can produce opposite outcomes depending on the engineer using it — one uses it to move faster on work they understand deeply, the other uses it to avoid understanding the work at all, and the loop cannot tell the difference. That is why it is the deepest risk: every other failure mode has a detector, but cognitive surrender is invisible to the system and only the engineer can prevent it. The cure is using the loop with genuine judgment; the failure is using it to avoid having to think.

Answer 14.8

Design the human-approval flow for an agent that can run destructive database commands. What are the requirements?

The approval must be fail-closed, per-action, re-evaluated against current state, and auditable. Concretely: destructive commands (drops, mass deletes, schema changes) are blocked by default and require explicit human approval for that specific action, not a blanket session approval (Chapter 3's destructive-action guard); the approval is checked at the moment of execution against the current state, because state changed since any earlier approval (the **StepGuard** pattern); credentials are never directly accessible to the loop, so even a hijacked instruction cannot widen its authority; every approval decision, the action, its reasoning, and outcome are logged durably for audit; and a kill switch can halt the loop independently. The model must never be able to self-authorise by asserting it has approval in its own text — approval comes from an out-of-band human action that populates the approved set. These requirements hold whether the loop is interactive or unattended; unattended only raises the stakes.

Answer 14.9

Production case: a team scaled to 20 parallel coding agents; velocity rose, then a series of subtle bugs shipped that all passed CI. What happened?

Three of this chapter's risks compounded. The orchestration tax was exceeded — 20 streams of output overran the team's review bandwidth, so changes got rubber-stamped rather than reviewed. That produced comprehension debt: code passing CI accumulated faster than anyone understood it, and "CI passed" substituted for comprehension. And the subtle bugs passing CI are overfitting to the verifier (Chapter 7) at fleet scale — the test suite did not cover the behaviour that broke, and with no human truly reviewing, nothing caught it. The fix is not more agents or a bigger model: cap concurrency to what review can actually absorb; strengthen the CI signal so "green" means more (broaden coverage, add the missed behaviours); require genuine review-for-understanding on agent changes, especially risky ones; and treat the green check as necessary but not sufficient. The root cause is scaling the tooling past the human review capacity it depends on — exactly the orchestration tax made real.

Answer 14.10

Two engineers build the identical loop and get opposite results. What determines the difference, and what does that imply for how you design oversight?

The difference is entirely in how the engineer uses it: one uses the loop to move faster on work they understand and continue to review with judgment; the other uses it to avoid understanding the work at all and accepts outputs uncritically. The loop is identical and cannot tell the two uses apart — the human has to. The implication for oversight design is that you cannot rely on technical controls alone to produce good outcomes, because the deepest risks (cognitive surrender, comprehension debt) are human-factors risks that grow as the loop gets *better*, not worse, and have no automated detector. So oversight has to be designed to keep the human in the judgment seat by construction: review-for-understanding as a required step, explicit ownership boundaries for decisions that stay human, approval gates on the irreversible, and a culture that treats the loop as leverage for understood work rather than a substitute for understanding. Build the loop like someone who intends to remain the engineer, not merely the person who presses go.

Cost, Latency, and Efficiency Engineering

Subscribe → aiengineeringinsider.substack.com/subscribe

Reliability decides whether a loop works. Economics decides whether it gets to stay.

A loop that solves the task for four dollars when the task is worth one has not automated anything. It has invented a more expensive way to lose money. This is the chapter where loop engineering meets the invoice — where a verification pass that doubles your cost is either the smartest spend you make or pure waste, depending entirely on what being wrong would have cost you. Reliability got the loop working. Economics decides whether it gets to stay.

A loop that works reliably but costs more than the value of the task it completes is not a solution, it is a liability. This chapter covers the economics of running loops in production, where usage patterns and resulting costs can vary wildly depending on how aggressively a team scales parallel execution.

15.1 Cost-Per-Iteration Accounting

Total cost for a single loop run is the sum across every iteration: every model call, every tool invocation, every verification pass. Tracking this at the level of an individual run, not just in aggregate across a fleet, is what makes it possible to identify which specific loop *pattern* (Chapter 9) or which specific task category is disproportionately expensive relative to the value it delivers.

! Pitfall / Warning

Token costs in particular can vary wildly between teams and even between individual engineers on the same team, depending on how token-rich or token-constrained their usage patterns are. A loop design that is comfortably affordable for one team's usage pattern can be a meaningful cost driver for another. Measure your own actual cost per run; do not assume a published benchmark transfers directly.

15.2 Latency Optimization

- **Minimize round trips** — batch independent tool calls within an iteration, as established in Chapter 6, rather than issuing them one at a time.
- **Caching** — cache tool results and intermediate computations that are likely to be reused across nearby iterations or across separate runs against the same underlying resource.
- **Prompt caching** — where the underlying model API supports it, structure prompts so that the large, stable portion of context (system instructions, project skills) is cached across iterations, and only the smaller, changing portion is re-processed each time.

15.3 Adaptive Loop Depth: Model Routing

Not every iteration needs the same model. A common and effective pattern: use a cheaper, faster model for routine iterations — straightforward tool calls, simple reasoning steps — and escalate to a stronger, more expensive model specifically when the loop detects it is stuck (Chapter 7) or when a verification step (Chapter 8) fails repeatedly.

✓ Best Practice

Route by difficulty, not by default. A long debugging loop can consume a large number of model calls; a quick, well-specified copy edit does not need a frontier reasoning model for every step. Matching model cost to the actual difficulty of each iteration, rather than routing every workflow through the same fixed model, is one of the highest-leverage cost optimizations available.

15.4 Batching and Parallelization

Where sub-tasks are genuinely independent — as established in the multi-agent architectures of Chapter 10 — running them in parallel reduces wall-clock latency even when it does not reduce total token cost. This is a latency-versus-cost tradeoff specifically, not a free optimization: parallel execution generally increases total spend (more concurrent model calls) in exchange for a faster result.

15.5 Context Pruning

The state-compaction techniques from Chapter 4 — rolling summaries, sliding windows, scratchpad isolation — directly reduce per-iteration token cost by keeping the context sent to the model on each call as small as it can be while remaining sufficient. Every unnecessary token carried forward into a new iteration is paid for again on that iteration, and the one after, for as long as it remains in context.

15.6 Tool-Call Budgets as a Failure Signal

★ Key Insight

Tool-call budgets, introduced in Chapter 3 as a cost control, serve a second purpose worth emphasizing here: an agent exhausting its budget without making progress is itself a failure signal, not merely a stopping point. Treat budget exhaustion as a trigger to escalate or change strategy, the same way Chapter 3 established for loop budgets generally — not simply as the mechanical end of a run that gets logged and forgotten.

15.7 Putting It Together: A Cost-Aware Loop

A cost-aware production loop typically combines several of these techniques at once: a cheap, fast model handles routine iterations; context is kept compact through rolling summarization; independent sub-tasks run in parallel where the latency benefit justifies the added spend; and the loop escalates to a stronger model, or to a human, the moment its tool-call budget is exhausted without measurable progress — rather than continuing to spend at the same rate indefinitely.

From the Field

The single biggest cost mistake I see in early-stage loop deployments is not an inefficient model choice — it is a missing budget-exhaustion-as-signal mechanism. Teams cap the budget correctly, the loop stops correctly, and then the same stuck task gets resubmitted and re-runs the entire expensive loop from scratch with no memory that this exact failure already happened. The fix is cheap: persist the failed-attempt record (Chapter 4) so the next attempt starts smarter, not from zero.

15.8 A Cost-Aware Loop




Figure 15.1: A cost-aware loop routes routine steps to a cheap model and escalates to a strong model (or a human) only when stuck or when verification keeps failing.

15.9 Cost Levers Compared

Lever	Reduces	Trade-off
Model routing	Cost per routine step	Routing logic + a misroute risk
Prompt caching	Re-processing stable context	API support required
Context pruning	Tokens carried forward	Risk of dropping needed state
Parallelization	Wall-clock latency	Increases total spend
Budget-as-signal	Wasted spend on stuck loops	Needs a no-progress detector

Table 15.1: Most levers cut cost or latency; parallelization is the one that buys latency *at* the cost of total spend.

15.10 Reference Implementation: Routing by Difficulty

The highest-leverage cost optimisation is matching model cost to actual difficulty — cheap model for routine steps, strong model only when the loop is struggling. `loop_engineering_lab/economics/router.py` pairs that router with a per-iteration cost meter.

`loop_engineering_lab/economics/router.py` — route by difficulty, not by default

```
@dataclass
class ModelRouter:
    cheap: str
    strong: str
    escalate_after_failures: int = 2

    def route(self, *, stuck: bool = False, verifier_failures: int = 0) -> str:
        # Pay frontier prices only when the loop is actually struggling.
        if stuck or verifier_failures >= self.escalate_after_failures:
            return self.strong
        return self.cheap

@dataclass
class CostMeter:
    prices_per_1k: dict[str, float]
    total_usd: float = 0.0

    def charge(self, model, prompt_tokens, completion_tokens) -> float:
        cost = (prompt_tokens + completion_tokens) / 1000 *
            ↪ self.prices_per_1k.get(model, 0)
        self.total_usd += cost
        return cost
```

The meter exists so cost is attributable per run and per model, which is what lets you find the specific loop pattern or task category that is expensive relative to the value it delivers — the prerequisite for deciding where routing or pruning will actually pay off.

15.11 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 15.1

Why account for cost per individual run rather than only in fleet aggregate?

Hint. Per-run accounting attributes spend to a pattern or category.

Exercise 15.2

Explain difficulty-based model routing and why it is high-leverage.

Hint. Cheap model by default, strong model only when struggling.

Exercise 15.3

How do caching and prompt caching reduce cost and latency, and what are their limits?

Hint. Isolate the stable prefix; mind invalidation and API support.

Exercise 15.4

Why is parallelization not a free optimization?

Hint. It buys latency at the cost of total spend.

Exercise 15.5

How does context pruning directly reduce cost, and what is the risk of over-pruning?

Hint. Every carried token is paid again; do not drop needed state.

Exercise 15.6

Budget exhaustion is a cost control. What is its second, more valuable purpose?

Hint. Exhaustion without progress is a signal to escalate.

Exercise 15.7

Walk through the components of a cost-aware production loop.

Hint. Routing, compaction, caching, selective parallelism, budget-as-signal.

Exercise 15.8

Production case: a loop's cost per task tripled month-over-month while success rate held steady. How do you find and fix it?

Hint. A path-level cost regression hidden by stable success — measure the trajectory.

Exercise 15.9

When does adding a verification loop save money rather than cost money, despite the extra calls?

Hint. When the cost of being wrong exceeds the extra verification calls.

Exercise 15.10

How do you set token and tool-call budgets to balance cost against task completion?

Hint. From the eval distribution, above the 95th percentile of successes.

15.12 Chapter Summary

- Track cost per individual run, not just in fleet-wide aggregate, to identify which loop patterns or task categories are disproportionately expensive.
- Reduce latency through parallel independent tool calls, result caching, and prompt caching for stable context.
- Route iterations to cheaper or stronger models based on actual difficulty, escalating specifically when the loop is stuck or verification repeatedly fails.
- Parallelization trades increased total cost for reduced wall-clock latency; it is not a free optimization.

- Context pruning directly reduces per-iteration cost, since every token carried forward is paid for again on every subsequent call.
- Treat tool-call budget exhaustion as a failure signal that should trigger escalation, not just a stopping point that ends the run with no lasting record.

15.13 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 15.1

Why account for cost per individual run rather than only in fleet aggregate?

Because aggregate cost tells you the total bill but not where it comes from, and you cannot optimise what you cannot attribute. Per-run accounting — summing every model call, tool invocation, and verification pass for a single run — lets you identify which specific loop pattern (Chapter 9) or task category is disproportionately expensive relative to the value it delivers, which is exactly the signal that tells you where routing, pruning, or a better verifier will pay off. Fleet aggregate hides a category that is 5% of runs but 40% of spend. The reference **CostMeter** attributes spend per model and per run for this reason. There is also a measurement caveat: token costs vary wildly between teams and even engineers depending on usage patterns, so a design affordable for one team can be a real cost driver for another — measure your own actual cost per run rather than assuming a published benchmark transfers.

Answer 15.2

Explain difficulty-based model routing and why it is high-leverage.

Difficulty-based routing uses a cheaper, faster model for routine iterations — straightforward tool calls, simple reasoning — and escalates to a stronger, more expensive model specifically when the loop detects it is stuck (Chapter 7) or when verification fails repeatedly (Chapter 8). It is high-leverage because a long debugging loop can consume many model calls, and routing every step through a frontier model overpays massively for the routine majority; matching model cost to the actual difficulty of each iteration captures most of the quality where it matters while paying cheap prices everywhere else. The reference **ModelRouter** encodes exactly this — cheap by default, strong on **stuck** or repeated verifier failures. The discipline is “route by difficulty, not by default”; the failure mode is a single fixed model for every workflow, which is either needlessly expensive (frontier everywhere) or needlessly unreliable (cheap everywhere). Routing is one of the few optimisations that improves cost without sacrificing quality on the hard steps.

Answer 15.3

How do caching and prompt caching reduce cost and latency, and what are their limits?

Result caching stores tool results and intermediate computations likely to be reused across nearby iterations or across runs against the same resource, avoiding repeated work. Prompt caching, where the model API supports it, structures the prompt so the large, stable portion (system instructions, project skills) is cached across iterations and only the smaller changing portion is re-processed each call — cutting both token cost and latency on every iteration after the first. The limits: result caching requires correct invalidation (a stale cache hit is worse than a miss), and prompt caching requires API support and a prompt structured to keep the stable prefix actually stable — if the cached prefix changes every call, you get no benefit. Both reward designing the context so the expensive, reusable part is isolated from the cheap, changing part, which is also exactly what context pruning and scratchpad isolation (Chapter 4) encourage.

Answer 15.4

Why is parallelization not a free optimization?

Because it trades increased total cost for reduced wall-clock latency — it is a latency-versus-cost trade, not a pure win. Running genuinely independent sub-tasks (Chapter 10) in parallel returns a result faster, but it generally increases total spend because you are issuing more concurrent model calls than a sequential run would at any instant; you do not save tokens, you just spend them at

once. So you parallelize when the latency benefit justifies the added spend — an interactive task where a user is waiting, or a deadline — and you do not when throughput and cost matter more than wall-clock time, where sequential execution is cheaper for the same work. Confusing parallelization with a cost optimization is a common mistake; it optimizes time, and time is sometimes worth more than the marginal tokens, but the tokens are still spent.

Answer 15.5

How does context pruning directly reduce cost, and what is the risk of over-pruning?

Every token carried forward into a new iteration is paid for again on that iteration, and the next, for as long as it stays in context — so keeping the context small directly reduces per-iteration cost, and the savings compound over a long loop. The state-compaction techniques from Chapter 4 (rolling summaries, sliding windows, scratchpad isolation) are the mechanism: keep the context sent on each call as small as it can be while remaining sufficient. The risk of over-pruning is dropping state the loop actually needs — a constraint, a prior failed attempt, an earlier decision — which causes the loop to violate something it “knew” or to repeat a failure whose record was pruned away. So pruning is a fidelity-versus-cost trade: prune aggressively enough to keep cost flat, but preserve the substance (especially failed-attempt signatures) that prevents the loop from spinning. The art is compacting noise while keeping signal.

Answer 15.6

Budget exhaustion is a cost control. What is its second, more valuable purpose?

A tool-call or token budget caps spend, but an agent exhausting its budget *without making progress* is itself a failure signal, not merely a stopping point — it is evidence the current strategy is wrong. The more valuable use is to treat budget exhaustion as a trigger to escalate (to a stronger model or a human) or to change strategy, the way Chapter 3 establishes for loop budgets generally, rather than as the mechanical end of a run that gets logged and forgotten. The biggest cost mistake in early deployments is precisely missing this: the budget caps correctly, the loop stops correctly, and then the same stuck task is resubmitted and re-runs the entire expensive loop from scratch with no memory that this exact failure already happened. The cheap fix is to persist the failed-attempt record (Chapter 4) so the next attempt starts smarter, not from zero. Budget exhaustion should make the system learn, not just stop.

Answer 15.7

Walk through the components of a cost-aware production loop.

It combines several techniques at once. A cheap, fast model handles routine iterations, with escalation to a stronger model only when the loop is stuck or verification repeatedly fails (difficulty routing). Context is kept compact through rolling summarisation so per-iteration token cost stays roughly flat instead of growing. Prompt caching keeps the stable context prefix from being re-processed each call. Independent sub-tasks run in parallel where the latency benefit justifies the added spend, not by default. And budget exhaustion without progress escalates to a stronger model or a human — and persists the failed-attempt record — rather than spending at the same rate indefinitely or re-running from scratch later. The reference project assembles these from `ModelRouter`, `CostMeter`, the Chapter 4 scratchpad, and the Chapter 3 budgets. No single lever dominates; cost-awareness is the composition of routing, compaction, caching, selective parallelism, and budget-as-signal.

Answer 15.8

Production case: a loop’s cost per task tripled month-over-month while success rate held steady. How do you find and fix it?

Steady success with rising cost points to growing inefficiency in the *path*, not the outcome — exactly what per-run cost accounting and trajectory evaluation (Chapter 13) are for. I would attribute cost per run, per model, and per task category to find which category or loop pattern inflated; check trajectory length for successful runs (are they taking more steps to reach the same answer?); and look for runaway-near-misses (runs hitting high iteration counts before barely succeeding). Common causes: context bloat raising per-iteration token cost (fix with compaction and prompt caching), a regression that made some category thrash before converging (fix the verifier or the prompt), or a model/router change sending routine steps to a frontier model (fix routing). The reference **CostMeter** per-model breakdown usually localises it fast. The key is that a stable success rate hides path-level cost regressions, so you measure the trajectory and the per-category cost, not just the outcome.

Answer 15.9

When does adding a verification loop save money rather than cost money, despite the extra calls?

When the cost of shipping a wrong output exceeds the cost of the extra verification calls, which is most production cases. Verification adds a model call (or more, on retries) and some tool calls per run, but it prevents the far larger downstream cost of a wrong output reaching production — remediation, a broken deploy, lost trust, a human cleaning it up. So the right comparison is not cost-per-attempt but cost-per-*successful, correct* output, which folds in the avoided cost of failures the verifier catches. Verification saves money on high-stakes, hard-to-reverse, or expensive-to-remediate work; it costs money with no offsetting benefit on cheap, easily-reversible, low-stakes work, which is why Chapter 8 says reserve it for where being wrong is more expensive than checking. The economic framing is that verification is insurance — worth its premium exactly when the loss it prevents is large relative to the premium.

Answer 15.10

How do you set token and tool-call budgets to balance cost against task completion?

Empirically, from the eval distribution, not by guessing. Look at the token and iteration histograms for *successful* runs and set the budget above the 95th percentile of successful completions — high enough not to truncate legitimately hard tasks, low enough that a non-converging loop does not waste serious money before stopping. Pair the budget with a no-progress detector that fires earlier on the specific pathology of a stuck loop, so you are not paying for eighteen wasted iterations before the ceiling hits. Tune per task class, because a quick copy edit and a long debugging session have very different legitimate budgets. And make budget exhaustion a signal that escalates and persists the failed attempt, so the spend buys learning rather than a silent stop. The balance is: generous enough that the budget rarely truncates a real success, tight enough that runaway cost is bounded, and instrumented so exhaustion improves the next attempt instead of repeating this one.

Loop Engineering as an Organizational Discipline

Subscribe → aiengineeringinsider.substack.com/subscribe

Everything that made your loop safe lived in your head. Scale is what forces it onto paper.

One engineer with a great loop is a productivity story. A hundred engineers with a hundred personal loops, each carrying its conventions in one person's head, is an incident waiting to be scheduled. Everything that made a loop safe when you ran it — your judgment about when to stop it, your sense of which commands were dangerous, your memory of why the team banned that pattern last year — evaporates the moment the loop belongs to everyone and no one. The last discipline of loop engineering is not technical. It is organisational: making the implicit explicit before scale forces the issue.

Everything so far in this book has been about designing one loop well. This final chapter is about what changes when an entire team adopts loops across many workflows, many repositories, and many people — because what works for one engineer's personal automation does not automatically scale to a team standard.

16.1 Repository Instructions

A repository-level instructions file — conventions, validation commands, architectural decisions, and the specific reasons certain patterns are avoided after past incidents — is the team-scale version of the skills concept introduced in Chapter 4. The difference in scale matters: an individual skill serves one engineer's personal loops; a repository instruction file serves every loop, run by every engineer and every automated trigger, that touches that codebase.

✓ Best Practice

Write repository instructions as if every future loop — run by a colleague, by an automated trigger, by a hill-climbing-generated harness change — will read them cold, with no other context. Vague instructions degrade gracefully for an experienced human; they degrade far less gracefully for an automated loop with no surrounding judgment to compensate.

16.2 Branch and Pull Request Rules for Agent-Generated Changes

Teams running loops at scale generally need explicit rules distinguishing agent-generated changes from human-authored ones: which branches a loop may write to directly, whether agent-authored pull requests require an additional reviewer beyond what a human-authored change would need, and how agent-generated commits are labeled so reviewers know what they are looking at before they start reading the diff.

16.3 Safe Defaults Around Destructive Commands and Credentials

! Pitfall / Warning

Individual developers can rely on personal judgment to stop or redirect a loop before it does something destructive. Teams cannot rely on that same judgment being applied consistently across every engineer, every repository, and every automated trigger. Team-level loop governance needs safe defaults — destructive commands blocked by default, credentials never directly accessible to a loop — that do not depend on every individual remembering to apply the same caution every time.

16.4 Standard, Reusable Loops as Organizational Assets

Once a particular loop design proves reliable for a recurring task — a dependency upgrade, an incident response, pull-request comment resolution, UI polish, test repair — it is worth converting that working design into a standardized, shared pattern, rather than leaving every engineer to reconstruct prompts, validation commands, and stopping rules from scratch each time they hit a similar task.

★ Key Insight

Reusable loops reduce onboarding friction specifically because a new team member does not need to reinvent prompts, validation commands, and stopping rules for tasks the team has already solved. The first time a team builds a reliable test-repair loop, that is one engineer's win. The moment it becomes a standard, reusable pattern, it is the team's asset.

16.5 CI Checks Both Agents and Humans Trust

A continuous-integration check that an automated loop treats as authoritative ground truth needs to be at least as reliable as the check a human reviewer trusts. A flaky test suite is a tolerable annoyance for a human who can apply judgment about which failures to ignore; it is a serious problem for a loop using that same suite as its verification signal (Chapter 8), because the loop has no equivalent judgment to fall back on and will either chase a phantom failure or learn to ignore real ones.

16.6 Clear Ownership for Final Decisions

Ownership Boundary

An explicit statement of which decisions remain with a human regardless of how capable or reliable the surrounding loops become — typically product direction, architectural tradeoffs, and anything where the cost of being systematically wrong is high and difficult to reverse.

The goal of team-level loop governance is not to remove humans from engineering work. It is to let humans spend more time on judgment, architecture, product quality, and review, while loops handle more of the repetitive mechanics — which only works if the boundary between the two is explicit and consistently enforced, rather than left to drift informally over time.

16.7 Capturing Reusable Loops

✓ Best Practice

When a loop design works well, capture it deliberately as a named, documented pattern before moving on to the next task. The capture step is easy to skip under deadline pressure, and it is exactly the step that determines whether the team compounds its loop-engineering investment over time or re-derives the same working design repeatedly from scratch.

16.8 The Loop Engineering Maturity Model

It helps to name where a team actually is. Most organisations climb the same five rungs, and most pain comes from operating at one level while believing they are at the next — shipping autonomous loops (Level 3) without the verification of Level 2, or scaling parallel agents without the human boundary that makes Level 3 safe.




Figure 16.1: The Loop Engineering Maturity Model. Each rung adds one of the four surfaces; skipping a rung is where production incidents come from.

Level 0 — One-shot prompting

No loop. A human prompts, reads, and re-prompts. Fine for isolated answers; hits a ceiling the moment a task needs to verify itself.

Level 1 — Agent loop

Observe–act–observe with tools, but no real verifier and few budgets. Useful and dangerous: it will confidently ship wrong output because nothing contradicts it.

Level 2 — Verified loop

An external verifier and explicit budgets wrap the agent loop. This is the first level you can trust with anything that matters.

Level 3 — Scaled & autonomous

Multi-agent coordination, event-driven triggers, and observability, with the human boundary and isolation that make unattended scale safe.

Level 4 — Self-improving discipline

A hill-climbing loop improves the harness from aggregated traces, governed by team-wide standards and human-reviewed changes. The loop now compounds.

16.9 The Future of Loop Engineering

As coding agents and broader agentic systems continue to improve, the most valuable individual skill is unlikely to remain writing a precisely tuned single prompt. It is far more likely to be designing reliable systems of context, action, observation, and correction — the full discipline this book has covered, from the five-stage anatomy in Chapter 1 through the organizational practices in this final chapter.

Teams will increasingly evaluate the tools and platforms they adopt not purely on the quality of the underlying model, but on how well those tools support safe iteration, repository-level awareness, reliable verification, meaningful human review, and team-wide governance — the full stack this book has laid out, not just the model at its center.

From the Field

I think this is a genuine preview of how engineering work is going to evolve, and I remain genuinely glad to be building this way. But the honest caveat is this: if I were not reviewing the output myself, or if I relied entirely on automated loops to catch what I missed, my own work’s quality would suffer, predictably and progressively. Build the loop. Build it like someone who intends to remain the engineer, not merely the person who presses go.

16.10 From Personal Script to Team Standard




Figure 16.2: Team-scale loop governance: shared repository instructions and reusable loop patterns feed every team loop, gated by safe defaults and an explicit human ownership boundary.

16.11 Individual Skill vs. Repository Instruction

Dimension	Individual skill (Ch. 4)	Repository instruction (Ch. 16)
Scope	One engineer's loops	Every loop touching the repo
Readers	A human + their agent	All engineers, all automated triggers
Failure of vagueness	Degrades gracefully (human fills gaps)	Degrades hard (no surrounding judgment)
Maintenance	Personal	A shared, owned team asset

Table 16.1: Repository instructions are the team-scale skill — and because automated loops read them cold, vagueness costs far more than it does for a human reader.

16.12 Reference Implementation: Reusable Loops as Org Assets

A proven loop design should become a named, documented pattern rather than something each engineer re-derives. The reference project's `loop_engineering_lab/patterns/catalog.py` keeps a registry of builders — the code-level form of “a team standard, not a personal script.”

`loop_engineering_lab/patterns/catalog.py` — capturing a loop as a shared asset

```
# Each named pattern is the same engine wired to a different feedback signal.
PATTERNS = {
    "retry": retry_loop,
    "test_driven": test_driven_loop,
    "plan_execute_verify": plan_execute_verify_loop,
}

# Capturing a proven loop is registering it once, so the whole team reuses it
# instead of reconstructing prompts, validation commands, and stopping rules.
def register_pattern(name: str, builder) -> None:
    if name in PATTERNS:
        raise ValueError(f"pattern '{name}' already captured")
    PATTERNS[name] = builder # now an organizational asset, not a one-off
```

The point is organisational, not technical: the first time a team builds a reliable test-repair loop it is one engineer's win; the moment it is captured as a named, reusable entry every engineer invokes, it becomes the team's compounding asset. The capture step is what turns a personal script into a standard.

16.13 Exercises

These exercises consolidate the chapter and are written for self-study and reference rather than rote drill. Each states a prompt and a hint; the full worked answers appear in the *Answers* section that follows.

Exercise 16.1

What changes when loops move from one engineer's tool to a team-wide practice?

Hint. The implicit context and judgment must become explicit and enforced.

Exercise 16.2

Why must repository instructions be written for an automated reader, not just a human?

Hint. An automated reader takes vague instructions literally.

Exercise 16.3

What explicit rules do teams need for agent-generated changes, and why?

Hint. Branch limits, an extra reviewer, and provenance labelling.

Exercise 16.4

Why can't team-level safety rely on individual judgment, and what replaces it?

Hint. Consistency comes from safe defaults, not individual memory.

Exercise 16.5

What makes a reusable loop an organizational asset, and what is the cost of not capturing it?

Hint. Capture it once as a named pattern, or pay to re-derive it each time.

Exercise 16.6

Why must CI checks a loop trusts be at least as reliable as the checks a human reviewer trusts?

Hint. The loop has no judgment to discount a flaky signal.

Exercise 16.7

What is an ownership boundary, and which decisions belong on the human side of it?

Hint. High-stakes, hard-to-reverse decisions stay with a human.

Exercise 16.8

Production case: a team adopts loops widely; six months later code quality has quietly dropped despite green CI. What governance gaps explain it?

Hint. Thin instructions, trusted-but-weak CI, and unscaled review.

Exercise 16.9

What do you think becomes the most valuable engineering skill as agentic systems improve, and why?

Hint. Designing the system around the model, not the single prompt.

Exercise 16.10

How do you keep an organization from sliding into cognitive surrender as loops become more reliable?

Hint. Make review-for-understanding and ownership structural, not optional.

16.14 Chapter Summary

- Repository instructions are the team-scale version of the individual skill concept from Chapter 4, and need to be written for an automated reader, not just a human one.
- Explicit branch and pull-request rules for agent-generated changes prevent ambiguity about review requirements and provenance.
- Safe defaults around destructive commands and credentials matter precisely because team-wide consistency cannot rely on individual judgment alone.
- Converting a proven loop design into a standardized, shared pattern compounds the team's investment instead of leaving each engineer to rebuild it independently.
- CI checks used as a loop's verification signal need to be at least as reliable as the checks a human reviewer already trusts.
- An explicit ownership boundary — which decisions stay with a human regardless of loop capability — is what makes scaled automation safe rather than merely fast.
- The durable skill going forward is designing the system around the model, not perfecting a single prompt into it.

This concludes the main text. The appendices that follow provide a condensed pattern reference, a glossary of named failure modes, and a conceptual mapping to common agentic frameworks, for use as ongoing references after a first full read.

16.15 Answers

Worked answers to the exercises above. Each restates its prompt for convenient reference.

Answer 16.1

What changes when loops move from one engineer's tool to a team-wide practice?

The thing that works for one engineer's personal automation does not automatically scale to a team standard, because the implicit context and judgment the individual supplied is no longer present. A personal loop relies on its author's knowledge of conventions, their judgment about when to stop it, and their willingness to review its output; at team scale you have many engineers, many repositories, and automated triggers, none of which share that implicit context. So the discipline shifts from "design one loop well" to governance: repository-level instructions that encode the context every loop needs, explicit rules for agent-generated changes, safe defaults that do not depend on individual caution, reusable patterns so the team compounds rather than re-derives, and clear ownership boundaries. The core realization is that everything implicit in a personal loop — context, caution, review — must be made explicit and enforced when the loop becomes a shared practice, because it can no longer rely on one person holding it in their head.

Answer 16.2

Why must repository instructions be written for an automated reader, not just a human?

Because a repository instruction file serves every loop, run by every engineer and every automated trigger, that touches the codebase — and an automated loop reads it cold, with no surrounding judgment to fill gaps. Vague instructions degrade gracefully for an experienced human, who infers intent from context and experience; they degrade far less gracefully for an automated loop, which takes them literally and has no taste to compensate for ambiguity. So you write them as if every future reader — a colleague, an automated trigger, even a hill-climbing-generated harness change — will read them with zero other context: explicit conventions, exact validation commands, architectural decisions, and the specific reasons certain patterns are banned after past incidents. This is the team-scale version of the Chapter 4 skill concept, but the stakes of vagueness are higher precisely because the reader may be a machine with no judgment to recover from an unclear instruction.

Answer 16.3

What explicit rules do teams need for agent-generated changes, and why?

Three categories. Which branches a loop may write to directly (typically not protected branches), so an agent cannot push straight to production. Whether agent-authored pull requests require an additional reviewer beyond a human-authored change, reflecting that agent output warrants a different review bar. And how agent-generated commits are labelled, so reviewers know what they are looking at before they start reading the diff — reviewing agent-authored code well requires knowing it is agent-authored. These rules exist because provenance and review requirements become ambiguous at scale: without them, agent changes blend into the normal flow and get the same (or less) scrutiny as human changes despite different risk characteristics, and reviewers cannot calibrate their attention. Making provenance and review policy explicit is what keeps agent contributions auditable and appropriately scrutinised rather than silently mixed into the codebase.

Answer 16.4

Why can't team-level safety rely on individual judgment, and what replaces it?

An individual developer can rely on personal judgment to stop or redirect a loop before it does something destructive, but a team cannot rely on that same judgment being applied consistently across every engineer, every repository, and every automated trigger — someone will forget, a new

hire will not know, an unattended trigger has no judgment at all. So team-level governance needs safe defaults that do not depend on anyone remembering to apply caution: destructive commands blocked by default, credentials never directly accessible to a loop, sensitive actions gated behind explicit approval. These are structural guarantees rather than behavioural hopes. The principle is that consistency at scale comes from defaults and enforcement, not from every individual independently exercising the same care every time — because the one time someone does not is exactly when the destructive action happens, and at team scale that one time is statistically inevitable.

Answer 16.5

What makes a reusable loop an organizational asset, and what is the cost of not capturing it?

A reusable loop becomes an asset when a proven design for a recurring task — a dependency upgrade, incident response, PR-comment resolution, test repair — is converted into a standardized, shared, documented pattern that any engineer invokes, rather than each engineer reconstructing the prompts, validation commands, and stopping rules from scratch. It reduces onboarding friction (a new member does not reinvent solved tasks) and compounds the team's investment. The cost of not capturing it is re-derivation: the same working design is rebuilt repeatedly, inconsistently, with each engineer's version drifting and each hitting the same pitfalls the first solver already solved. The capture step is easy to skip under deadline pressure and is exactly what determines whether the team compounds its loop-engineering investment or pays the same cost again every time. The reference PATTERNS registry is the minimal mechanism: capturing a loop is registering it once so it is reused, not rebuilt.

Answer 16.6

Why must CI checks a loop trusts be at least as reliable as the checks a human reviewer trusts?

Because the loop uses CI as its verification signal (Chapter 8) and has no judgment to fall back on. A flaky test suite is a tolerable annoyance for a human reviewer, who can apply judgment about which failures to ignore as noise; it is a serious problem for a loop, which either chases a phantom failure (wasting budget fixing a flake) or, worse, learns to ignore failures and so misses real ones. The loop treats the check as authoritative ground truth, so the check has to actually be authoritative. This raises the organizational bar: investing in a reliable, non-flaky CI suite stops being just good hygiene and becomes a prerequisite for safe automation, because every loop's correctness is downstream of it. A team that lets its CI rot cannot safely run verification loops against it — the verifier is only as trustworthy as the signal it is built on, and a flaky signal makes the whole loop unreliable.

Answer 16.7

What is an ownership boundary, and which decisions belong on the human side of it?

An ownership boundary is an explicit statement of which decisions remain with a human regardless of how capable or reliable the surrounding loops become — typically product direction, architectural trade-offs, and anything where the cost of being systematically wrong is high and difficult to reverse. It belongs on the human side because these are judgment calls where a loop has no basis to decide and where a wrong decision compounds expensively: a loop can implement a feature, but whether to build it, how to architect the system it lives in, and what trade-offs to accept are human calls. The boundary must be explicit and consistently enforced rather than left to drift informally, because the goal of governance is not to remove humans from engineering but to let them spend more time on judgment, architecture, and review while loops handle repetitive mechanics — which only works if everyone knows exactly which decisions are never delegated. An implicit boundary erodes silently as loops get more capable; an explicit one holds.

Answer 16.8

Production case: a team adopts loops widely; six months later code quality has quietly dropped despite green CI. What governance gaps explain it?

Several governance gaps compound into this. Repository instructions were probably thin or stale, so loops operated without the conventions and past-incident knowledge they needed, producing technically-passing but contextually-wrong changes. CI was likely treated as sufficient rather than necessary — green meant shipped, but the suite did not cover what degraded (overfitting to the verifier at org scale). Review practice did not scale with loop volume (the orchestration tax from Chapter 14), so agent changes got rubber-stamped, accumulating comprehension debt. And reusable patterns and ownership boundaries were probably informal, so quality varied by engineer and no one owned the systemic drift. The fixes are governance, not model choice: strengthen and maintain repository instructions, broaden CI so green means more, scale review-for-understanding and cap concurrency to review capacity, capture reliable patterns as standards, and enforce explicit ownership. Green CI with falling quality is the signature of trusting automation past the governance that should bound it.

Answer 16.9

What do you think becomes the most valuable engineering skill as agentic systems improve, and why?

Designing reliable systems of context, action, observation, and correction — the full loop-engineering discipline — rather than writing a single precisely-tuned prompt. As models improve, the marginal value of perfecting one prompt falls, because the model handles more of the single-shot quality itself; what does not get automatically solved is the system around the model: how work is found, context assembled, results verified, failures fed back, humans kept in the judgment seat, and the whole thing governed at team scale. That is the durable skill because it is about architecture and judgment, which compound, rather than about coaxing one good output, which the model increasingly does on its own. Teams will increasingly evaluate tools and platforms not purely on model quality but on how well they support safe iteration, repository awareness, reliable verification, meaningful human review, and team governance — the full stack, not just the model at its center. The engineer who designs the loop well remains valuable as the model inside it improves; the one who only writes prompts is automating away their own niche.

Answer 16.10

How do you keep an organization from sliding into cognitive surrender as loops become more reliable?

By making review-for-understanding and explicit ownership structural rather than optional, because the risk grows precisely as loops get better. Concretely: require genuine review of agent output — not just a glance at green CI — with the reviewer expected to understand the change, and resource that review by capping concurrency to actual review capacity (the orchestration tax). Keep ownership boundaries explicit so humans retain the high-stakes, hard-to-reverse decisions regardless of loop capability. Treat “the verifier passed” as necessary but not sufficient, and build a culture that uses loops as leverage for understood work rather than a substitute for understanding. Measure comprehension debt indirectly — incident rates on agent-heavy areas, how long root-cause takes — and respond when it rises. The deepest point is that two teams can adopt the identical tooling and diverge: one moves faster on work it understands, the other stops understanding the work, and the tooling cannot tell the difference. Governance exists to keep the organization on the first path by construction, because individual discipline alone will not hold at scale.

PART VI

Loop Engineering in Production

Theory meets the market. This bonus part reverse-engineers ten of the most influential agentic products shipping today — Devin, Claude Code, Cursor, Codex, Factory, Sweep, Replit, Gemini Code Assist, Atlan, and OpenHands — and reads each one back through the four surfaces. The differences between them are, almost entirely, differences in where the verifier lives, what grounds the loop in reality, and how a human re-enters when it is stuck.

Bonus: Ten Companies Running Loop Engineering in Production

Subscribe → aiengineeringinsider.substack.com/subscribe

What the patterns in this book look like at the scale of a shipping product.

The preceding chapters built loop engineering from first principles. This bonus chapter does the opposite: it reverse-engineers ten production systems — the most influential agentic products in the market as of this writing — and reads each one back through the vocabulary of the book. The goal is not marketing. It is to show that “observe, act, verify, adjust, repeat,” a verification contract, an explicit budget, and a human escape hatch are not academic abstractions. They are the load-bearing structure of every system below, and the differences between these products are largely differences in *where* the verifier lives, *what* grounds the loop in reality, and *how* a human re-enters when the loop is stuck.

★ Key Insight

Every system in this chapter is, at its core, the same primitive from Chapter 2 — a model calling tools in a loop until a verifier is satisfied. What distinguishes a developer toy from a production platform is everything *around* that primitive: the sandbox, the verification contract, the budget, the observability, and the human-in-the-loop boundary. Read each case study as a study in those four surfaces.

! Pitfall / Warning

Vendors iterate quickly and rarely publish full internal architectures. The descriptions below are system-design reconstructions based on public documentation, talks, product behaviour, and the engineering patterns the products clearly exhibit — not internal blueprints. Treat them as “how a competent team would build this, consistent with what is observable,” and verify specifics against current vendor documentation before depending on them.

17.1 The Ten at a Glance

Product	Loop locus	Defining loop characteristic
Cognition Devin	Cloud sandbox VM	Long-horizon autonomous dev loop with its own shell, browser, and editor
Anthropic Claude Code	Local terminal	CLI agent loop over local files and tools, human in the terminal
Cursor Agent Mode	IDE workspace	Editor-native multi-file loop fed by compiler/linter feedback
OpenAI Codex	Stateful interpreter	Write-run-observe loop inside an isolated Python/container runtime
Factory Droids	Orchestrated roles	Orchestrator/Worker/Validator loop with explicit verification contracts
Sweep AI	GitHub CI	Issue-triggered loop that iterates against the CI pipeline on a branch
Replit Agent	Cloud runtime	Concept-to-deploy loop pinging live ports for service health
Gemini Code Assist	IDE + terminal	Context-grounded loop using live repo and build environment
Atlan AI Agents	Governed metadata	Policy-wrapped loop over a data-lineage schema, not raw code
OpenHands	Sandboxed Docker	Open framework for custom dev loops with HITL escape hatches

Table 17.1: Ten production loops, distinguished by where the loop runs and what grounds it in reality.

The remaining sections treat each system in turn. For each, we cover the system-design intent, the architecture (with a diagram), the component breakdown, the verification and control surfaces, and the engineering lessons a reader can carry into their own design.

17.2 Cognition AI — Devin

System-design intent

Devin is positioned as an autonomous software engineer rather than an autocomplete assistant. The defining design decision is *environmental*: Devin operates inside its own cloud-hosted, sandboxed virtual machine equipped with a shell, a code editor, and a web browser. That choice moves Devin from the “suggest an edit” regime into the “do the whole task” regime, because the loop now has somewhere to run code, read the real `stderr`, install dependencies, and iterate without a human mediating every command. The loop’s value comes precisely from closing on real execution feedback: a failing unit test or a non-zero exit code is not an error to surface to a human, it is the next observation that drives the next action.

The most cited demonstration of this design is the work with Nubank, where Devin was applied to refactor and migrate a very large legacy ETL codebase reported in the millions of lines. A migration of that size is the ideal proving ground for a long-horizon loop: the task decomposes into thousands of mechanically similar but individually verifiable units, each of which the loop can attempt, run, check, and either accept or retry — exactly the structure a single autonomous loop with a strong verifier excels at, and exactly the structure that defeats one-shot generation.

Architecture




Figure 17.1: Devin: a long-horizon autonomous loop grounded in a sandbox VM and verified by the customer’s own test suite.

Component breakdown

The **planner** converts a high-level intent (“migrate this module to the new framework”) into a task graph of smaller, individually verifiable steps — the plan-and-execute strategy from Chapter 2, with re-planning when a step falsifies an assumption. The **agent loop** is the inner reason–act cycle that issues concrete commands. The **sandbox VM** is the grounding surface: by giving the agent a real environment rather than a simulated one, Devin ensures every observation is a true consequence of a true action, which is what makes the verification signal trustworthy. The **verifier** is the project’s own test suite and build, run after each unit of work; a failure is fed back as context rather than escalated. **Run state and memory** persist progress so a long migration can survive interruption and resume (Chapter 3’s resumability requirement at industrial scale). The **human** surface is asynchronous oversight: the engineer reviews progress, answers blocking questions, and steers, rather than driving each step.

Verification and control

The architecturally interesting property is that Devin’s verifier is the *customer’s own test suite*. That is the strongest possible verification contract because it is the same signal the customer already trusts to gate merges. The control surface is a budget over a very long horizon: the loop must distinguish “hard but progressing” from “stuck,” or a large migration becomes an unbounded spend. The no-progress detector and per-task budgets from Chapter 3 are what keep a thousand-step migration economical.

Lessons for your own design

First, ground the loop in a real environment whenever the cost allows; a sandbox that runs the actual tests turns the verification problem from “judge plausibility” into “read an exit code.” Second, decompose long-horizon work into individually verifiable units so the loop has a dense reward signal and resumability falls out naturally. Third, make the customer’s existing test suite the verifier rather than inventing a new one — adoption is far easier when the agent is gated by the same checks the team already trusts.

17.3 Anthropic — Claude Code

System-design intent

Claude Code is a command-line agent that runs where the developer already works: the local terminal, against the local filesystem and the project’s real toolchain. The design thesis is that the loop should be *thin and transparent* — the agent reads files, runs local commands (tests, linters, git), updates a running log of steps, and performs verification passes, while the human watches and intervenes in the same terminal. Rather than a remote autonomous VM, the locus is the developer’s own machine, which trades some autonomy for trust, auditability, and zero data egress.

The loop mechanics are explicit and inspectable. The agent maintains a working context of the task, decides on a tool call (read a file, run a command, edit a file), observes the result, and continues until the objective is satisfied or it needs the human. Because everything runs locally and synchronously by default, the human can see each action as it happens and stop the loop at any point — the synchronous-first principle from Chapter 2 made a product default.

Architecture




Figure 17.2: Claude Code: a transparent, local, synchronous loop over the developer’s real toolchain, with permissioned actions.

Component breakdown

The **CLI orchestrator** is the step-orchestration loop: it assembles context, presents tool schemas to the model, dispatches the chosen tool, and appends results. **Local tools** are the real project toolchain — the filesystem, shell commands, the test runner, git — so observations are genuine. The **context and step log** is the working memory of the run, and exposing it to the developer is what makes the loop auditable. The **verification pass** re-runs the relevant checks (tests, types, lint) before the agent claims completion. The **developer** is in the loop continuously: approving potentially destructive commands, answering questions, and redirecting — the destructive-action guard from Chapter 3 realised as an interactive approval prompt.

Verification and control

Claude Code’s control model is built around *permissioned action*: side-effecting or destructive commands prompt for explicit human approval, which is the interactive form of the fail-closed guard from Chapter 3. Verification is delegated to the project’s own tests and type-checks. Because the loop is local and synchronous, the budget that matters most is the human’s attention and the token cost of context assembly; effective use depends on giving the agent a tight, well-scoped task with a checkable definition of done.

Lessons for your own design

A transparent, local, synchronous loop is often the right *first* product even when full autonomy is the eventual goal — it builds user trust and is dramatically easier to debug. Make the step log a first-class, visible artifact. And route every irreversible action through an explicit approval, rather than trying to make the model “careful enough” to be trusted unsupervised.

17.4 Anyosphere — Cursor Agent Mode

System-design intent

Cursor’s Agent Mode embeds the loop inside the IDE. The defining design decision is *proximity to the editor’s own signals*: the agent executes multi-file edits directly in the local workspace, then leans on exactly the feedback the IDE already computes continuously — compiler diagnostics, language-server type errors, and linter output — to detect and fix its own mistakes in near real time. Where Devin grounds its loop in a remote sandbox and Claude Code in the terminal, Cursor grounds its loop in the editor’s live program model, which is the fastest and cheapest verification signal available for code because it is computed incrementally as files change.

This makes Cursor’s loop unusually tight. After a multi-file edit, the agent does not wait for a full test run to learn it broke something; the language server surfaces a type error within milliseconds, the agent observes it, and revises. The result is a loop whose iteration latency is dominated by model inference rather than by running an external build — a deliberate trade that favours fast, local, structural feedback.

Architecture




Figure 17.3: Cursor Agent Mode: a fast structural verifier (the language server) under a slower test tier, editing the live workspace.

Component breakdown

The **codebase index** (semantic + lexical retrieval over the repository) is what lets the agent assemble relevant context without the human pasting files — the Context stage of the five-stage anatomy, automated. The **multi-file edit engine** applies coordinated changes across files in one logical action, which matters because real changes rarely touch a single file. The **language-server / compiler** is the primary verifier: type errors and diagnostics are read directly from the IDE’s program model. The **terminal** provides the slower, deeper verification tier — actually running the code and tests — used when structural checks

are insufficient. The loop layers a fast cheap verifier (LSP) under a slow expensive one (test run), a tiered-verification pattern worth copying.

Verification and control

Cursor’s verification is explicitly tiered: structural checks first (cheap, instant, local), runtime checks second (expensive, slower). The control challenge is scope: a multi-file edit engine can sprawl across a codebase, so scope guards (Chapter 3) that keep edits within the task boundary are essential to avoid the agent “helpfully” refactoring unrelated code. Because edits land directly in the working tree, version control is the safety net — the human reviews a diff, and rejecting is a revert.

Lessons for your own design

Use the cheapest, fastest verifier you have as the inner loop and reserve expensive verifiers for the outer loop. If your domain has an equivalent of a language server — a fast static checker that catches a large fraction of errors structurally — put it on the critical path of every iteration. And treat version control as the undo stack for an agent that edits in place.

17.5 OpenAI — Codex & Code Interpreter Loops

System-design intent

OpenAI’s code-execution agents run the model inside a *stateful* execution loop: the model writes code, the system executes it inside an isolated interpreter (a sandboxed Python kernel or container), the model observes the console output, returned values, or data frames, and adjusts. The defining property is statefulness — variables, imported libraries, and intermediate results persist across turns, so the agent builds up a working environment incrementally rather than re-deriving it each step. This turns the interpreter into an extension of the model’s working memory.

This pattern is the canonical realisation of “execute to verify.” For data analysis, the loop is unbeatable: the model proposes a transformation, runs it, sees the actual shape and statistics of the result, and corrects course on real data rather than on its assumptions about the data. An exception traceback is not a failure to hide; it is the highest-signal observation the loop can receive, because it points directly at the line and cause.

Architecture




Figure 17.4: OpenAI Codex / Code Interpreter: a stateful write-run-observe loop inside an isolated interpreter.

Component breakdown

The **isolated interpreter** is a network-restricted, resource-capped sandbox — the security boundary that makes it safe to execute model-written code. The **kernel state** is the persistent namespace; its persistence is the feature, but it is also a liability, because a corrupted variable from an early step silently poisons later steps (a state-management hazard from Chapter 4). The **observation** channel normalises stdout, return values, rich objects (data frames, plots), and exceptions into the message stream the model reasons over. The control surface is a per-execution timeout and resource ceiling, plus an overall turn budget on the conversation.

Verification and control

Verification here is *execution itself*: the ground truth is whether the code ran and what it produced. The subtlety is that “it ran without error” is necessary but not sufficient — code can run and produce a wrong answer — so robust deployments add an outer verifier (assertions, expected-shape checks, or a second model reviewing the result against the goal). The sandbox’s resource limits are the budget; without them, a single `while True` or an accidental $O(n^2)$ over a large frame can hang the loop.

Lessons for your own design

A stateful interpreter is a superpower for data and analysis loops, but persistent state cuts both ways: design a way to reset or checkpoint the kernel when the loop detects it has gone down a bad path, rather than dragging corrupted state forward. And never conflate “executed cleanly” with “correct” — pair execution with an assertion-level verifier that checks the result against the goal.

17.6 Factory AI — Droids

System-design intent

Factory’s Droids platform is built around an explicit *role-separated* loop: rather than one monolithic agent, the work flows through Orchestrators that decompose and route, Workers that generate code, and Validators that enforce contracts before anything is approved. This is the multi-agent loop architecture

of Chapter 10 applied as a product thesis: separating the actor from the validator structurally, so the entity that checks the work is not the same entity (or the same context) that produced it. That separation attacks the single biggest weakness of self-grading loops — a model is a lenient judge of its own output.

The Validators enforce explicit *verification contracts*: styling, security, and policy standards expressed as checks that must pass before generated code advances. This makes the loop’s definition of “done” a property of the organisation, not of the model’s confidence — exactly the discipline Chapter 3 argues for, scaled to an enterprise where “looks good” is unacceptable and “passes the security and style contract” is the bar.

Architecture




Figure 17.5: Factory Droids: a role-separated Orchestrator–Worker–Validator loop enforcing an explicit verification contract.

Component breakdown

The **Orchestrator** owns the outer planning loop: it breaks an objective into work items, routes them to Workers, and tracks completion. **Workers** are the actors — focused agents that generate code for a single work item with their own bounded context and budget. **Validators** are dedicated verification agents and checks that run the contract: static analysis, security scanning, style enforcement, and policy rules. The **approval gate** is the human-in-the-loop boundary for changes that clear the automated contract but still warrant sign-off. The architecture is a textbook orchestrator–worker pattern with the validator promoted to a first-class role rather than an afterthought.

Verification and control

The verification contract is the heart of the design: a machine-checkable specification of what “acceptable” means, enforced by an agent that did not write the code. Control is distributed — each Worker has its own budget, and the Orchestrator has a global budget over the whole objective — which prevents one stuck work item from consuming the entire run. This is the nested-budget discipline of Chapter 3 made organisational.

Lessons for your own design

Separate the actor from the verifier whenever the stakes justify it; a validator with fresh context and an explicit contract catches errors a self-grading loop rationalises away. Express your definition of done as a machine-checkable contract, not a prompt asking the model to “make sure it is high quality.” And give each sub-agent its own budget so the failure of one cannot starve the whole system.

17.7 Sweep AI — Sweep

System-design intent

Sweep is a GitHub-native loop triggered by repository events: create or label an issue, and the agent picks it up, clones the repository on an isolated branch, modifies files, opens a pull request, and then — the defining behaviour — iterates against the project’s CI/CD pipeline, reading failure logs and refining the PR until the checks go green. The locus of the loop is the *version-control and CI platform itself*, which is a shrewd design choice: it inherits, for free, the verification, isolation, and review machinery a mature engineering org has already built.

The verifier is the existing CI suite. The isolation boundary is the feature branch. The human review surface is the standard PR review. By aligning the loop with the GitHub workflow rather than inventing a parallel one, Sweep turns “an AI changed our code” into “a contributor opened a PR that passes CI and is awaiting review” — a far lower-friction integration. The loop’s convergence criterion is simply “CI is green,” which is about as crisp and trustworthy a termination condition as exists in software.

Architecture




Figure 17.6: Sweep: a GitHub-native, issue-triggered loop that iterates against CI until the checks pass.

Component breakdown

The **event trigger** (issue creation or labelling) is what makes Sweep an *event-driven* loop in the sense of Chapter 11 — the loop starts in response to a platform event, not a human chat turn. The **isolated branch** is the sandbox; nothing the agent does touches `main` until a human merges. The **CI pipeline** is the verifier, and crucially it is the team’s own pipeline, so a green run means the same thing it means for any human contributor. The **PR review** is the human-in-the-loop gate: reviewers can request changes, and the agent treats review comments as new observations driving another iteration. **Failure-log ingestion** closes the loop — the agent parses CI output to localise and fix the failure.

Verification and control

The verification contract is “CI passes,” inherited wholesale from the repository. Control comes from the platform: branch protection, required checks, and review approval are all guards the org already configured. The loop’s own budget governs how many CI iterations it will attempt before giving up and escalating to a human — important because each CI run costs real minutes and money, so an unbounded retry against a flaky or fundamentally-failing test is a genuine cost risk.

Lessons for your own design

Trigger loops from the events of the system you already run, and inherit that system’s verification and isolation rather than rebuilding them. If your organisation has CI, it already has a world-class verifier and a review gate — wire the agent into them. And bound CI iterations explicitly, because the cheapest-looking verifier (someone else’s pipeline) still has a real per-run cost.

17.8 Replit — Replit Agent

System-design intent

Replit Agent targets the full arc from concept to a running, deployed application inside Replit’s cloud runtime. The loop does not just write code; it designs schemas, installs packages, configures databases, builds features, and — the grounding behaviour — continuously pings local ports to verify the server is actually up and healthy. The design intent is end-to-end delivery for users who may not be able to debug a stack trace themselves, which raises the bar on the loop’s autonomy and on its verification: the agent must confirm the app *works*, not merely that the code compiles.

The port-health check is the architecturally telling detail. It is a liveness probe used as a verifier: the loop takes an action (start the server), then probes the runtime to confirm the expected effect (the port responds). This is verification grounded in the running system’s observable behaviour rather than in static analysis — the right verifier when the deliverable is “a working app” rather than “correct source.”

Architecture




Figure 17.7: Replit Agent: a concept-to-deploy loop verified by a liveness probe against the running service.

Component breakdown

The **cloud runtime** provides the app process, the database, and the network surface; it is both the execution environment and the verification surface. The **environment manager** handles package installation and schema configuration — the unglamorous Context-stage work that determines whether the app can even start. The **health probe** is the verifier: a liveness check against the expected port. The **deploy** step promotes a verified build to a public URL, and the **user preview** is the human-in-the-loop signal: the user looks at the running app and gives natural-language feedback, which the loop treats as the next intent. The loop thus spans the entire SDLC, with verification at the “does it run” layer rather than the “does it compile” layer.

Verification and control

Verification is behavioural (the server responds) layered over structural (the build succeeds). The control challenge is that an end-to-end build has many places to get stuck — a missing dependency, a migration failure, a port conflict — so the loop needs per-stage verification and clear escalation when a non-code blocker (a credential, a quota) appears, because no amount of code iteration fixes a missing API key. That is the blocker-escalation stopping rule from Chapter 3, essential when the user cannot debug.

Lessons for your own design

When the deliverable is a running system, verify behaviour, not just compilation — a liveness probe is a cheap, decisive verifier. Manage the environment explicitly; many “the agent failed” incidents are really “the environment was never set up correctly.” And design clear escalation for non-code blockers, especially when your users cannot diagnose them.

17.9 Google — Gemini Code Assist & Developer Agents

System-design intent

Gemini Code Assist couples the model with the IDE and terminal, using live environmental context to discover repository files, resolve compilation failures, and update test coverage without the developer hand-feeding context through prompts. The design emphasis is on *grounding the loop in the developer’s actual environment* — the open project, the build system, the test runner — so the agent’s observations come from the real toolchain rather than from a description of it. This is the Context and Observation stages automated against live state.

The distinguishing thesis is environmental awareness at enterprise scale: large repositories, established build systems, and existing test suites. The agent must navigate a big codebase it did not write, form hypotheses about where a compilation failure originates, fix it, and confirm by rebuilding — the same observe-act-verify cycle, but stress-tested against the messiness of a mature enterprise codebase rather than a greenfield project.

Architecture




Figure 17.8: Gemini Code Assist: a context-grounded loop using live repository discovery and the real build/test environment.

Component breakdown

The **context engine** performs automatic repository discovery — finding the files, symbols, and build targets relevant to the task, so the loop is grounded in the real project graph. The **build and test environment** is the verifier: a failed compilation or a coverage gap is a concrete, machine-readable observation. The **repo index** supports retrieval over a large codebase. The **developer** remains the approver and guide, particularly important in enterprises where changes pass through code review and compliance. The architecture’s emphasis is less on novel loop mechanics and more on making the Context and Observation stages robust at the scale of a corporate monorepo.

Verification and control

Verification leans on the enterprise’s existing build and test infrastructure — the same gate human engineers face. Control includes the organisation’s review and compliance processes, into which the agent’s output flows as ordinary changes. The hard problem Gemini’s design targets is *context at scale*: in a million-file monorepo, assembling the right context cheaply is the difference between a loop that converges and one that flails, which is why the context engine is a first-class component rather than a retrieval afterthought.

Lessons for your own design

At enterprise scale, the Context stage is where loops succeed or fail — invest in robust, automatic context discovery before optimising the model. Ground verification in the existing build and test systems. And accept that in a large organisation the human review and compliance loop is part of your system whether you designed for it or not, so make the agent’s output flow cleanly into it.

17.10 Atlan — Atlan AI Agents

System-design intent

Atlan’s agents are the outlier in this chapter: the loop does not act on code, it acts on *metadata, lineage, and governance*. Built for data cataloging and governance, the loop wraps every action in a cross-platform data-lineage schema and continuously verifies metadata policies and compliance parameters, so that an agent improving a catalog cannot silently violate a governance rule or corrupt lineage. The defining design decision is that the loop’s action space is constrained by a policy model — the agent operates inside guardrails expressed as data-governance rules rather than executing freely and checking afterward.

This is a profound shift in where verification sits. In a coding loop, the agent acts and a verifier judges the result. In Atlan’s design, the policy and lineage schema are *preconditions* on the action itself: a proposed change is evaluated against governance constraints before it is applied, because in a governed data estate an incorrect metadata change can have compliance and downstream-pipeline consequences that are expensive or impossible to reverse. It is the destructive-action guard of Chapter 3 generalised into a continuous, policy-driven control plane.

Architecture




Figure 17.9: Atlan AI Agents: a policy-and-lineage-gated loop that verifies the precondition *before* acting.

Component breakdown

The **lineage schema** is the world model: a graph of how data flows across platforms, against which any proposed change is checked for downstream impact. The **policy engine** is the precondition verifier — compliance rules, access policies, and quality constraints that a candidate change must satisfy before application. The **catalog** is the system of record the agent ultimately writes to. The **data steward** is the human-in-the-loop authority for changes that require judgment or override. The loop is deliberately conservative: the verifier runs *before* the action, not after, because the cost asymmetry (a wrong governance change is far more expensive than a wrong code suggestion) justifies erring toward blocking.

Verification and control

Verification is preventive rather than corrective: policy compliance and lineage integrity are checked as guards on the action. Control is the governance model itself — the agent’s autonomy is bounded by what the policies permit, and anything outside that boundary escalates to a steward. This inverts the usual coding-loop posture (act fast, verify after) into (verify the precondition, then act), which is the correct posture whenever actions are hard to reverse and the blast radius is large.

Lessons for your own design

When actions are irreversible or high-blast-radius, move verification *before* the action as a precondition rather than after it as a judgment. Encode your domain’s constraints as an explicit policy model the loop must satisfy, not as hopeful instructions in a prompt. And maintain a world model (here, lineage) rich enough to reason about downstream impact, so the guard can catch consequences the agent would not otherwise see.

17.11 OpenHands — OpenHands Framework

System-design intent

OpenHands (formerly OpenDevin) is an open-source platform for building agentic developer loops, providing customisable, sandboxed Docker execution environments and, importantly, first-class human-in-the-loop escape hatches. Where the other nine entries are products with opinionated loops, OpenHands is a *framework* for designing your own — which makes it the most instructive entry for a loop engineer, because it exposes the seams the products hide. It is the closest production-grade analogue to the reference project in this book, scaled up.

The design thesis is configurability with safety. The sandboxed Docker runtime gives each agent an isolated, reproducible environment to act in. The HITL escape hatches let a human pause, inspect, take over, and hand back — acknowledging that no autonomous loop is reliable enough to run unsupervised on every task, so the boundary between agent and human must be a designed, first-class control rather than an afterthought. This makes OpenHands a laboratory for exactly the patterns this book formalises.

Architecture




Figure 17.10: OpenHands: an open framework modelling the loop as a typed event stream with a first-class human escape hatch.

Component breakdown

The **event stream and action API** is the orchestration backbone: agents emit typed actions (run a command, edit a file, browse), the runtime executes them, and observations flow back as events — a clean realisation of the observe-act loop with every step recorded. The **Docker sandbox** is the isolated, reproducible execution environment. The **pluggable agent** interface lets engineers swap reasoning strategies, which is the framework’s reason for existing. The **session state** supports resumability. The **HITL escape hatch** is the designed boundary where a human can intervene and resume, the safety control of Chapter 14 exposed as a core API rather than bolted on.

Verification and control

Because OpenHands is a framework, verification and control are *configurable* — you bring your own verifiers, budgets, and stopping rules, much as the reference project lets you inject a **TerminationPolicy**. The platform’s contribution is the substrate: a sandboxed runtime, a recorded event stream for observability and replay, and a HITL boundary. This is the right division of responsibility for a framework — provide the mechanism (isolation, event log, human boundary) and let the engineer supply the policy (what to verify, when to stop, what to escalate).

Lessons for your own design

Model your loop as a typed event stream of actions and observations; it makes the system replayable, debuggable, and testable, exactly as Chapter 13 argues. Make the human-in-the-loop boundary a first-class, designed API — pause, inspect, take over, resume — rather than an emergency stop. And separate mechanism from policy: a reusable loop substrate plus injectable verifiers and budgets is what lets one core serve many tasks.

17.12 Cross-Cutting Lessons

Reading the ten systems together, the same handful of design forces recur, and they map cleanly onto the chapters of this book.

Design force	How the leaders handle it	Chapter
Grounding	A real sandbox/runtime, so observations are true consequences	2, 6
Verification	The customer’s own tests, CI, types, or liveness probes	8
Verifier independence	Separate validator role or fresh-context check (Factory, Atlan)	8, 10
Budgets	Per-task and global ceilings; no-progress detection	3
Precondition guards	Verify-before-act for irreversible actions (Atlan)	3, 14
Human boundary	Designed HITL escape hatches and approval gates	14
Observability	Recorded event/step streams for replay and audit	13

Table 17.2: The ten products differ in emphasis, but every one is an instance of the same small set of loop-engineering forces.

✓ Best Practice

If you take one thing from this chapter: the winning systems did not win by having a smarter model in the loop. They won by engineering the four surfaces *around* the model — grounding, verification, budget, and the human boundary — with more care than their competitors. That is the entire thesis of loop engineering, validated at production scale.

References

Cognition AI. (2024). *Introducing Devin, the first AI software engineer*. <https://www.cognition.ai/blog/introducing-devin>

Cognition AI. (2025). *Devin & Nubank: Refactoring a large-scale codebase* [Case study]. <https://www.cognition.ai/>

Anthropic. (2025). *Claude Code: Agentic coding from the terminal* [Product documentation]. <https://docs.anthropic.com/en/docs/claude-code>

Anthropic. (2024). *Building effective agents*. <https://www.anthropic.com/research/building-effective-agents>

Anyosphere. (2024). *Cursor: The AI code editor — Agent* [Product documentation]. <https://docs.cursor.com/>

OpenAI. (2024). *Code Interpreter and tool use in the Assistants API* [Product documentation]. <https://platform.openai.com/docs/>

OpenAI. (2025). *Codex: A cloud-based software engineering agent* [Product announcement]. <https://openai.com/>

Factory AI. (2024). *Bringing autonomy to software engineering with Droids*. <https://www.factory.ai/>

- Sweep AI. (2023). *Sweep: AI-powered junior developer for bug fixes and features*. <https://docs.sweep.dev/>
- Replit. (2024). *Replit Agent documentation*. <https://docs.replit.com/replitai/agent>
- Google. (2024). *Gemini Code Assist documentation*. <https://cloud.google.com/gemini/docs/codeassist/overview>
- Atlan. (2024). *Atlan AI: The metadata control plane for the modern data stack*. <https://atlan.com/>
- Wang, X., Li, B., Song, Y., Xu, F. F., Tang, X., Zhuge, M., ... Neubig, G. (2025). OpenHands: An open platform for AI software developers as generalist agents. In *Proceedings of ICLR 2025*. <https://github.com/All-Hands-AI/OpenHands>

The Loop Engineering Pre-Launch Checklist

Subscribe → aiengineeringinsider.substack.com/subscribe

If you can tick every box, the loop is ready for production. If you cannot, you have found your next task.

Run this checklist before any loop touches production traffic. It is organised around the four surfaces from the Preface — grounding, verification, budget, and the human boundary — plus the operational concerns that decide whether a working loop stays working. Each item maps to a chapter; if an item is unclear, the chapter is the long answer.

Surface 1 — Grounding

- ☐ The loop acts in a real (or faithfully sandboxed) environment, so every observation is a true consequence of a true action. (Ch. 2, 6)
- ☐ Code execution and any side-effecting tools run inside a network-isolated, resource-capped sandbox. (Ch. 14)
- ☐ Tool results are treated as untrusted data; no single observation can widen the loop’s authority (prompt-injection defence). (Ch. 6)

Surface 2 — Verification

- ☐ “Done” is defined as a programmatically checkable condition, not “the output looks good.” (Ch. 3, 8)
- ☐ Success depends on an *external* verifier, not on the model declaring itself finished. (Ch. 8)
- ☐ The entity that grades the output is not the entity that produced it (maker/checker split). (Ch. 8)
- ☐ Verification failures are fed back as structured, actionable context for the next attempt. (Ch. 6, 8)
- ☐ The verifier covers the real, user-visible requirement — not just a narrow proxy that can be gamed. (Ch. 7, 8)

Surface 3 — Budget

- ☐ Hard ceilings are set on iterations, tokens, tool calls, wall-clock time, and dollar cost. (Ch. 3, 15)
- ☐ A no-progress detector stops a stuck loop earlier than the budget ceiling would. (Ch. 3, 7)
- ☐ Budget exhaustion escalates or changes strategy and persists the failed attempt — it does not silently re-run from zero. (Ch. 4, 15)
- ☐ Budgets are tuned per task class from the eval distribution, not guessed. (Ch. 3, 15)

Surface 4 — Human boundary

- ☐ Destructive and irreversible actions are blocked by default and require explicit, per-action human approval. (Ch. 3, 14)
- ☐ Permission and safety checks are re-evaluated against the *current* state every iteration, not once at the start. (Ch. 14)
- ☐ An external kill switch can halt the loop independently of its own logic. (Ch. 14)
- ☐ An explicit ownership boundary names the decisions that always stay with a human. (Ch. 16)
- ☐ Concurrency is bounded to what human review can actually absorb (the orchestration tax). (Ch. 14)

State, robustness, and resumability

- ☐ Working state (plan, attempt history, scratchpad) is compacted so context cost stays roughly flat. (Ch. 4)
- ☐ An attempt log is consulted *before* acting, so the loop cannot spin on a repeated failure. (Ch. 4, 7)
- ☐ Every error is classified as recoverable or hard-blocker before the loop decides how to respond. (Ch. 7)
- ☐ Side-effecting actions are idempotent, or completed effects are tracked in a durable ledger. (Ch. 6)
- ☐ Loop state is serialised at iteration boundaries so a crash resumes instead of restarting or double-acting. (Ch. 3, 4)
- ☐ Every tool call and the loop as a whole enforce a maximum duration. (Ch. 7)

Observability and economics

- ☐ Complete per-step traces record the action, the observation, *and* the reasoning. (Ch. 13)
- ☐ Run-level metrics separate success, blocked, and runaway (budget-exhausted) outcomes. (Ch. 13)
- ☐ Cost is attributable per run and per model, so expensive patterns are visible. (Ch. 15)
- ☐ The loop's own execution is monitored, so a silent unattended failure is itself alertable. (Ch. 11, 13)
- ☐ Iterations route to a cheaper model by default and escalate to a stronger one only when stuck. (Ch. 15)

Autonomy and team readiness (if applicable)

- ☐ Event handling is idempotent and concurrency-safe (per-resource locking or serialisation). (Ch. 11)
- ☐ Findings land in a reviewable inbox or pull request; the loop proposes irreversible actions rather than taking them. (Ch. 11, 14)
- ☐ Repository instructions encode conventions for an automated reader, not just a human one. (Ch. 16)
- ☐ The CI signal the loop trusts is at least as reliable as the one a human reviewer trusts. (Ch. 16)
- ☐ Any self-modifying (hill-climbing) change is evaluated offline and human-reviewed before rollout. (Ch. 12)

★ Key Insight

A loop fails in production on the surface you skipped, not the one you polished. The value of this list is not the boxes you tick easily — it is the first box you cannot tick honestly.

Loop Pattern Quick-Reference

Subscribe → aiengineeringinsider.substack.com/subscribe

All nine patterns from Chapter 9, condensed to a single reference.

Pattern	When to Use	Watch Out For
Retry Loop	Short, atomic tasks with clear pass/fail criteria	Unmodified retries with no strategy change
Plan-Execute-Verify	Multi-step, order-dependent tasks	Over-commitment to a bad plan; missing replanning trigger
Explore-Narrow	Unknown solution space, debugging, optimization	Context explosion from unpruned parallel branches
Test-Driven	Bug fixes, regression prevention	Tests too narrow to capture real requirement
Compiler-Driven	Migrations, dependency upgrades, refactors	Compiler success is not proof of product correctness
Review-Driven	Mechanical follow-through on human feedback	Treating every comment as a literal instruction
Runtime Debugging	Failures only visible at runtime	Pattern-matching to a plausible cause instead of evidence
Product Iteration	Copy, UI states, design constraints	Skipping accessibility and responsive checks
Human-in-the-Loop	High-cost-of-error or under-specified tasks	Interrupting too often, eliminating the time savings

Table 17.3: Condensed from Chapter 9. Match the pattern to the dominant feedback signal available for the task.

17.13 Matching Patterns to Feedback Signals

Available Feedback Signal	Recommended Pattern
Pass/fail test or check	Retry, Test-Driven
Type checker / compiler output	Compiler-Driven
Human review comments	Review-Driven
Runtime behavior (logs, traces)	Runtime Debugging
Unknown solution space	Explore-Narrow
Multi-step, order-dependent task	Plan-Execute-Verify
Design or product constraints	Product Iteration
Genuine ambiguity requiring judgment	Human-in-the-Loop

★ Key Insight

Most production systems compose several of these patterns in sequence rather than selecting exactly one. A typical bug-fix loop chains Runtime Debugging into Test-Driven into Review-Driven, with a Human-in-the-Loop pause before production. Use this table to identify the right starting pattern for each segment of a larger workflow.

Failure Mode Glossary

Subscribe → aiengineeringinsider.substack.com/subscribe

Named failure modes from across this book, gathered in one place.

Thrashing

The agent repeatedly changes code or state without converging on a solution. Usually caused by an unclear goal, a noisy validation signal, or editing too much at once in a single iteration. See Chapter 7.

Fix: Narrow the objective, shrink the diff size per iteration, use a more reliable observation source.

Overfitting to the Verifier

The agent satisfies its validation signal — tests pass, linter is clean — while missing the actual underlying requirement. Common when tests are too narrow or user-visible behavior is never directly checked. See Chapter 7.

Fix: Combine automated checks with requirement review and, where relevant, direct manual or browser-based verification.

Context Drift

The agent continues acting on assumptions that are now stale — missing a concurrent edit, ignoring a new failure, or following a plan that no longer matches the current state of the system. See Chapter 7.

Fix: Refresh context after every meaningful observation; never treat the initial plan as sacred.

Oscillation

The loop alternates between two or more states without converging to either, distinct from thrashing in that the pattern is cyclical rather than unbounded drift. See Chapter 13.

Fix: Detect cyclical state repetition explicitly, the same mechanism used for infinite-loop detection in Chapter 7, and force an escalation or strategy change on detection.

Runaway Loop

A loop that exhausts its entire budget without ever reaching a terminal state, often invisible in aggregate success-rate metrics because it looks similar to a deliberately bounded run. See Chapter 13.

Fix: Track budget-exhaustion rate as an explicit, separate metric from success rate.

Intent Debt

The recurring cost paid when an agent starts a session with no memory of project-specific context and fills that gap with a confident but possibly wrong guess. See Chapter 4.

Fix: Externalize project conventions and decisions into skills or repository instructions, read fresh every run.

The Orchestration Tax

The principle that a human's review bandwidth, not available tooling, is the real ceiling on how many parallel loops can be run responsibly. See Chapter 14.

Fix: Scale review practice and reviewer capacity alongside, not after, scaling parallel loop execution.

Comprehension Debt

The growing gap between what exists in a system and what the responsible engineer actually understands about it, as loop-driven work ships faster than a human can absorb and verify it. See Chapter 14.

Fix: Deliberately read and understand loop output; do not treat verifier success as equivalent to personal understanding.

Cognitive Surrender

The temptation to stop forming an independent opinion about a loop's output once the loop runs reliably, accepting whatever it returns without genuine review. See Chapter 14.

Fix: No technical fix exists; this is a discipline question for the individual engineer and the team's review culture.

Framework and Primitive Mapping

Subscribe → aiengineeringinsider.substack.com/subscribe

Conceptual mapping from this book's vocabulary to common agentic frameworks. Verify current names and behavior against each framework's own documentation before relying on specifics; tooling in this space changes quickly.

This appendix maps the concepts covered in this book to the rough shape of primitives offered by common agentic frameworks and platforms, at a conceptual level only. It deliberately omits version-specific API details, which change quickly and are better verified directly against current documentation at the time of use.

17.14 Core Concepts to Framework Primitives

This Book's Concept	Typically Appears As
Core agent loop (Ch. 2)	A model-plus-tools agent constructor; the base building block of most agent frameworks
Verification loop (Ch. 8)	Middleware or a post-run hook that scores output against a rubric and triggers a retry
Loop stacking (Ch. 12)	Graph-based orchestration where nodes can represent entire sub-loops, with conditional edges controlling re-entry
Event-driven loop (Ch. 11)	Scheduled tasks, cron triggers, and webhook-based automations, often paired with a condition-based “run until done” primitive
Worktree isolation (Ch. 10)	Per-agent or per-thread isolated working directories, sometimes a native platform feature and sometimes a manual git-worktree convention
Skills (Ch. 4)	A structured, named file of project-specific instructions and conventions, invoked by name or matched automatically by description
Sub-agents (Ch. 10)	Named, separately configured agents defined in a project-level file, each with its own instructions and optionally its own model
Trace analysis (Ch. 12)	A dedicated observability and evaluation platform that ingests run traces and supports aggregate analysis across many runs

! Pitfall / Warning

Treat this table as a map between concepts and categories of tooling, not as a guarantee that any specific named feature exists in any specific product at the time you are reading this. Agentic frameworks and platforms are evolving quickly; verify current capability against the framework's own documentation before depending on a specific primitive.

17.15 Building Without a Framework

Every concept in this book can be implemented directly, without any framework, using nothing more than an API client, a state-management approach of your own design, and the control logic from Chapters 3 and 7. The reference implementations referenced throughout this book under `loop_engineering_lab` take exactly this approach deliberately: understanding the primitives a framework abstracts away is what makes the framework's abstractions legible, rather than a black box you are simply trusting.

★ Key Insight

A framework can meaningfully accelerate building a loop. It cannot substitute for understanding what the loop is actually doing at each stage. Build at least one loop from scratch, following Chapters 2 through 8 directly against a raw model API, before adopting a framework's abstractions for that same category of work.

References

Subscribe → aiengineeringinsider.substack.com/subscribe

- Amazon Web Services. (2023). *Idempotency in event-driven architectures* [Whitepaper]. <https://docs.aws.amazon.com/lambda/latest/dg/lambda-concepts.html>
- Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete problems in AI safety. *arXiv preprint* arXiv:1606.06565.
- Anthropic. (2024). *Building effective agents*. <https://www.anthropic.com/research/building-effective-agents>
- Anthropic. (2025). *Claude Code: Agentic coding from the terminal* [Product documentation]. <https://docs.anthropic.com/en/docs/claude-code>
- Anthropic. (2025). *Claude Code best practices for agentic coding*. <https://www.anthropic.com/engineering/claude-code-best-practices>
- Anthropic. (2025). *How we built our multi-agent research system*. <https://www.anthropic.com/engineering/multi-agent-research-system>
- Anthropic. (2024). *Prompt caching with Claude* [Product documentation]. <https://docs.anthropic.com/en/docs/build-with-claude/prompt-caching>
- Bainbridge, L. (1983). Ironies of automation. *Automatica*, 19(6), 775–779. [https://doi.org/10.1016/0005-1098\(83\)90046-8](https://doi.org/10.1016/0005-1098(83)90046-8)
- Bornstein, M., & Radovanovic, R. (2024). Emerging architectures for LLM applications. *Andreessen Horowitz*. <https://a16z.com/emerging-architectures-for-llm-applications/>
- Boyd, J. R. (1987). *A discourse on winning and losing* [Briefing slides]. Air University Library, Maxwell AFB.
- Cheng, Y., Zhang, C., Zhang, Z., Meng, X., Hong, S., Li, W., . . . Liu, Y. (2024). Exploring large language model based intelligent agents: Definitions, methods, and prospects. *arXiv preprint* arXiv:2401.03428.
- Chen, L., Zaharia, M., & Zou, J. (2023). FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint* arXiv:2305.05176.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., . . . Zaremba, W. (2021). Evaluating large language models trained on code. *arXiv preprint* arXiv:2107.03374.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., . . . Schulman, J. (2021). Training verifiers to solve math word problems. *arXiv preprint* arXiv:2110.14168.
- Dann, J., & Patel, P. (2024). The economics of LLM inference at scale. *ACM Queue*, 22(3), 40–61.
- Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., & Mordatch, I. (2023). Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint* arXiv:2305.14325.
- Forsgren, N., Humble, J., & Kim, G. (2018). *Accelerate: The science of lean software and DevOps*. IT Revolution Press.
- Fowler, M. (2014). *CircuitBreaker*. [martinfowler.com. https://martinfowler.com/bliki/CircuitBreaker.html](https://martinfowler.com/bliki/CircuitBreaker.html)
- Google. (2023). *Site reliability engineering: How Google runs production systems*. <https://sre.google/books/>
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., & Fritz, M. (2023). Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec)*.

- Helmer, M., & Geihs, K. (2023). Reliability patterns for long-running autonomous processes. *Journal of Systems and Software*, 199, 111620.
- Hohpe, G., & Woolf, B. (2003). *Enterprise integration patterns: Designing, building, and deploying messaging solutions*. Addison-Wesley.
- Hong, S., Zhuge, M., Chen, J., Zheng, X., Cheng, Y., Zhang, C., ... Wu, C. (2024). MetaGPT: Meta programming for a multi-agent collaborative framework. In *Proceedings of ICLR 2024*.
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., & Narasimhan, K. (2024). SWE-bench: Can language models resolve real-world GitHub issues? In *Proceedings of ICLR 2024*.
- Khattab, O., Singhvi, A., Maheshwari, P., Zhang, Z., Santhanam, K., Vardhamanan, S., ... Potts, C. (2024). DSPy: Compiling declarative language model calls into self-improving pipelines. In *Proceedings of ICLR 2024*.
- Kilo. (2026). *What is loop engineering? AI feedback loops*. Kilo Articles.
- Kleppmann, M. (2017). *Designing data-intensive applications*. O'Reilly Media.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., ... Stoica, I. (2023). Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP)*.
- LangChain. (2026). *The art of loop engineering*. LangChain Blog.
- Liu, J., Xia, C. S., Wang, Y., & Zhang, L. (2023). Is your code generated by ChatGPT really correct? Rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 157–173. <https://doi.org/10.1162/tacl.a-00638>
- Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., ... Tang, J. (2023). AgentBench: Evaluating LLMs as agents. *arXiv preprint arXiv:2308.03688*.
- Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., ... Clark, P. (2023). Self-Refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.
- Majumder, B. P., et al. (2024). Observability and evaluation practices for LLM agents. *arXiv preprint arXiv:2407.00121*.
- MindStudio. (2026). *What is loop engineering? The new meta for AI coding agents*. MindStudio Blog.
- Nygard, M. T. (2018). *Release it! Design and deploy production-ready software* (2nd ed.). Pragmatic Bookshelf.
- Osmani, A. (2026). *Loop engineering*. addyosmani.com.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., ... Lowe, R. (2022). Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)* 35.
- Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S. G., Stoica, I., & Gonzalez, J. E. (2024). MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*.
- Park, J. S., O'Brien, J., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th ACM Symposium on User Interface Software and Technology (UIST)*.
- Patil, S. G., Zhang, T., Wang, X., & Gonzalez, J. E. (2023). Gorilla: Large language model connected with massive APIs. *arXiv preprint arXiv:2305.15334*.
- Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., ... Sun, M. (2024). ToolLLM: Facilitating large language models to master 16000+ real-world APIs. In *Proceedings of ICLR 2024*.
- Russell, S., & Norvig, P. (2021). *Artificial intelligence: A modern approach* (4th ed.). Pearson.

- Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., ... Scialom, T. (2023). Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.
- Shavit, Y., Agarwal, S., Brundage, M., Adler, S., O’Keefe, C., Campbell, R., ... Robinson, D. G. (2023). *Practices for governing agentic AI systems*. OpenAI.
- Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., & Yao, S. (2023). Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.
- Sigelman, B. H., Barroso, L. A., Burrows, M., Stephenson, P., Plakal, M., Beaver, D., ... Shanbhag, C. (2010). *Dapper, a large-scale distributed systems tracing infrastructure* (Tech. Rep.). Google Research.
- Sweep AI. (2023). *Sweep: AI-powered junior developer for bug fixes and features*. <https://docs.sweep.dev/>
- Wang, L., Ma, C., Feng, X., Zhang, Z., Yang, H., Zhang, J., ... Wen, J.-R. (2024). A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6), 186345. <https://doi.org/10.1007/s11704-024-40231-1>
- Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., ... Zhou, D. (2023). Self-consistency improves chain-of-thought reasoning in language models. In *Proceedings of ICLR 2023*.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., ... Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)* 35.
- Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., ... Wang, C. (2023). AutoGen: Enabling next-gen LLM applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155*.
- Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S., Narasimhan, K., & Press, O. (2024). SWE-agent: Agent-computer interfaces enable automated software engineering. In *Advances in Neural Information Processing Systems (NeurIPS)* 37.
- Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T. L., Cao, Y., & Narasimhan, K. (2023). Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR)*.
- Yuan, W., Pang, R. Y., Cho, K., Sukhbaatar, S., Xu, J., & Weston, J. (2024). Self-rewarding language models. *arXiv preprint arXiv:2401.10020*.
- Zelikman, E., Wu, Y., Mu, J., & Goodman, N. D. (2022). STaR: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems (NeurIPS)* 35.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., ... Stoica, I. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS)* 36.

About the Author

Subscribe → aiengineeringinsider.substack.com/subscribe

Lamhot Siagian is an AI engineer who writes about the design of production-grade agentic AI systems — loop engineering, multi-agent architectures, verification, observability, and the operational discipline of running these systems reliably at scale. He publishes under **AI Engineering Insider**.

Connect and follow for new releases:

- Substack: aiengineeringinsider.substack.com
- LinkedIn: linkedin.com/in/lamhotsiagian
- Beacons: beacons.ai/aiengineeringinsider

If this book changed how you design agentic systems, the highest compliment you can pay it is sharing it with one other engineer who is still prompting their agent one turn at a time.

- A/B testing, 110
- adjustment, 2
- agent loop, 11
- agent-generated changes, 137
- agentic grader, 66
- asynchronous execution, 11
- Atlan, 148
- audit logging, 119
- autonomous loop, 93

- batching, 128
- budget as failure signal, 128
- budget-exhaustion rate, 110

- chain (vs loop), 2
- chain-of-thought, 39
- checklist, 164
- CI checks, 137
- circuit breaker, 57
- Claude Code, 148
- code execution, 47
- Codex (OpenAI), 148
- cognitive surrender, 119
- compiler-driven loop, 75
- comprehension debt, 119, 137
- condition-based looping, 93
- context (loop stage), 2
- context drift, 57
- context pruning, 128
- context window, 30
- cost engineering, 128
- cost per iteration, 128
- cron trigger, 93
- Cursor, 148

- destructive-action guard, 20
- deterministic grader, 66
- Devin (Cognition), 148

- early exit, 20
- ephemeral state, 30
- episodic memory, 30
- error taxonomy, 57
- event-driven loop, 93

- explore-narrow loop, 75
- exponential backoff, 57
- external memory, 30

- Factory Droids, 148
- failure handling, 57
- feedback injection, 66
- feedback loop, 2
- five-stage anatomy, 2
- four surfaces, 148, 164

- Gemini Code Assist, 148
- git worktree, 85
- graceful degradation, 57
- grader, 66
- graph-of-thought, 39
- grounding, 2
- guardrail, 119

- hallucinated action, 57
- harness, 101
- heartbeat, 93
- hierarchical decomposition, 85
- hill-climbing loop, 101
- human oversight, 119
- human-in-the-loop, 75

- idempotency, 47
- idempotent event handling, 93
- intent, 2
- intent debt, 30
- isolated state, 85

- kill switch, 119

- latency optimization, 128
- LLM-as-judge, 66
- LLMClient, 11
- loop (definition), 2
- loop budget, 20
- loop engineering, 2
- loop pattern, 75
- loop stacking, 101
- loop synchronization, 85
- loop tracing, 110

- maker/checker split, 66
- mock LLM, 11
- model routing, 128
- multi-agent systems, 85

- nested loops, 20
- no-progress detector, 20

- observability, 110
- observation, 2
- observe-reason-act, 11
- OODA loop, 2
- OpenHands, 148
- orchestration tax, 119
- orchestrator-worker, 85
- organizational discipline, 137
- oscillation, 39
- overfitting to the verifier, 57
- ownership boundary, 137

- parallel tool calls, 47
- peer-to-peer agents, 85
- permission boundary, 119
- persistent state, 30
- plan revision, 39
- plan-act-verify, 2
- plan-and-execute, 11
- plan-execute-verify loop, 75
- product iteration loop, 75
- prompt caching, 128
- prompt engineering, 2
- prompt injection, 47

- ReAct, 2, 11
- reactive loop, 11
- reasoning (in loop), 39
- reflection, 39
- reinforcement learning fine-tuning, 101
- repetition detection, 57
- replay, 110
- Replit Agent, 148
- repository instructions, 137
- resumability, 20
- retry, 47
- retry loop, 75
- retry-with-modification, 57
- reusable loops, 137
- review-driven loop, 75
- reviewer agent, 85

- robustness, 57
- robustness testing, 110
- rolling summary, 30
- rubric scoring, 66
- runaway loop, 110
- runtime debugging loop, 75

- safe defaults, 137
- safety, 119
- sandbox, 47
- sandboxing, 119
- scope guard, 20
- scratchpad, 30
- self-consistency, 39
- self-critique, 39
- self-improvement, 101
- serialization, 30
- shared state, 85
- skills, 30
- sliding window, 30
- state and memory, 30
- state compaction, 30
- state machine, 11
- step-level evaluation, 110
- step-orchestration loop, 11
- stop condition, 66
- stopping rule, 20
- structured feedback, 47
- stuck-state detection, 57
- Sweep, 148
- synchronous execution, 11

- terminal states, 11
- termination condition, 20
- TerminationPolicy, 20
- test-driven loop, 75
- thrashing, 57
- time budget, 20
- time-based looping, 93
- timeout, 57
- token budget, 20
- token-generation loop, 11
- tool error handling, 47
- tool registry, 47
- tool selection, 47
- tool use, 47
- tool-call budget, 20
- tool-call cycle, 47

trace (artifact), [101](#)
trace analysis, [101](#)
trajectory evaluation, [110](#)
tree-of-thought, [39](#)
trigger reliability, [93](#)
truncation, [30](#)

verification, [2](#)
verification contract, [148](#)
verification loop, [66](#)
verifier-in-the-loop, [39](#)

webhook trigger, [93](#)
working memory, [30](#)