




AI ENGINEERING INSIDER PODCAST

Anthropic Launches Claude Fab...

0:00



20:55




[Listen now](#)

Anthropic Launches Claude Fable 5 and Mythos 5: The Next Step Toward Long-Horizon AI Agents



AI ENGINEERING INSIDER

JUN 09, 2026



Claude Fable 5 & Mythos 5

The Next Step Toward Long-Horizon AI Agents

Anthropic just announced **Claude Fable 5** and **Claude Mythos 5**, two new frontier models designed for complex, long-running AI work. The big story is not just benchmark improvement. It is that Claude is moving deeper into the era of **autonomous agents**: systems that can handle multi-step software engineering, scientific research, vision-based reasoning, and extended knowledge work with less human scaffolding.

Claude **Fable 5** is the generally available version. Anthropic says it is the most capable model they have released broadly, with strong performance across software engineering, knowledge work, vision, scientific reasoning, and long-context tasks. The company claims the

model's advantage becomes more visible as tasks become longer and more complex, which is exactly the direction enterprise AI is moving: from simple chatbots to agents that execute workflows.

The more restricted model, **Claude Mythos 5**, uses the same underlying model but has some safeguards lifted for trusted users, especially cyberdefenders, infrastructure providers, and selected research organizations. Anthropic says Mythos 5 will initially be deployed through **Project Glasswing** in collaboration with the U.S. government, with broader trusted access planned later.

	Claude Mythos 5 / Fable 5	Claude Mythos Preview	Claude Opus 4.8	GPT 5.5	Gemini 3.1 Pro
Agentic coding SWE-Bench Pro	80.3%	77.8%	69.2%	58.6%	54.2%
Agentic coding FrontierCode (Diamond)	29.3% xhigh	—	13.4% xhigh	5.7% xhigh	—
Knowledge work GDPval-AA	1932	—	1890	1769	1314
Knowledge work vision GDPpdf	29.8% no tools	—	22.5% no tools	24.9% no tools	16.7% no tools
Spatial reasoning Blueprint-Bench 2	38.6%	—	14.5%	36.2%	26.5%
Tool use AutomationBench	17.4%	—	15.5%	12.9%	9.6%
Computer use OSWorld-Verified	85.0%	85.4%	83.4%	78.7%	76.2%
Legal Legal Agent Benchmark	13.3%	—	10.4%	2.1%	0.0%
Multidisciplinary reasoning Humanity's Last Exam	59.0%* no tools	56.8% no tools	49.8% no tools	41.4% no tools	44.4% no tools
	64.5%* with tools	64.7% with tools	57.9% with tools	52.2% with tools	51.4% with tools
Biology BioMysteryBench	46.1%* hard	29.6% hard	40.0% hard	—	—
	83.9%* human solved	82.6% human solved	80.4% human solved	—	—
Agentic coding Terminal-Bench 2.1	88.0%*	—	82.7%	83.4% Codex CLI	70.7% Gemini CLI
Cybersecurity ExploitBench (Cap%)	78.0%*	69.0%	40.0%	34.0%	—
Health HealthBench Professional	66.0%*	64.7%	56.9%	51.8%	—

Methodology: Reported scores are within a 1-3 percentage point difference for Claude Mythos 5 and Claude Fable 5. This table shows the higher score of the two. Starred (*) benchmarks show a larger difference due to our blocking safeguards for cybersecurity and biology-related questions. For these benchmarks, Claude Fable 5 performs closer to Claude Opus 4.8 due to fallbacks. See the system card for details.

Why This Matters

The launch signals a major shift in frontier AI: capability and safety are now being packaged as different access tiers.

For most users, Fable 5 provides advanced reasoning with conservative safeguards. For trusted organizations, Mythos 5 unlocks more powerful dual-use capabilities in cybersecurity and biology. This model-access strategy reflects a growing industry pattern: the most powerful AI

systems may not be released in a single public version. Instead, they may be deployed through **permissioned access**, depending on user identity, use case, and risk profile.

Capabilities

Anthropic highlights several major improvements:

Software Engineering: Fable 5 is positioned as a stronger coding and agentic engineering model. Anthropic cites early testing in which Stripe reportedly used it for a large, codebase-wide migration that would have taken a team months to do by hand.

Knowledge Work: The model shows stronger document reasoning, chart interpretation, finance analysis, and problem-solving. This matters for enterprise workflows where models must analyze messy documents, spreadsheets, and business data rather than answer isolated questions.

Vision: Anthropic describes Fable 5 as state-of-the-art for vision tasks, including extracting numbers from scientific figures and rebuilding web application source code from screenshots. This is important because many real-world workflows are multimodal, involving screenshots, diagrams, PDFs, dashboards, and visual interfaces.

Memory and Long Context: Fable 5 is designed to stay focused across millions of tokens and improve its own work using notes and persistent memory. This is one of the most important directions for agentic AI because long-running tasks require continuity, planning, and self-correction.

Scientific Research: Anthropic claims Mythos 5 helped accelerate drug-design workflows and produced novel molecular biology hypotheses. They also describe autonomous genomics research in which the model operated for more than a week with limited human input. These claims are significant, but some results are still pending

publication, so they should be treated as early evidence rather than independently validated consensus.

The Safety Architecture

The most interesting part of this announcement is the safety design. Fable 5 includes classifiers that detect risky requests across cybersecurity, biology/chemistry, and model distillation. When these classifiers trigger, the user may receive a response from Claude Opus 4.8 instead of Fable 5. Anthropic says these safeguards trigger in fewer than 5% of sessions on average, though they may still produce false positives.

This is a practical compromise: instead of refusing every risky query, Anthropic routes some requests to a less capable but still useful model. For AI product teams, this is a useful design pattern: **capability routing based on risk classification**.

Pricing and Availability

Both Fable 5 and Mythos 5 are priced at **\$10 per million input tokens** and **\$50 per million output tokens**. Fable 5 is available through the Claude API as `claude-fable-5`, while Mythos 5 remains restricted to selected trusted-access users.

AI Engineering Takeaway

Claude Fable 5 and Mythos 5 show where frontier AI systems are heading:

AI is no longer only about better chat responses. The frontier is shifting toward **long-horizon execution, multimodal reasoning, tool use, memory, autonomous research, and risk-aware deployment**.

For AI engineers, this launch is a reminder that the next generation of AI platforms will need more than model APIs. They will need:

- risk classifiers
- permissioned access layers
- fallback models
- audit logging
- long-context orchestration
- memory systems
- evaluation harnesses
- human-in-the-loop review

The future AI stack is not just “call the best model.” It is **routing the right task to the right capability level with the right safety boundary.**