




▶ Watch now

Fine-Tuning & PEFT LoRA, QLoRA, RLHF, DPO, ORPO, Adapter Layers

Fine-tuning is not just “training the model more.” In production AI engineering, it is a decision about behavior adaptation, cost, latency, data quality, safety, evaluation, deployment, and rollbacks



AI ENGINEERING INSIDER

MAY 26, 2026

1. What's fine-tuning?

A strong AI engineer should explain fine-tuning like this:

- ↳ **Prompting** changes behavior at inference time.
- ↳ **RAG** adds external knowledge without changing model weights.

- ↳ **Fine-tuning** updates the model's behavior by training on task- or domain-specific examples.
- ↳ **PEFT** updates only a small subset of parameters rather than the entire model.
- ↳ **Preference optimization** aligns the model toward preferred outputs using chosen/rejected responses.

LoRA freezes the base model and trains small, low-rank matrices injected into transformer layers, reducing trainable parameters and memory usage compared to full fine-tuning. QLoRA extends this idea by training LoRA adapters using a frozen 4-bit-quantized base model, employing techniques such as NF4 quantization, double quantization, and paged optimizers to reduce memory usage.

2. Key Concepts You Must Know

Fine-Tuning

Fine-tuning means taking a pretrained model and continuing training it on a smaller, task-specific dataset.

Use fine-tuning when:

- ↳ You need a consistent response format.
- ↳ You need a domain-specific tone or style.
- ↳ You need the model to follow specialized task patterns.
- ↳ You have many high-quality examples.
- ↳ Prompting alone becomes too long, brittle, or expensive.

Do not use fine-tuning when:

- ↳ The issue is missing external knowledge. Use RAG.
- ↳ The data changes frequently. Use retrieval or tools.

- ↳ You do not have clean labels. Improve data first.
 - ↳ You need guaranteed factuality. Fine-tuning does not guarantee truth.
-

PEFT: Parameter-Efficient Fine-Tuning

PEFT means adapting a large model by training only a small number of additional or selected parameters.

Common PEFT methods:

- ↳ **LoRA** — train low-rank update matrices.
- ↳ **QLoRA** — LoRA on top of a quantized base model.
- ↳ **Adapters** — insert small trainable modules inside transformer blocks.
- ↳ **Prefix tuning / prompt tuning** — train soft prompt vectors.
- ↳ **IA3** — scale internal activations with learned vectors.

Adapter modules were introduced as a compact transfer-learning method where the original network remains fixed and only small task-specific modules are trained. The original adapter work showed near state-of-the-art performance on GLUE while adding only a small percentage of task-specific parameters.

3. LoRA Interview Explanation

What is LoRA?

LoRA stands for **Low-Rank Adaptation**.

Instead of updating the full weight matrix of the model, LoRA represents the weight update as a low-rank decomposition.

Simple explanation:

- ↳ The base model weights stay frozen.
- ↳ Small trainable matrices are added to selected layers.
- ↳ During training, only those small matrices are updated.
- ↳ During inference, LoRA weights can often be merged into the base model for efficient serving.

Senior-Level Answer

“In full fine-tuning, every parameter of the base model is updated, which is expensive and hard to serve at scale. LoRA assumes that the task-specific weight update has low intrinsic rank. So instead of learning a full dense update, we learn two smaller matrices whose product approximates the update. This dramatically reduces the number of trainable parameters, memory usage, and storage per task. It is especially useful when we need multiple task-specific variants of a large model.”

Where is LoRA usually applied?

- ↳ Query projection
- ↳ Key projection
- ↳ Value projection
- ↳ Output projection
- ↳ Feed-forward layers
- ↳ Sometimes embedding layers or lm_head, depending on task

Important LoRA Hyperparameters

Area	LoRA	Adapter Layers
Mechanism	Low-rank weight update	Extra trainable modules
Base model	Frozen	Frozen
Inference latency	Often no extra latency if merged	Can add latency
Modularity	Strong	Strong
Common use	LLM fine-tuning	Multi-task/domain adaptation
Deployment	Merge or load adapter	Load adapter modules

4. QLoRA Interview Explanation

What is QLoRA?

QLoRA is a memory-efficient method for fine-tuning large language models.

It works by:

- ↳ Loading the base model in 4-bit quantized form.
- ↳ Freezing the quantized base model.
- ↳ Training LoRA adapters on top of it.
- ↳ Backpropagating gradients through the frozen quantized model into LoRA weights.

QLoRA showed that very large models could be fine-tuned with much lower GPU memory while preserving strong performance compared with full 16-bit fine-tuning.

Senior-Level Answer

“QLoRA is useful when the base model is too large for standard fine-tuning. Instead of training all model weights, we quantize the frozen base model to 4-bit precision and train LoRA adapters on top. The key

engineering benefit is GPU memory reduction. The tradeoff is that quantization introduces additional complexity, and training or serving may require careful kernel, hardware, and precision compatibility.”

QLoRA Production Benefits

- ↳ Lower GPU memory requirement
- ↳ Fine-tune larger models on limited hardware
- ↳ Lower adapter storage cost
- ↳ Easier experimentation across many datasets
- ↳ Good fit for domain-specific instruction tuning

QLoRA Risks

- ↳ Quantization instability
 - ↳ Hardware/kernel compatibility issues
 - ↳ Slower training in some setups
 - ↳ Harder debugging than standard LoRA
 - ↳ Quality degradation if quantization is poorly configured
-

5. Adapter Layers Interview Explanation

What are adapter layers?

Adapter layers are small neural modules inserted into a pretrained transformer. The base model is frozen, and only adapter parameters are trained.

A typical adapter looks like:

- ↳ Down-projection
- ↳ Nonlinear activation
- ↳ Up-projection
- ↳ Residual connection

Senior-Level Answer

“Adapters are a PEFT technique where we insert small task-specific trainable modules into the transformer. They are modular and useful when supporting many tasks or customers on the same base model. Compared with LoRA, adapters may add inference latency because they introduce additional layers in the forward pass, while LoRA can often be merged into the original weights.”

LoRA vs Adapter Layers

Area	LoRA	Adapter Layers	Mechanism	Low-rank weight update	Extra trainable modules
Base model	Frozen	Frozen	Inference latency	Often no extra latency if merged	Can add latency
Modularity	Strong	Strong	Common use	LLM fine-tuning	Multi-task/domain adaptation
Deployment	Merge or load adapter	Load adapter modules			

6. RLHF Interview Explanation

What is RLHF?

RLHF means **Reinforcement Learning from Human Feedback**.

A common RLHF pipeline:

- ↳ Step 1: Supervised fine-tuning on high-quality demonstrations.
- ↳ Step 2: Collect human preference rankings between model outputs.

- ↳ Step 3: Train a reward model to predict human preference.
- ↳ Step 4: Optimize the policy model using reinforcement learning, often PPO.
- ↳ Step 5: Evaluate helpfulness, harmlessness, truthfulness, and regression risk.

The InstructGPT work used human demonstrations and preference rankings to align GPT-style models with user intent, showing that fine-tuned models with human feedback could be preferred over much larger base models in human evaluations. PPO is a policy gradient method widely used in reinforcement learning and is often paired with classical RLHF pipelines.

Senior-Level Answer

“RLHF is a post-training alignment pipeline. We first train the model to imitate high-quality answers using supervised fine-tuning. Then we collect human rankings of candidate responses and train a reward model. Finally, we optimize the language model against that reward model while constraining it not to drift too far from the original policy. The main challenge is that RLHF is expensive, unstable, and sensitive to reward model quality.”

RLHF Failure Modes

- ↳ Reward hacking
 - ↳ Over-optimization against weak reward model
 - ↳ Loss of general capability
 - ↳ Verbosity bias
 - ↳ Safety over-refusal
 - ↳ Labeler disagreement
 - ↳ High cost of preference data
 - ↳ Training instability
-

7. DPO Interview Explanation

What is DPO?

DPO means **Direct Preference Optimization**.

DPO trains the model directly on preference pairs:

- ↳ Prompt
- ↳ Chosen response
- ↳ Rejected response

Unlike classical RLHF, DPO does not require a separate reward model or reinforcement learning loop. The DPO paper frames the language model itself as implicitly representing the reward model and uses a simpler classification-style objective.

Senior-Level Answer

“DPO is a simpler alternative to PPO-style RLHF. Instead of training a reward model and then running reinforcement learning, DPO directly optimizes the policy using preference pairs. It compares the probability assigned to the chosen response versus the rejected response, while using a reference model to control drift. In production, DPO is attractive because it is simpler, more stable, and easier to operationalize than RLHF.”

DPO Benefits

- ↳ Simpler than RLHF
- ↳ No separate reward model
- ↳ No online RL sampling loop
- ↳ More stable training
- ↳ Easier to reproduce
- ↳ Good fit for preference-alignment datasets

DPO Risks

- ↳ Requires high-quality preference pairs
 - ↳ Sensitive to noisy chosen/rejected labels
 - ↳ Can overfit preference style
 - ↳ May reduce diversity
 - ↳ Reference model selection matters
 - ↳ Does not solve factual grounding by itself
-

8. ORPO Interview Explanation

What is ORPO?

ORPO means **Odds Ratio Preference Optimization**.

ORPO is a reference-model-free preference optimization method. It combines supervised fine-tuning behavior with a preference optimization signal, using an odds-ratio term to increase the likelihood of chosen responses and penalize rejected responses.

Senior-Level Answer

“ORPO is designed to simplify preference alignment further. Unlike DPO, which typically compares the policy model against a reference model, ORPO is reference-model-free. It combines the supervised fine-tuning objective with a preference penalty in one training phase. The engineering benefit is lower memory and compute overhead because we do not need to load a separate reference model.”

ORPO Benefits

- ↳ Reference-model-free
- ↳ Lower memory overhead than DPO

- ↳ Combines SFT and preference optimization
- ↳ Simpler training pipeline
- ↳ Useful when compute is constrained

ORPO Risks

- ↳ Still needs strong preference data
- ↳ Less mature operational history than SFT or LoRA
- ↳ Can be sensitive to hyperparameters
- ↳ May over-align to narrow preference patterns
- ↳ Needs careful regression testing

9. LoRA vs QLoRA vs RLHF vs DPO vs ORPO

Method	Main Goal	Data Type	Trains What?	Best For	
Full fine-tuning	Task/domain adaptation	Input-output examples	All model weights	Maximum control, high compute	
LoRA	Efficient adaptation	Input-output examples	Low-rank adapters	Cost-efficient LLM tuning	
QLoRA	Memory-efficient LoRA	Input-output examples	LoRA on quantized base	Large models on limited GPU	
Adapters	Modular task adaptation	Input-output examples	Inserted modules	Multi-task or multi-tenant systems	
RLHF	Human preference alignment	Rankings/preferences	Reward model + policy	Complex alignment goals	
DPO	Simpler preference tuning	Chosen/rejected pairs	Policy model	Stable preference optimization	
ORPO	Reference-free preference tuning	Chosen/rejected pairs	Policy model	Lower-memory preference alignment	

10. Mock Interview Q&A

Q1. What is the difference between fine-tuning and RAG?

Strong Answer:

“Fine-tuning changes the model’s parameters, while RAG changes the context given to the model at inference time. I use fine-tuning when I want the model to learn a behavior, format, style, or task pattern. I use RAG when the model needs access to external or frequently changing knowledge. For example, if the model needs to answer from company policy documents, I would use RAG. If the model needs to consistently classify support tickets into internal categories, I may use fine-tuning.”

Follow-up:

What if the model gives outdated answers?

Answer:

“That is usually not a fine-tuning problem. I would use retrieval, tools, or database access because the knowledge changes over time.”

Q2. Why is PEFT important?

Strong Answer:

“PEFT is important because full fine-tuning large models is expensive in GPU memory, storage, and deployment complexity. With PEFT, we freeze the base model and train a small number of task-specific parameters. This allows faster experimentation, cheaper training, and easier multi-tenant deployment where each customer or task can have its own adapter.”

Q3. Explain LoRA to a non-expert.

Strong Answer:

“Imagine the base model is a very large expert. Instead of retraining the entire expert, LoRA adds small trainable correction layers that teach the expert how to behave for a specific task. We keep the original knowledge intact and only learn the small differences needed for the task.”

Q4. Explain LoRA technically.

Strong Answer:

“LoRA freezes the original model weights and learns a low-rank approximation of the weight update. Instead of updating a large dense matrix, it learns two smaller matrices. This reduces trainable parameters significantly and allows task-specific adaptation with lower memory and storage cost. In deployment, LoRA weights can often be merged into the base weights to avoid additional inference latency.”

Q5. When would you choose QLoRA over LoRA?

Strong Answer:

“I would choose QLoRA when the model is too large to fine-tune with normal precision on available hardware. QLoRA loads the base model in 4-bit quantized form and trains LoRA adapters on top. It is especially useful for experimenting with 13B, 34B, or larger models on constrained GPU memory. I would still evaluate whether quantization affects quality, latency, and production compatibility.”

Q6. What is the difference between SFT and preference tuning?

Strong Answer:

“SFT teaches the model to imitate ideal responses. Preference tuning teaches the model to prefer one response over another. SFT uses input-output examples. Preference optimization uses chosen/rejected pairs or ranked outputs. In practice, I often use SFT first to teach the task and then DPO or ORPO to improve answer quality, helpfulness, conciseness, or safety behavior.”

Q7. What is RLHF, and why is it hard?

Strong Answer:

“RLHF uses human feedback to align model behavior. The classical pipeline includes supervised fine-tuning, reward model training, and policy optimization using reinforcement learning. It is hard because it requires expensive preference data, stable reward modeling, careful PPO tuning, and strong evaluation to prevent reward hacking or capability regression.”

Q8. Why did DPO become popular?

Strong Answer:

“DPO became popular because it simplifies preference alignment. Instead of training a separate reward model and running PPO, DPO directly optimizes the model using chosen and rejected response pairs. It is easier to implement, more stable, and cheaper to operate than classical RLHF.”

Q9. What is the difference between DPO and ORPO?

Strong Answer:

“DPO usually uses a reference model to prevent the policy from drifting too far from the original model. ORPO removes the need for a separate reference model and combines supervised fine-tuning with a preference optimization signal. So ORPO can reduce memory and compute requirements, while DPO is more established and widely used in preference-alignment pipelines.”

Q10. What data do you need for fine-tuning?

Strong Answer:

“For SFT, I need high-quality prompt-response examples. For DPO or ORPO, I need preference pairs: a prompt, a chosen answer, and a rejected answer. The most important part is not just quantity but quality, diversity, and label consistency. Bad fine-tuning data can make the model worse than the base model.”

11. Senior System Design Interview Scenario

Prompt


“You need to fine-tune an LLM for an enterprise customer-support assistant. The model must answer in a consistent tone, follow company escalation rules, avoid hallucination, and support multiple customers.”

Strong Architecture Answer

“I would not start with fine-tuning immediately. First, I would separate the problem into knowledge, behavior, and alignment.”

- ↳ **Knowledge layer:** Use RAG for policy documents, product manuals, FAQs, and changing customer data.
- ↳ **Behavior layer:** Use SFT or LoRA to teach response format, tone, escalation behavior, and ticket classification patterns.
- ↳ **Preference layer:** Use DPO or ORPO on chosen/rejected examples to improve helpfulness, conciseness, and refusal behavior.
- ↳ **Safety layer:** Add output validation, policy checks, PII detection, and escalation rules.
- ↳ **Evaluation layer:** Use offline evals, human review, regression tests, hallucination tests, and production monitoring.
- ↳ **Deployment layer:** Serve one base model with customer-specific LoRA adapters or adapter routing.
- ↳ **Rollback layer:** Version every adapter, dataset, prompt template, and eval report.

Production Design



Base LLM + Customer LoRA Adapter
↓
Output Validator
↓
Citation / Grounding Check
↓
Escalation Decision
↓
Final Response
↓
Logging + Evaluation + Feedback Collection

What I Would Fine-Tune

- ↳ Ticket classification
- ↳ Response tone
- ↳ Escalation format
- ↳ Tool-use format
- ↳ Customer-service answer style
- ↳ Refusal and safe-completion behavior

What I Would Not Fine-Tune

- ↳ Frequently changing policy facts
- ↳ Real-time order status
- ↳ Inventory data

- ↳ Pricing
 - ↳ User-specific private information
 - ↳ Anything better handled by retrieval or tools
-

12. Evaluation Strategy

A strong answer should mention both model quality and operational quality.

Offline Evaluation

- ↳ Exact match for structured outputs
- ↳ Classification accuracy
- ↳ Preference win rate
- ↳ Human evaluation
- ↳ Safety violation rate
- ↳ Hallucination rate
- ↳ Groundedness score
- ↳ Refusal correctness
- ↳ Format compliance
- ↳ Regression against the base model

Online Evaluation

- ↳ User satisfaction
- ↳ Escalation rate
- ↳ Reopen rate
- ↳ Average handling time
- ↳ Deflection rate
- ↳ Human override rate

- ↳ Latency
 - ↳ Cost per request
 - ↳ Adapter-specific error rate
-

13. Common Interview Traps

Trap 1: “Fine-tuning will add new knowledge.”

Better answer:

“Fine-tuning can teach patterns, but it is not the best mechanism for dynamic factual knowledge. For changing knowledge, I would use RAG or tools.”

Trap 2: “More data is always better.”

Better answer:

“Fine-tuning quality depends heavily on data quality. A smaller clean dataset can outperform a larger noisy dataset.”

Trap 3: “RLHF always improves the model.”

Better answer:

“RLHF can improve alignment, but poor reward models can create reward hacking, over-refusal, or capability regression.”

Trap 4: “QLoRA is always better than LoRA.”

Better answer:

“QLoRA is better when memory is the bottleneck. If I have enough GPU memory and want simpler training, standard LoRA may be easier.”

Trap 5: “DPO solves hallucination.”

Better answer:

“DPO improves preference alignment, not factual grounding. For hallucination, I still need retrieval, citations, verification, and evals.”

14. Final Interview Cheat Sheet

Use this answer in interviews:

“Fine-tuning is best for changing model behavior, not for storing dynamic knowledge. For large LLMs, I usually prefer PEFT methods like LoRA or QLoRA because they reduce memory, storage, and deployment cost. LoRA trains low-rank adapters while freezing the base model. QLoRA further reduces memory by using a 4-bit quantized base model with LoRA adapters. For alignment, RLHF is powerful but complex, requiring reward modeling and reinforcement learning. DPO simplifies preference optimization by training directly on chosen/rejected pairs, while ORPO goes further by removing the reference model and combining SFT with preference optimization. In production, I would combine RAG for knowledge, LoRA or QLoRA for behavior adaptation, DPO or ORPO for preference alignment, and strong evals for safety, hallucination, regression, and cost.”